

APS1070

Foundations of Data Analytics and
Machine Learning

Winter 2022

Week 7:

- *Projections*
- *Matrix Decompositions*
- *Principal Component Analysis (PCA)*



Slide Attribution

These slides contain materials from various sources. Special thanks to the following authors:

- Scott Sanner
- Ali Hadi Zadeh
- Jason Riordon

Last Time

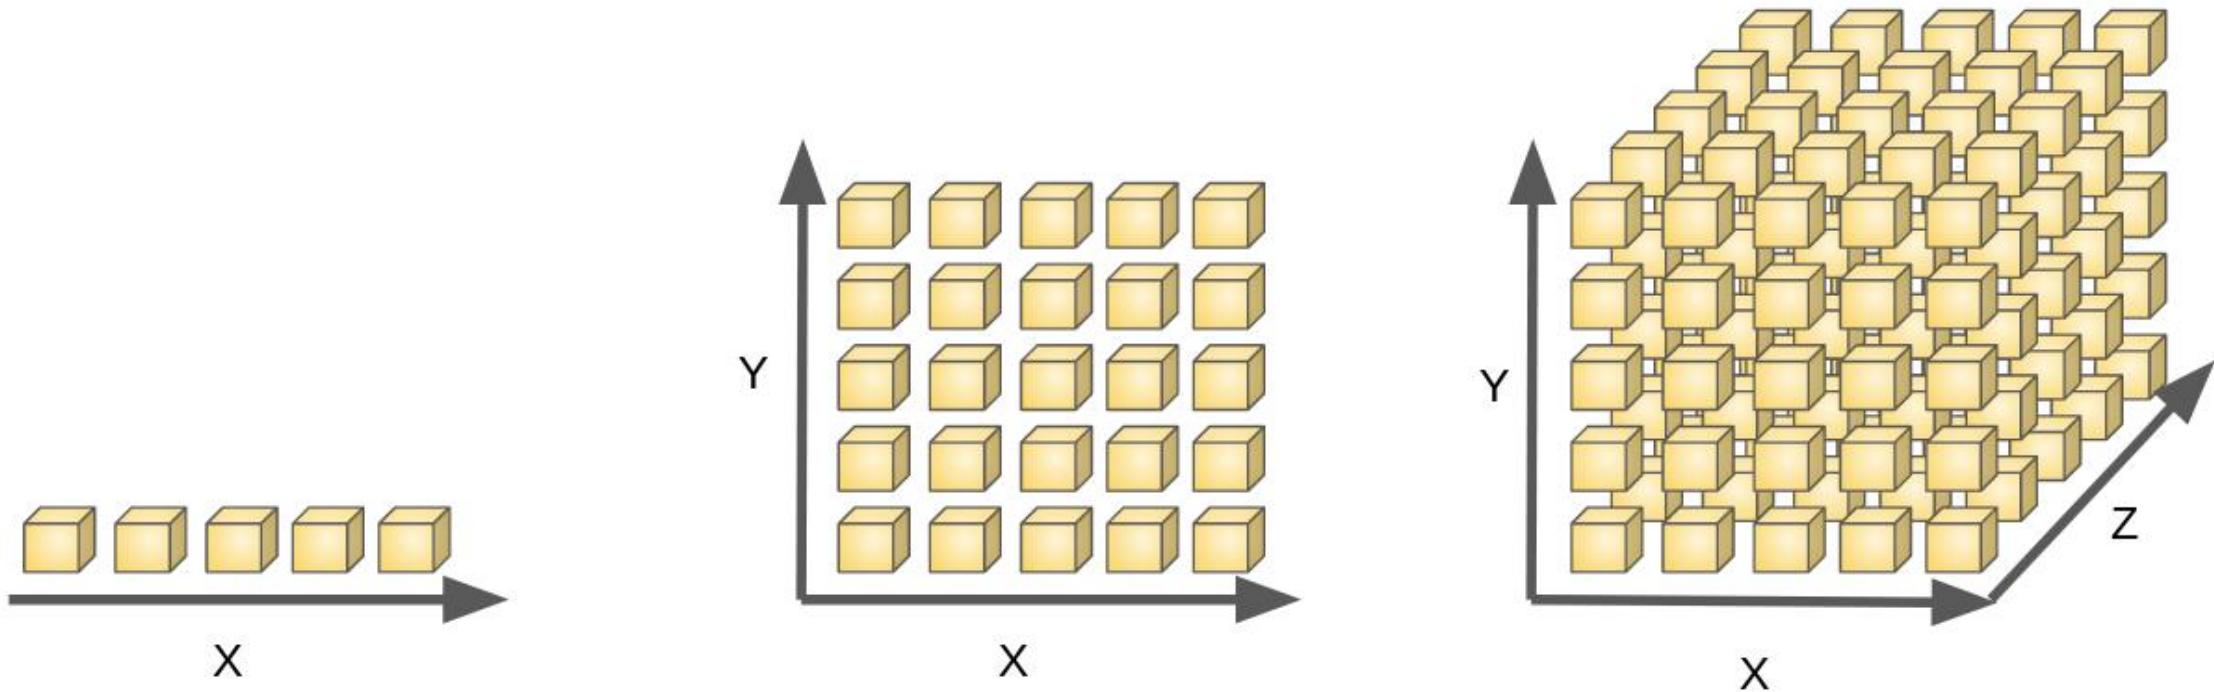
- Looked at linear algebra for data processing in the form of **Data Augmentation**
 - scalars, vectors and matrices
 - solving systems of equations
 - change of basis
 - analytical geometry
 - rotations and other transformations
- Today we will examine data processing in the form of **Dimensionality Reduction**

High-Dimensional Data



- Real world data is often high-dimensional!
- Challenge: difficult to **analyze**, **visualize** and **interpret**

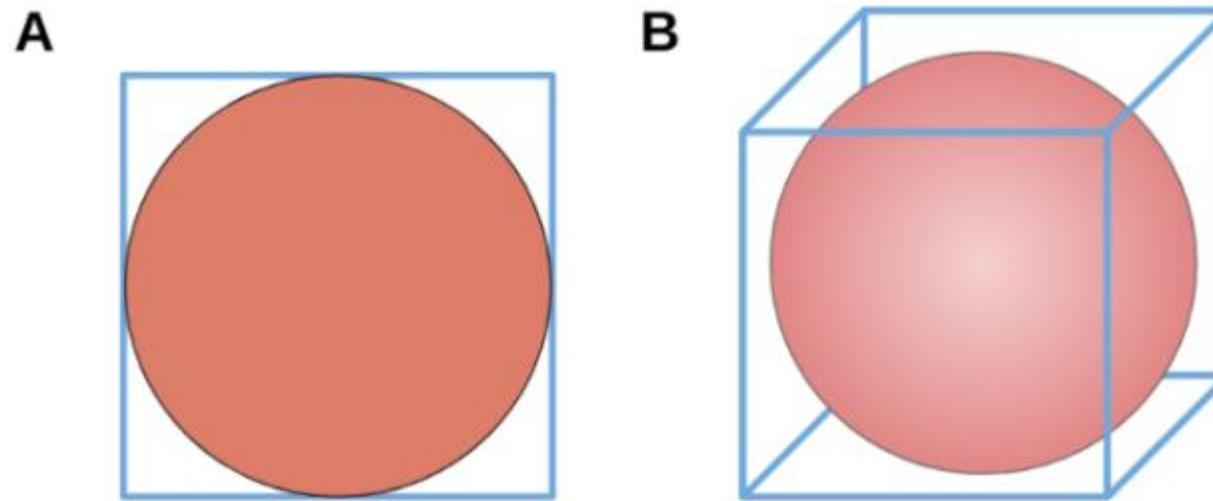
Curse of Dimensionality



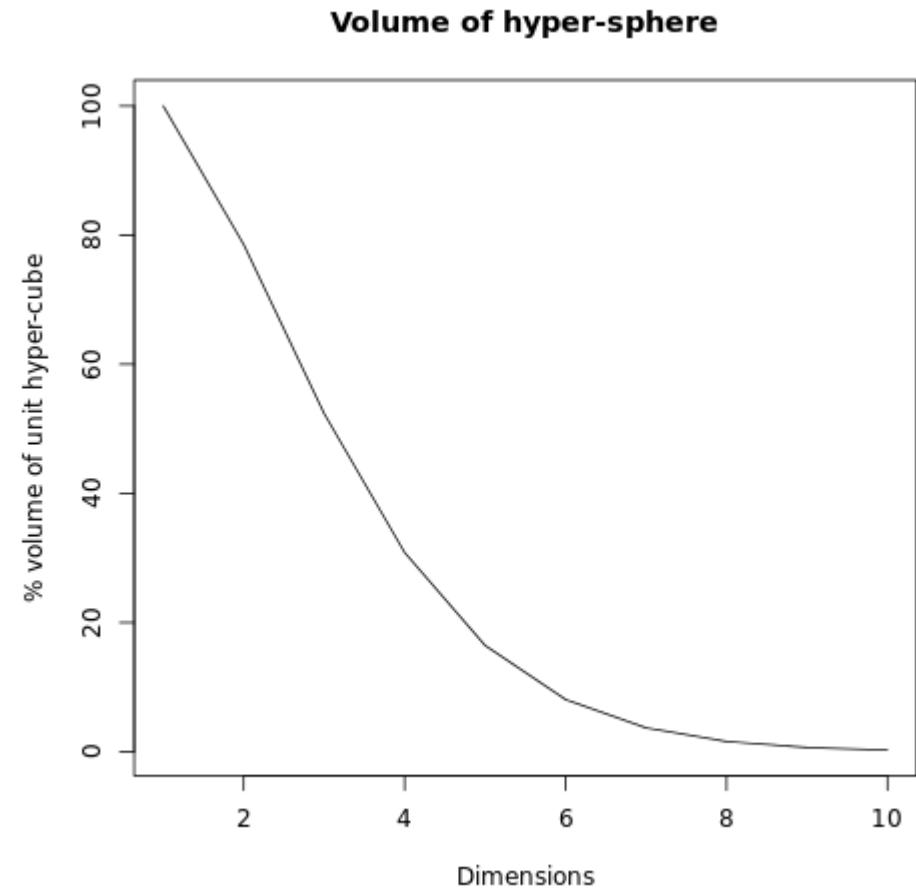
Source: [Peter Gleeson](#)

- As the number of feature or dimensions grows, the amount of data we need to generalize accurately grows exponentially.

Curse of Dimensionality



- Majority of the volume is outside of the sphere in higher dimensions

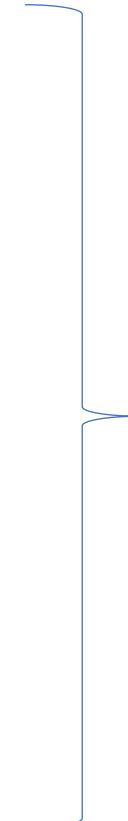


Curse of Dimensionality

- Many dimensions are unnecessary
- Data often lives on a low-dimensional manifold
- Dimensionality reduction finds the relevant dimensions

Agenda

- Projections
- Matrix Decompositions
 - Determinant
 - Trace
 - Eigenvalues and Eigenvectors
 - Eigendecomposition
- PCA
- Dimensionality Reduction



Theme:
Dimensionality Reduction
Part 1

Recap: Linear Algebra

- Given A and b, we want to solve for x:

$$Ax = b \quad \begin{bmatrix} 2 & -1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 1 \\ 5 \end{bmatrix}$$

- This can be given several interpretations:

- By rows:** x is the intersection of hyper-planes:

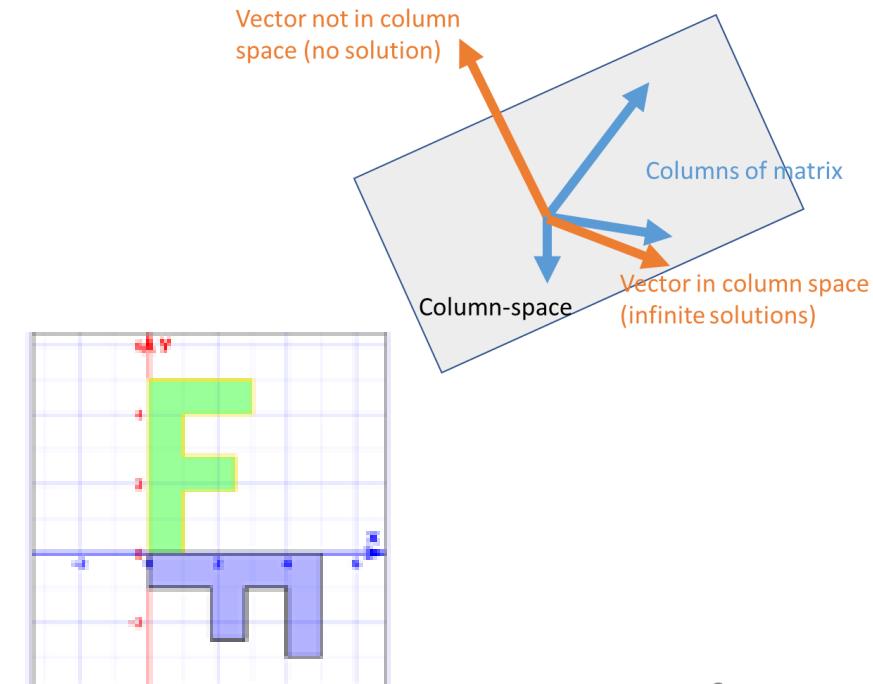
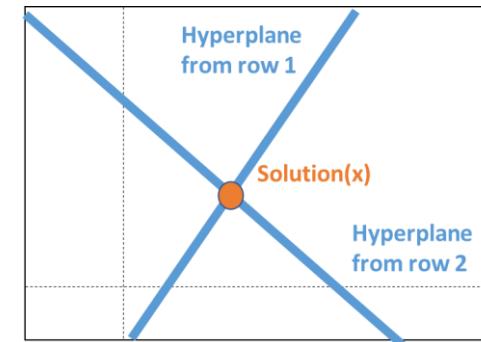
$$\begin{aligned} 2x - y &= 1 \\ x + y &= 5 \end{aligned}$$

- By columns:** x is the linear combination that gives b:

$$x \begin{bmatrix} 2 \\ 1 \end{bmatrix} + y \begin{bmatrix} -1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 5 \end{bmatrix}$$

- Transformation:** x is the vector transformed to b:

$$T(x) = b$$



Recap: Analytical Geometry

- Projections are linear transformations that project to lower dimensional feature space

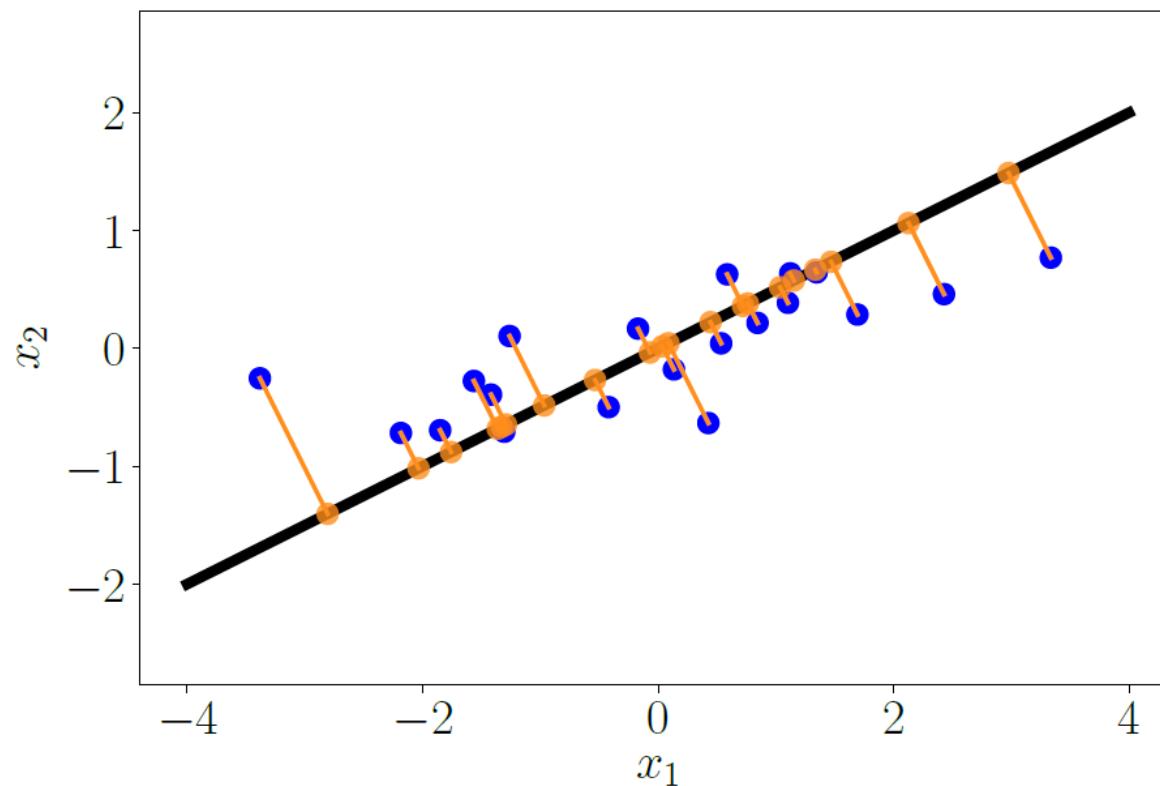
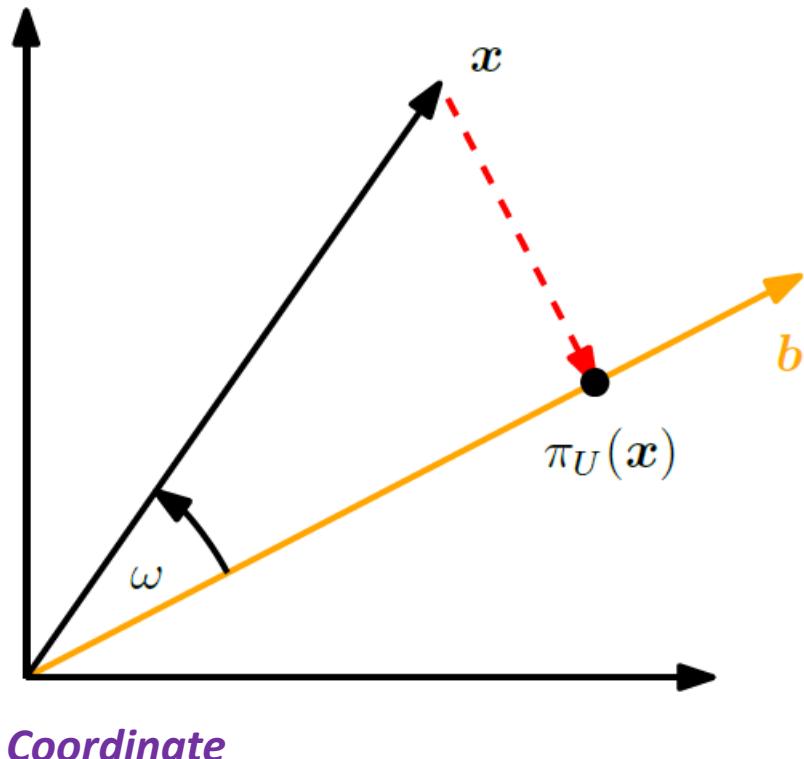


Figure 3.9
Orthogonal projection (orange dots) of a two-dimensional dataset (blue dots) onto a one-dimensional subspace (straight line).

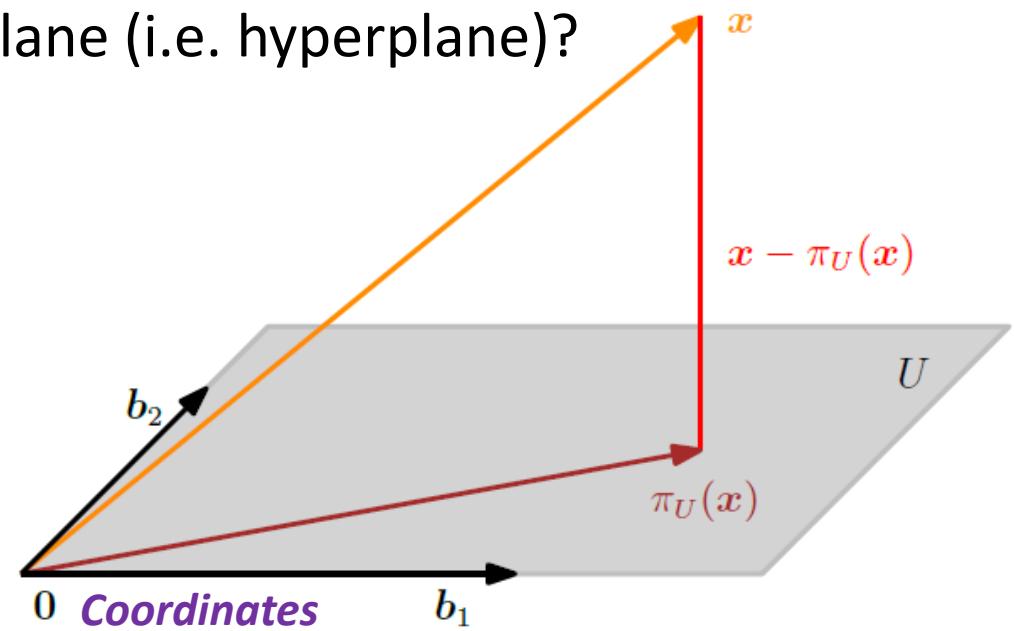
Recap: Orthogonal Projections

Project vector onto another vector



$$\pi_U(x) = \lambda b = b \frac{b^T x}{\|b\|^2} = \frac{b b^T}{\|b\|^2} x$$

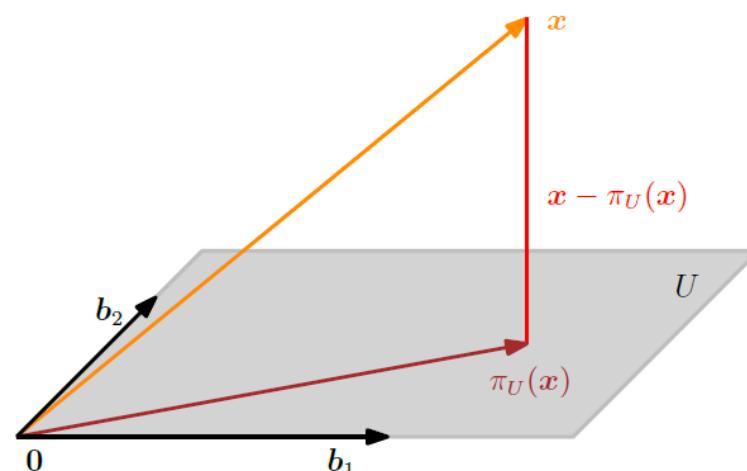
Now let's consider the case
of projecting a vector onto
a plane (i.e. hyperplane)?



$$\begin{aligned}\pi_U(x) &= B\lambda \\ B^T B \lambda &= B^T x\end{aligned}$$

Example

$$\mathbf{B}^T \mathbf{B} \boldsymbol{\lambda} = \mathbf{B}^T \mathbf{x}$$



Example 3.11 (Projection onto a Two-dimensional Subspace)

For a subspace $U = \text{span} \left[\begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix} \right] \subseteq \mathbb{R}^3$ and $\mathbf{x} = \begin{bmatrix} 6 \\ 0 \\ 0 \end{bmatrix} \in \mathbb{R}^3$ find the coordinates $\boldsymbol{\lambda}$ of \mathbf{x} in terms of the subspace U , the projection point $\pi_U(\mathbf{x})$ and the projection matrix P_π .

First, we see that the generating set of U is a basis (linear independence) and write the basis vectors of U into a matrix $B = \begin{bmatrix} 1 & 0 \\ 1 & 1 \\ 1 & 2 \end{bmatrix}$.

Second, we compute the matrix $B^T B$ and the vector $B^T \mathbf{x}$ as

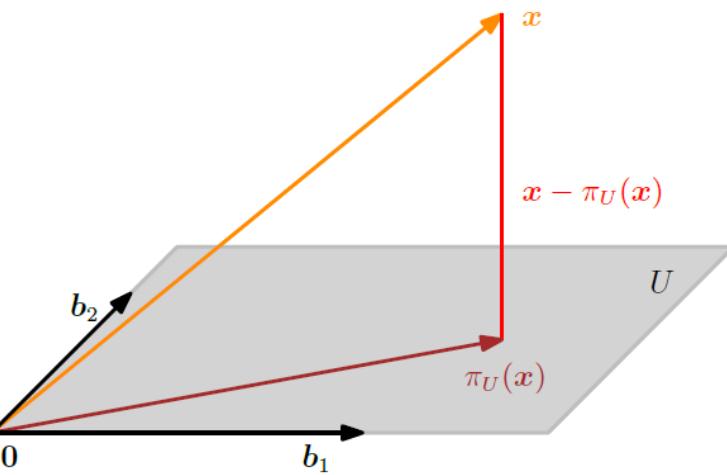
$$B^T B = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 1 & 1 \\ 1 & 2 \end{bmatrix} = \begin{bmatrix} 3 & 3 \\ 3 & 5 \end{bmatrix}, \quad B^T \mathbf{x} = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \end{bmatrix} \begin{bmatrix} 6 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 6 \\ 0 \\ 0 \end{bmatrix}. \quad (3.60)$$

Third, we solve the normal equation $B^T B \boldsymbol{\lambda} = B^T \mathbf{x}$ to find $\boldsymbol{\lambda}$:

$$\begin{bmatrix} 3 & 3 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \end{bmatrix} = \begin{bmatrix} 6 \\ 0 \end{bmatrix} \iff \boldsymbol{\lambda} = \begin{bmatrix} 5 \\ -3 \end{bmatrix}. \quad (3.61)$$

Example con't

$$B\lambda = B(B^T B)^{-1} B^T x$$



Third, we solve the normal equation $B^\top B \lambda = B^\top x$ to find λ :

$$\begin{bmatrix} 3 & 3 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \end{bmatrix} = \begin{bmatrix} 6 \\ 0 \end{bmatrix} \iff \lambda = \begin{bmatrix} 5 \\ -3 \end{bmatrix}. \quad (3.61)$$

Fourth, the projection $\pi_U(x)$ of x onto U , i.e., into the column space of B , can be directly computed via

$$\pi_U(x) = B\lambda = \begin{bmatrix} 5 \\ 2 \\ -1 \end{bmatrix}. \quad (3.62)$$

The corresponding *projection error* is the norm of the difference vector between the original vector and its projection onto U , i.e.,

$$\|x - \pi_U(x)\| = \left\| [1 \ -2 \ 1]^\top \right\| = \sqrt{6}. \quad (3.63)$$

Fifth, the projection matrix (for any $x \in \mathbb{R}^3$) is given by

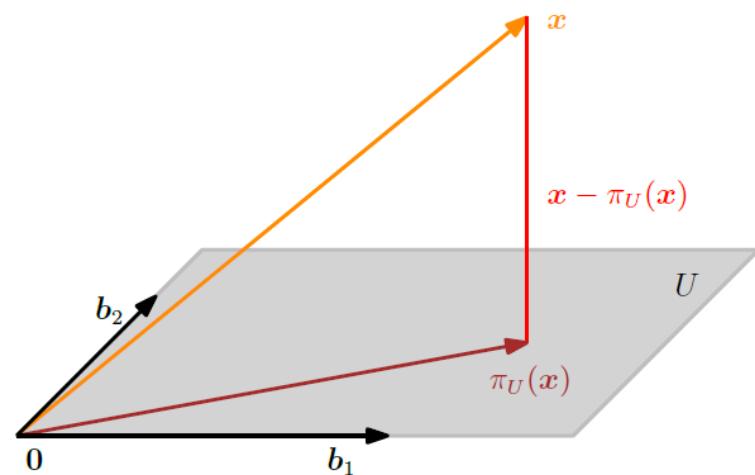
$$P_\pi = B(B^\top B)^{-1} B^\top = \frac{1}{6} \begin{bmatrix} 5 & 2 & -1 \\ 2 & 2 & 2 \\ -1 & 2 & 5 \end{bmatrix}. \quad (3.64)$$

To verify the results, we can (a) check whether the displacement vector $\pi_U(x) - x$ is orthogonal to all basis vectors of U , and (b) verify that $P_\pi = P_\pi^2$ (see Definition 3.10).

Projection onto a subspace with orthonormal bases

$$B\lambda = B(B^T B)^{-1} B^T x$$

If we have orthonormal bases (B), then this simplifies the calculation of the inverse because $B^T = B^{-1}$



$$B\lambda = BB^T x$$

Google Colab Code

“how to summarize matrices, how matrices can be decomposed, and how these decompositions can be used for matrix approximations”

Matrix Decompositions

Readings:

- Chapter 4.1-5 MML Textbook

Decomposing a Matrix - 1

- There are many ways we can describe a matrix
- Example, take a look at the following transformation:

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 5 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- Notice that we could write the same thing in terms of several transformations:

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 5 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ 5 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

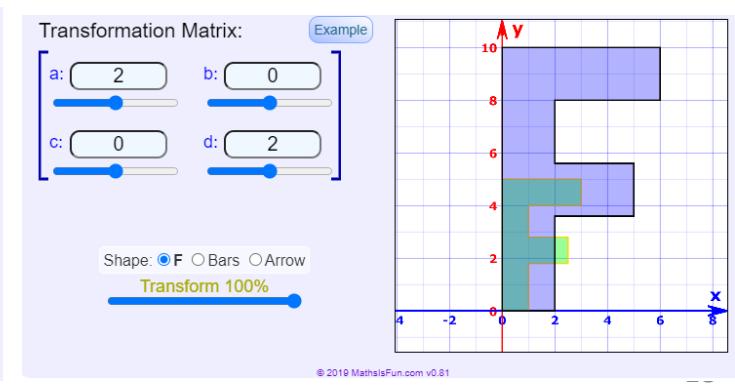
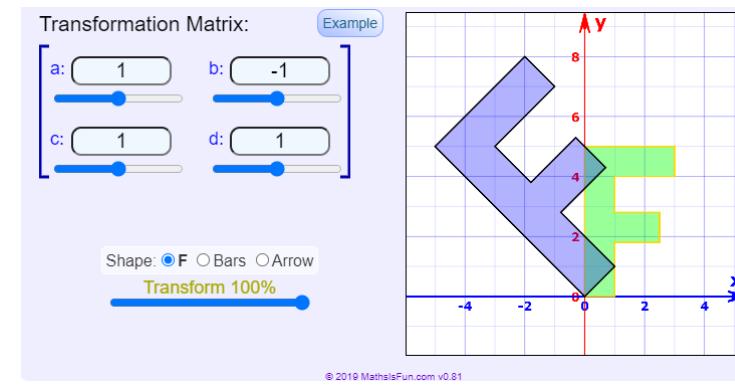
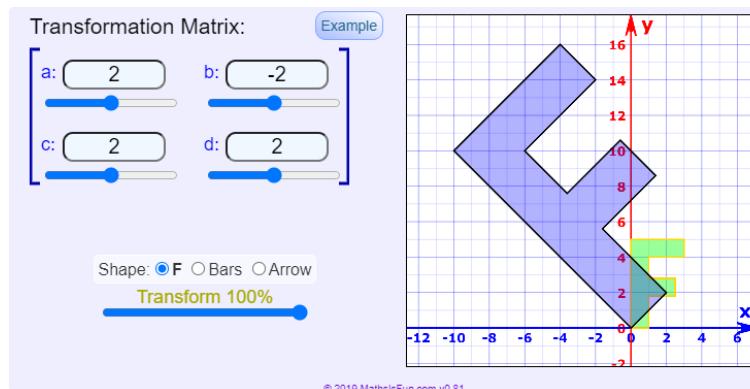
Decomposing a Matrix - 2

➤ There are many ways we can describe a matrix

➤ Example, take a look at the following transformation $A = \begin{bmatrix} 2 & -2 \\ 2 & 2 \end{bmatrix}$

➤ Notice that we could write the same thing in terms of several transformations:

$$A = \begin{bmatrix} 2 & -2 \\ 2 & 2 \end{bmatrix} = \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix}$$



Decomposing a Matrix

- This can be used for providing clarity on what the transformation is doing
- Simplifying and speeding up computations
- Dimensionality reduction as we'll see later

Key Concepts

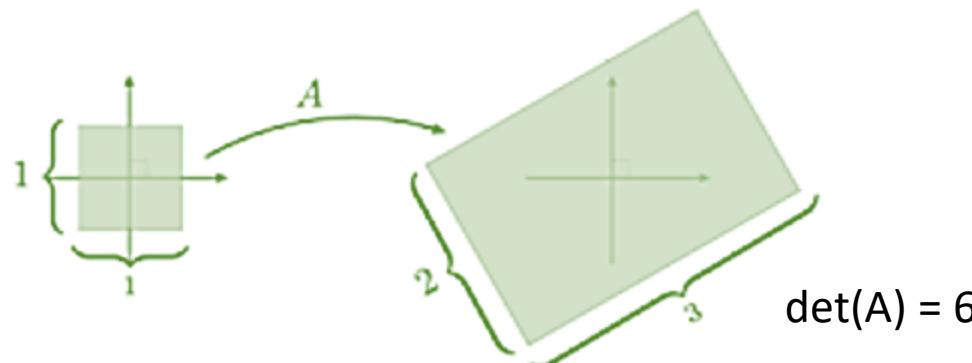
- Determinant
- Trace
- Diagonalization
- Eigenvectors
- Eigenvalues

Determinant

➤ The determinant of a square matrix is a function that maps A onto a real number.

➤ Simple geometric meaning.

- calculates how volume grows/shrinks under a linear transformation/mapping
- for example:



$$\det(A) = \begin{vmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{vmatrix}$$

Two vertical bars indicate determinant of matrix enclosed inside

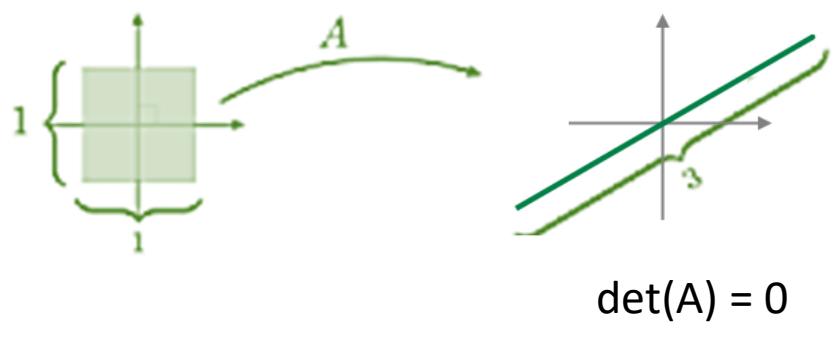
Determinant for Matrix Invertibility

➤ What does it mean when a matrix is not invertible?

➤ Simple geometric meaning.

- transformation to a subspace
- how does the volume grows/shrinks?
- for example:

Note:
If $\det(A) = 0$, then matrix A is singular.



Determining the Determinant

- The determinant can only be calculated on square matrices ($n \times n$):
- Determinant in 2d:

$$\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} = a_{11}a_{22} - a_{12}a_{21}$$

- Determinant in 3d:

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} = a_{11} \underbrace{\begin{bmatrix} a_{22} & a_{23} \\ a_{32} & a_{33} \end{bmatrix}}_{\text{recursively take alternating sums of sub-determinants}} - a_{12} \begin{bmatrix} a_{21} & a_{23} \\ a_{31} & a_{33} \end{bmatrix} + a_{13} \begin{bmatrix} a_{21} & a_{22} \\ a_{31} & a_{32} \end{bmatrix}$$
$$= a_{11}a_{22}a_{33} + a_{21}a_{32}a_{13} + a_{31}a_{12}a_{23} - a_{31}a_{22}a_{13} - a_{11}a_{32}a_{23} - a_{21}a_{12}a_{33}$$

Determining the Determinant

- Can you think of an easier way?
- Hint: what is the determinant of the following upper triangular matrix?

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{bmatrix}$$
$$\det(A) = a_{11}a_{22}a_{33} + a_{21}a_{32}a_{13} + a_{31}a_{12}a_{23} - a_{31}a_{22}a_{13} - a_{11}a_{32}a_{23} - a_{21}a_{12}a_{33}$$
$$= a_{11}a_{22}a_{33}$$

- Q: How can we obtain a triangular matrix?

Determining the Determinant

- Apply Gaussian Elimination to bring matrix into row Echelon form.
- Then determinant comes directly from the diagonal:

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{bmatrix} \quad \det(A) = \prod_i^n a_{ii}$$

for triangular matrix

- There are three things to consider:
 1. Row addition has no effect on determinant
 2. Each row switching operation changes sign of determinant
 3. Scaling a row scales the determinant

Trace

- The trace of a square matrix is the sum of the diagonal elements.
- It is defined as:

$$\text{tr}(A) = \sum_{i=1}^n a_{ii}$$

- Has some interesting identities that can help simplify equations:

- $\text{tr}(cA + dB) = c\text{tr}(A) + d\text{tr}(B)$
- $\text{tr}(A^T B) = \text{tr}(AB^T) = \text{tr}(B^T A)$
- $\text{tr}(ABC) = \text{tr}(ACB) = \text{tr}(BCA)$

Assuming that the dimensions work

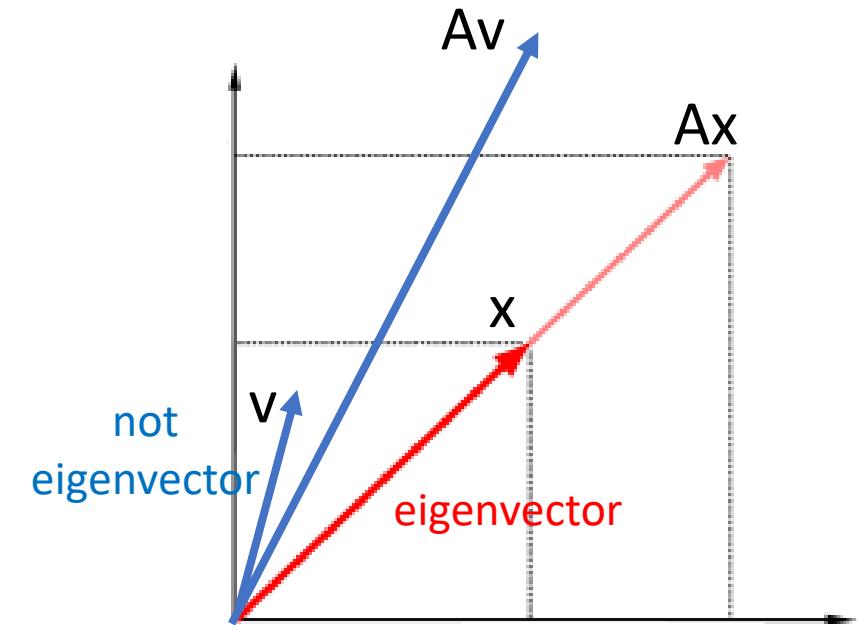
$$A = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix}$$

Eigenvectors

- Vectors that don't change their direction due to a transformation and can be represented as follows:

$$Ax = \lambda x$$

- Depending on the dimension there may be multiple independent vectors that satisfy the equation (**eigen****vectors**).
- Where λ is a scalar referred to as **eigenvalue**.

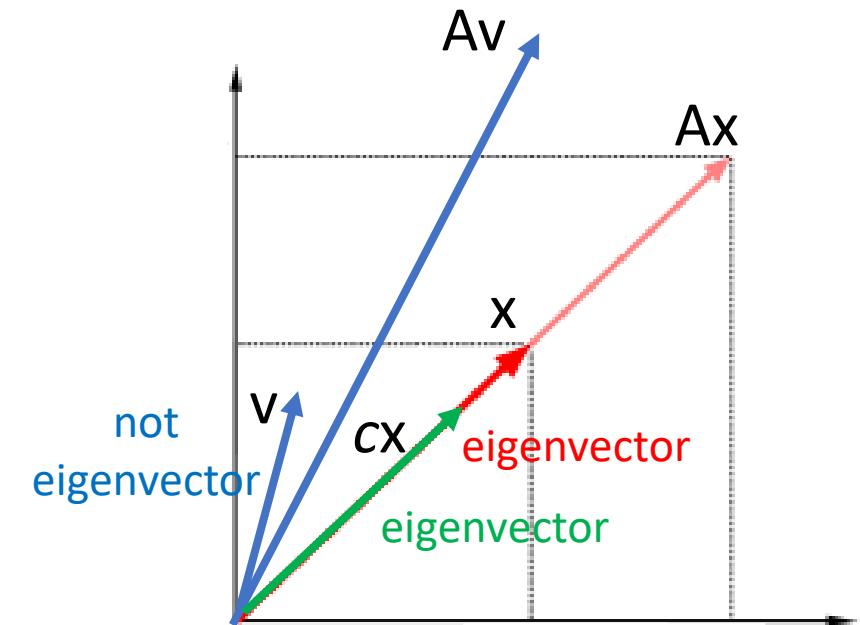


Non-Uniqueness of Eigenvectors

- If x is an eigenvector of A with corresponding eigenvalue λ then it holds that cx is also an eigenvector of A :

$$A(cx) = cAx = c\lambda x = \lambda(cx)$$

- All vectors that are collinear to x are also eigenvectors of A with the same eigenvalue λ



Finding Eigenvectors

- We can find the characteristic polynomial to find eigenvalues:

$$p_A(\lambda) = \det(A - \lambda I)$$

$$= \det \begin{bmatrix} a_{11} - \lambda & a_{12} & a_{13} \\ a_{21} & a_{22} - \lambda & a_{23} \\ a_{31} & a_{32} & a_{33} - \lambda \end{bmatrix}$$

$$= c_0\lambda + c_1\lambda^2 + \cdots + c_{n-1}\lambda^{n-1} + (-1)^n\lambda^n$$

We can describe a matrix A as a polynomial

Aside:

$$Ax = \lambda x$$

$$Ax - \lambda x = 0$$

$$(A - \lambda I)x = 0$$

solve for roots of the polynomial to obtain eigenvalues and eigenvectors!

Theorems

- There are a few theorems worth mentioning:
 1. λ_i is an eigenvalue if and only if λ_i is a root of a characteristic polynomial.
 2. eigenvectors of a square matrix ($n \times n$) with **n distinct eigenvalues** are linearly independent and form a basis of R^n .

Theorems

3. The determinant of a square matrix ($n \times n$) is the product of its eigenvalues.

$$p = \det(A) = \prod_i^n \lambda_i$$

4. The mean of the eigenvalues is the mean of the diagonal entries
i.e. the sum of the eigenvalues is the sum of the diagonal entries
i.e. the trace of an $n \times n$ matrix is the sum of its eigenvalues.

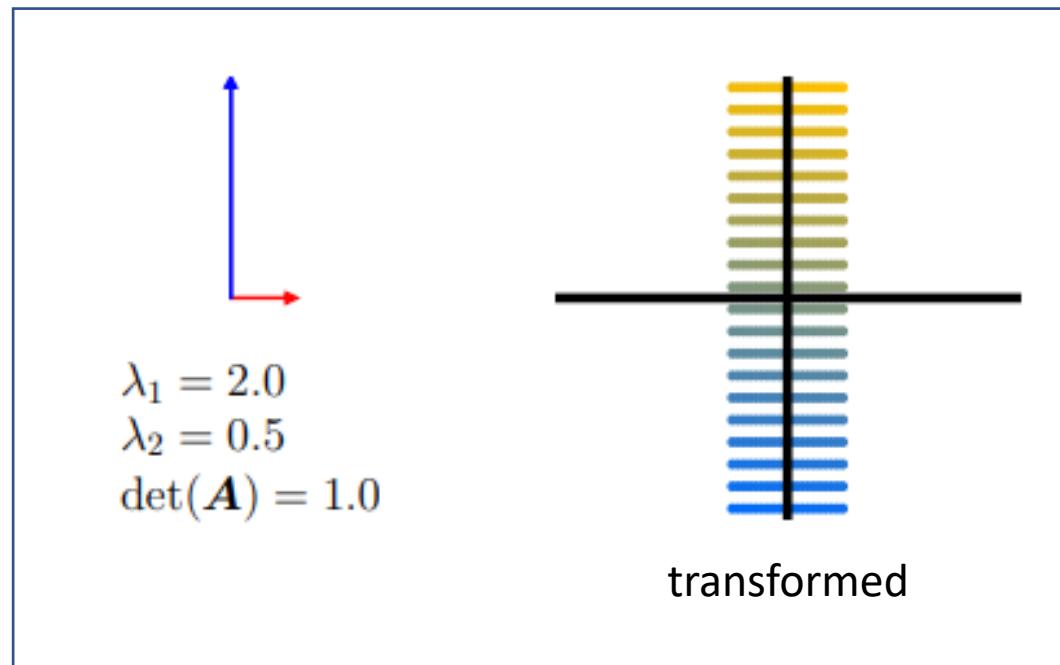
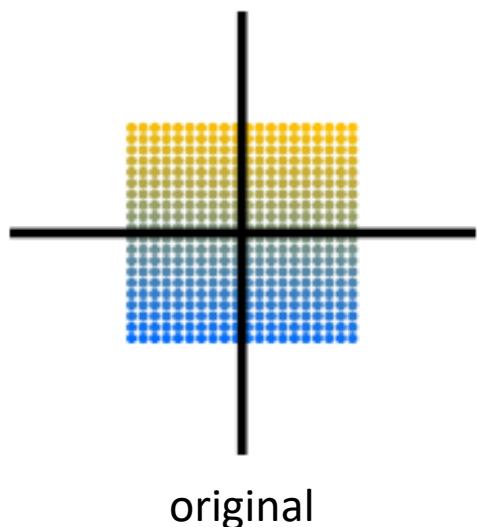
$$m = \text{tr}(A)/n = \sum_{i=1}^n \lambda_i / n$$

For 2×2 matrices, the eigenvalues are $\lambda_1, \lambda_2 = m \pm \sqrt{m^2 - p}$

Two numbers with a sum of 7 and a product of 10?

Graphical Intuition - Scaling

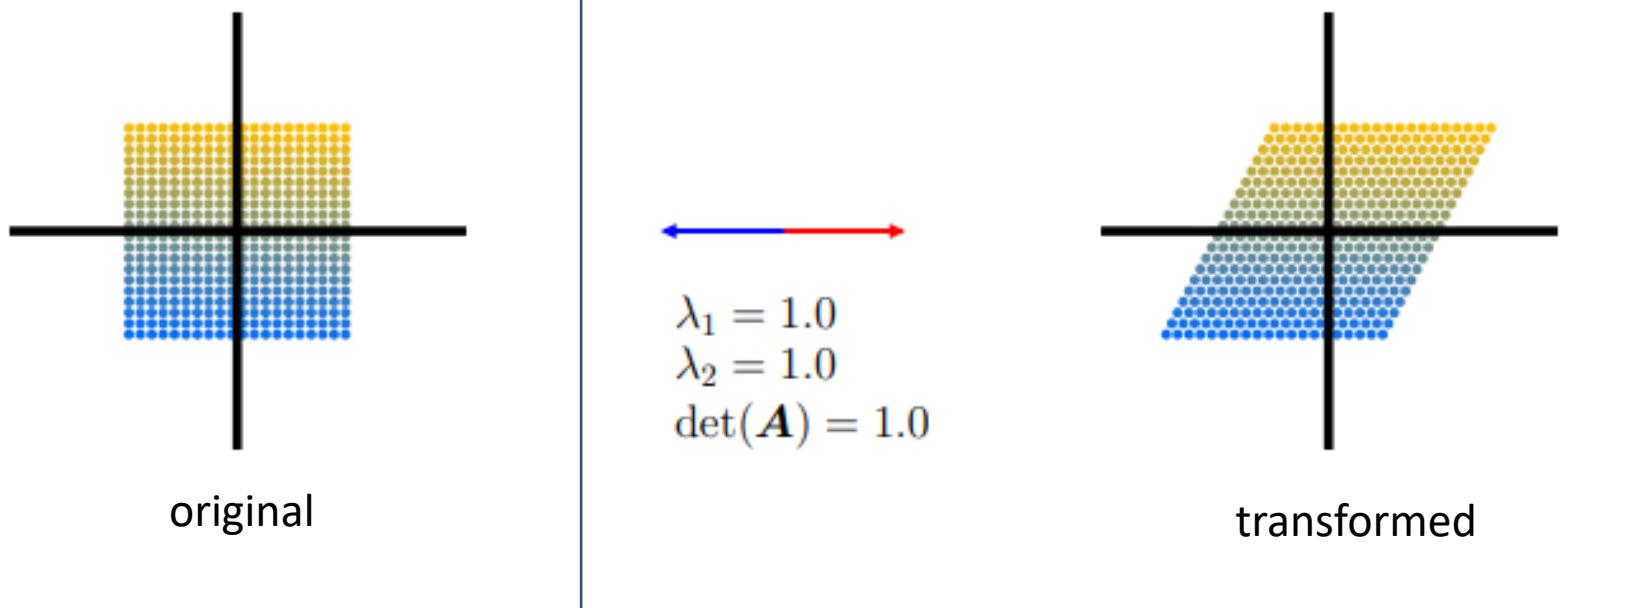
$$A = \begin{bmatrix} 1/2 & 0 \\ 0 & 2 \end{bmatrix}$$



➤ The vertical axis stretched by a factor of 2 and the horizontal axis is compressed by a factor of 1/2. The mapping is area preserving.

Graphical Intuition - Shear

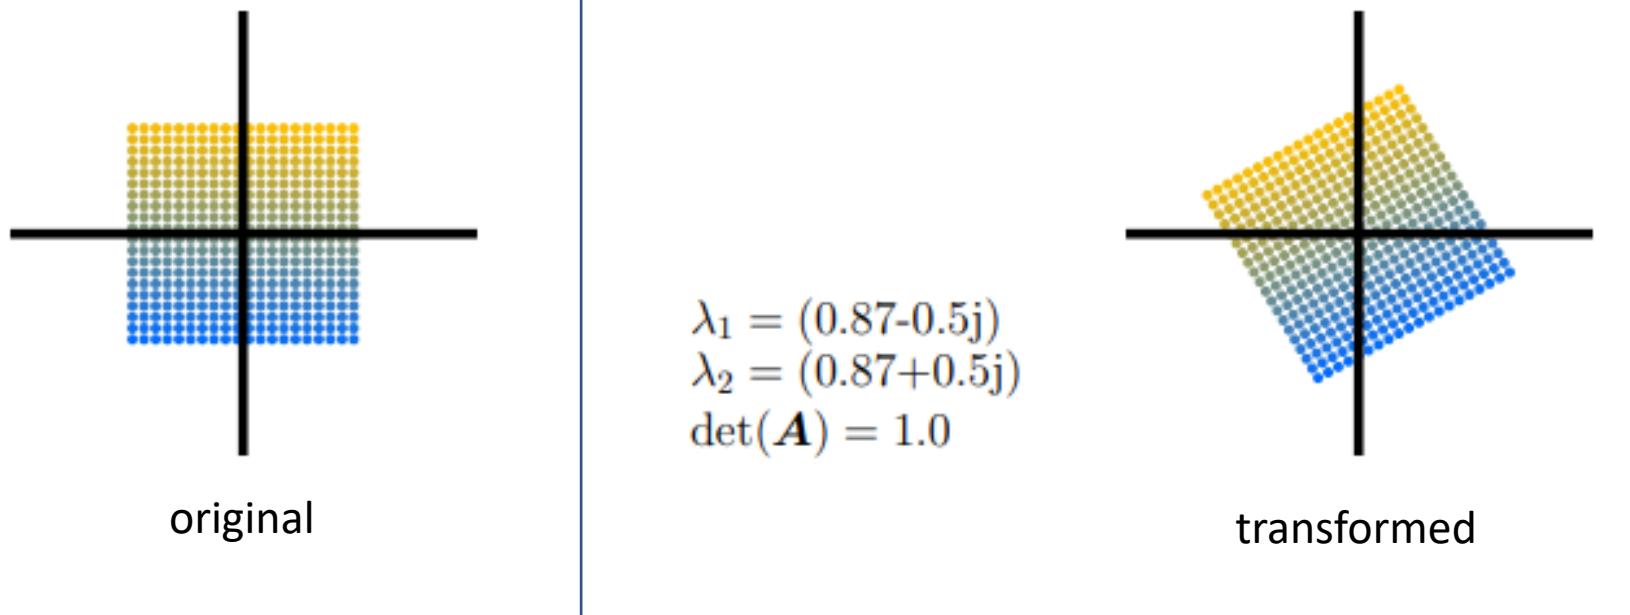
$$A = \begin{bmatrix} 1 & 1/2 \\ 0 & 1 \end{bmatrix}$$



➤ Shears the points along the horizontal axis to the right if on positive side of the vertical axis and left if on negative side. Area preserving.

Graphical Intuition - Rotation

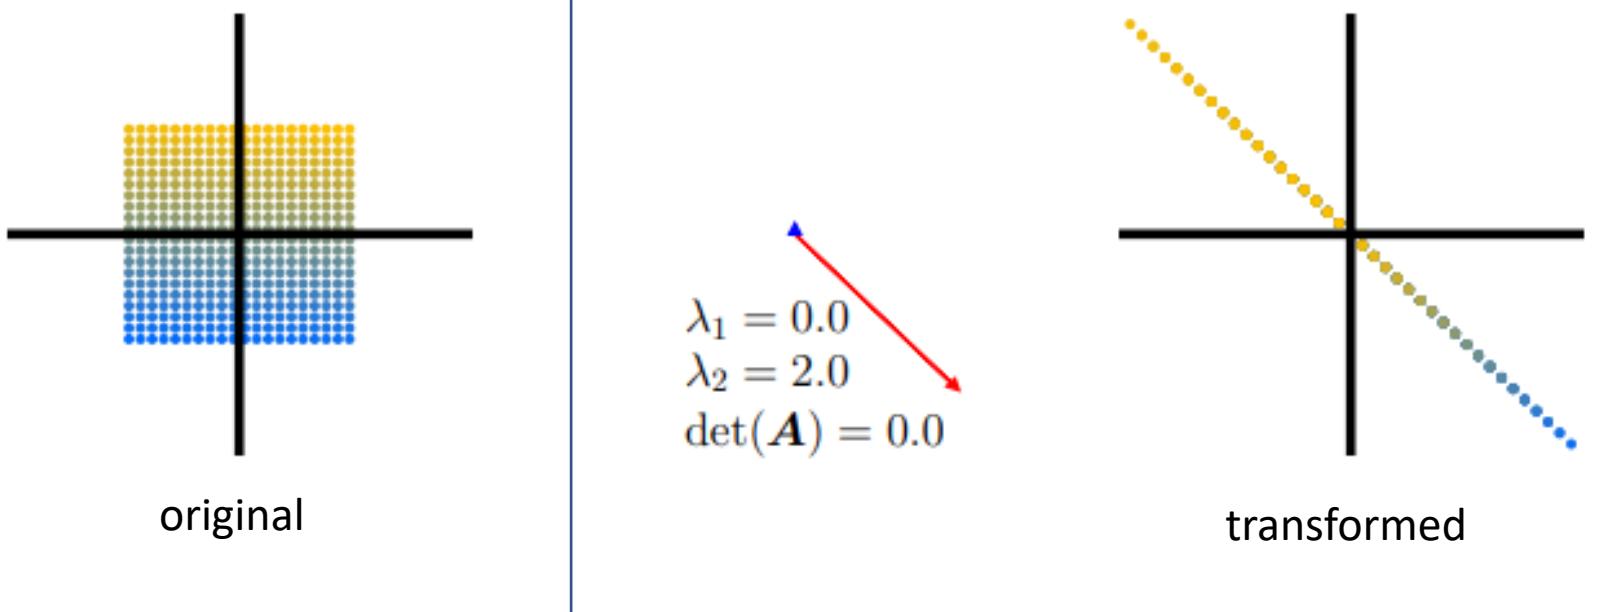
$$A = \begin{bmatrix} \cos(\pi/6) & -\sin(\pi/6) \\ \sin(\pi/6) & \cos(\pi/6) \end{bmatrix} = 0.5 \begin{bmatrix} \sqrt{3} & -1 \\ 1 & \sqrt{3} \end{bmatrix}$$



➤ The matrix rotates the points by $\pi/6$ rad (30 degrees) counterclockwise. The eigenvalues are complex, so cannot draw eigenvectors on real axis. The rotation is area preserving.

Graphical Intuition – Collapse to Subspace

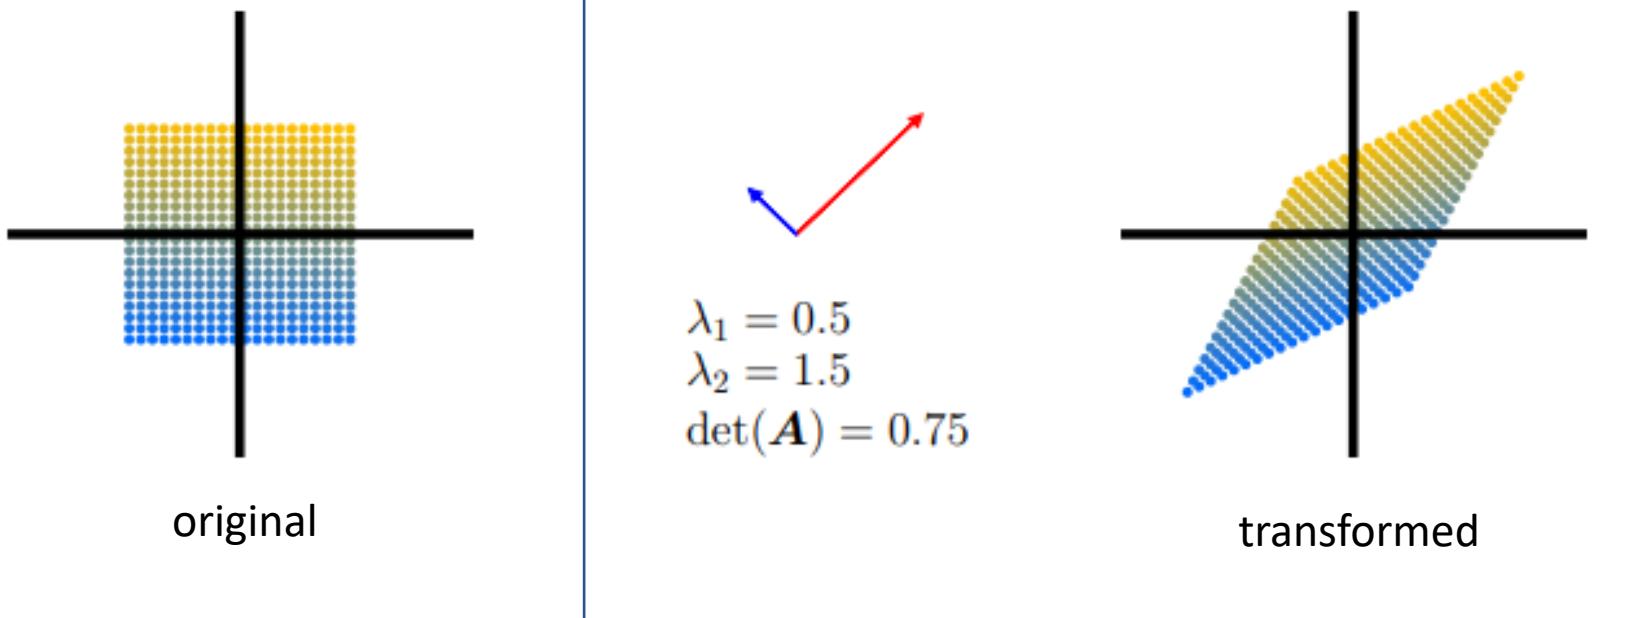
$$A = \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix}$$



- Represents a mapping that collapses a two-dimensional domain onto one dimension. The eigenvalue $\lambda_1 = 0$ collapses while the other stretches the space by 2. The area in this case is 0.

Graphical Intuition – Sheer and Stretch

$$A = \begin{bmatrix} 1 & 1/2 \\ 1/2 & 1 \end{bmatrix}$$



- Represents a shear and stretch mapping. It stretches along the red eigenvector by 1.5 and compresses along blue eigenvector by 0.5. That area is scaled by 75%.

Eigendecomposition

- A **square matrix** ($n \times n$) can be factored into:

$$A = PDP^{-1}$$

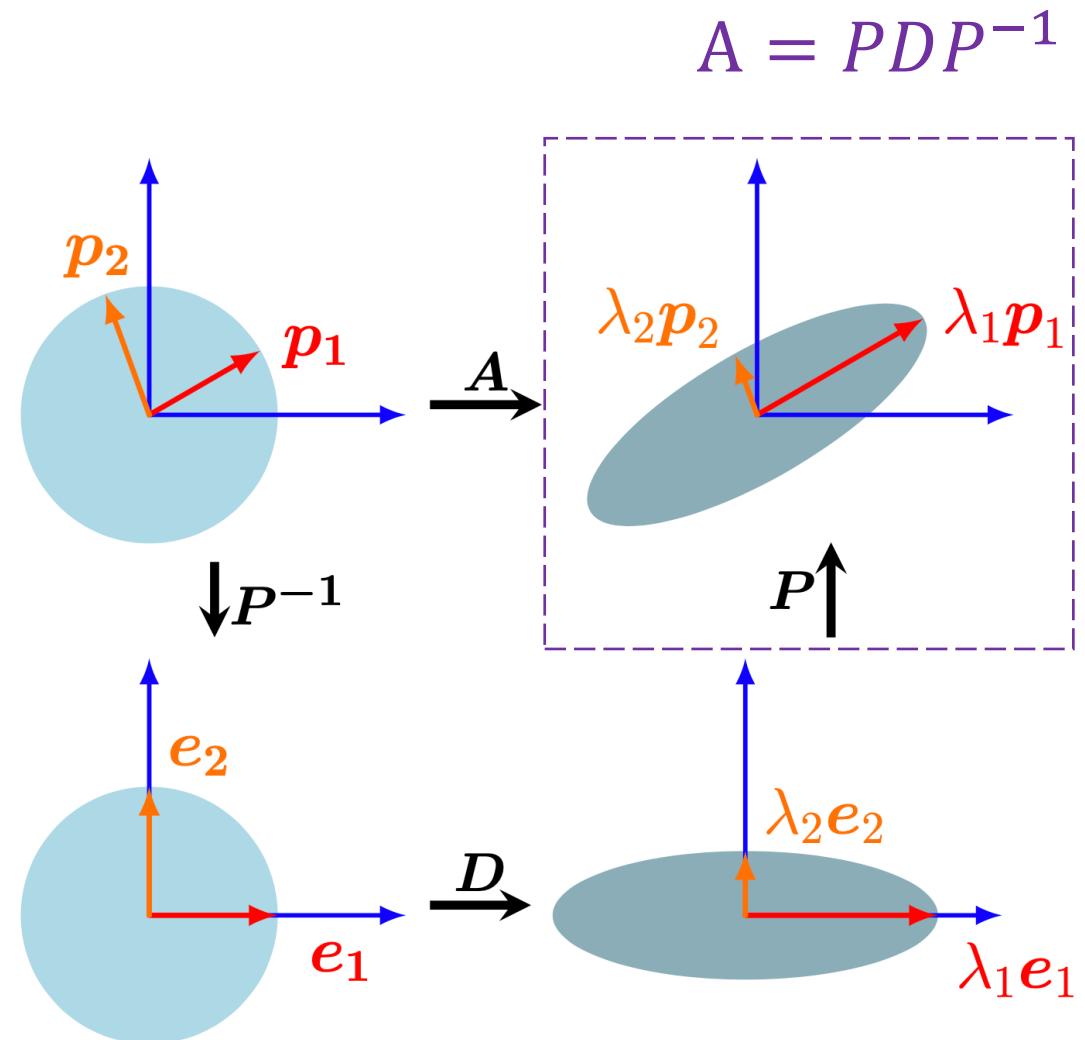
where P is $R^{n \times n}$ and D is a diagonal matrix whose diagonal entries are eigenvalues of A (if and only if the eigenvectors form the basis of R^n)

- A **square symmetric matrix** ($n \times n$) the eigenvectors form an orthonormal basis (with real eigenvalues) and A can be factored into:

$$A = PDP^T \quad \text{Spectral Theorem}$$

Geometric Intuition

- **Eigendecomposition** can be thought of as a sequence of transformations:
 1. P^{-1} performs a basis change depicted as a **rotation-like operation** from **standard basis** into the **eigenbasis**
 2. D performs a **scaling** along the **remapped orthogonal vectors** depicted by a circle being stretched
 3. P undoes the basis change depicted as a **reverse rotation** and restores the original coordinate frame.



Example 4.11 $A = \begin{bmatrix} 5/2 & -1 \\ -1 & 5/2 \end{bmatrix}$ Eigendecomposition of $A = PDP^{-1}$?

$$A = PDP^{-1} = \begin{bmatrix} 1 & 1 \\ -1 & 1 \end{bmatrix} \begin{bmatrix} 7/2 & 0 \\ 0 & 3/2 \end{bmatrix} \begin{bmatrix} 1/2 & -1/2 \\ 1/2 & 1/2 \end{bmatrix}$$

Principal Component Analysis

Readings:

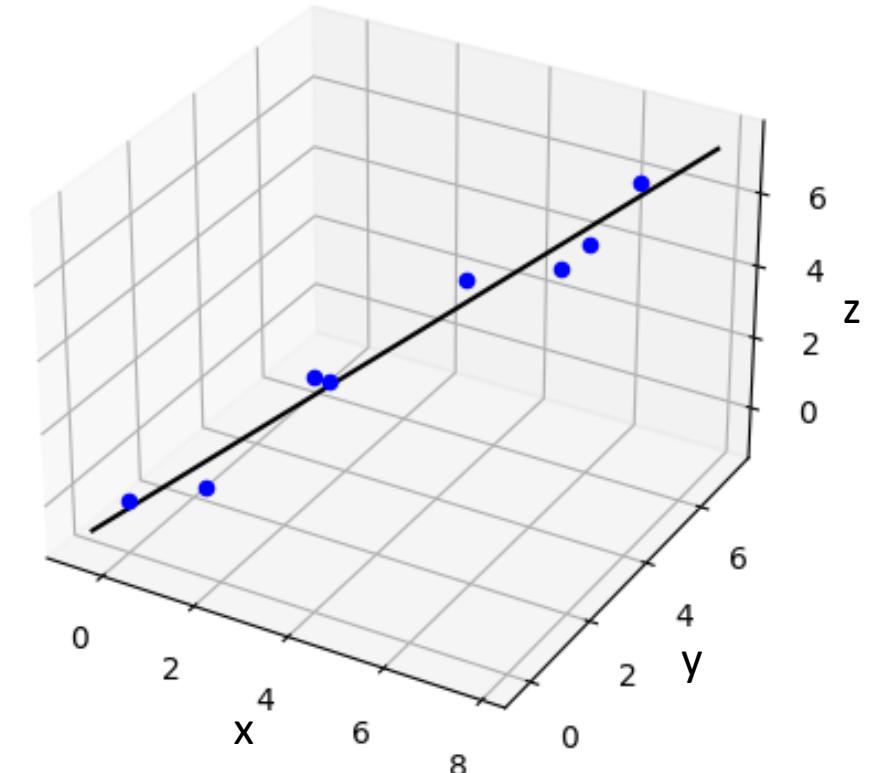
- Chapter 10 MML Textbook

PCA

- Principle component analysis (PCA) is one example where eigendecomposition is used frequently
- Decomposing data in terms of factors has many benefits:
 - data interpretation
 - dimensionality reduction

You can think of it as projections that maximize the variance in the data (i.e. information content).

Example



These slides are
adapted from

StatQuest

<https://www.youtube.com/watch?v=FgakZw6K1QQ>

Data

Samples!

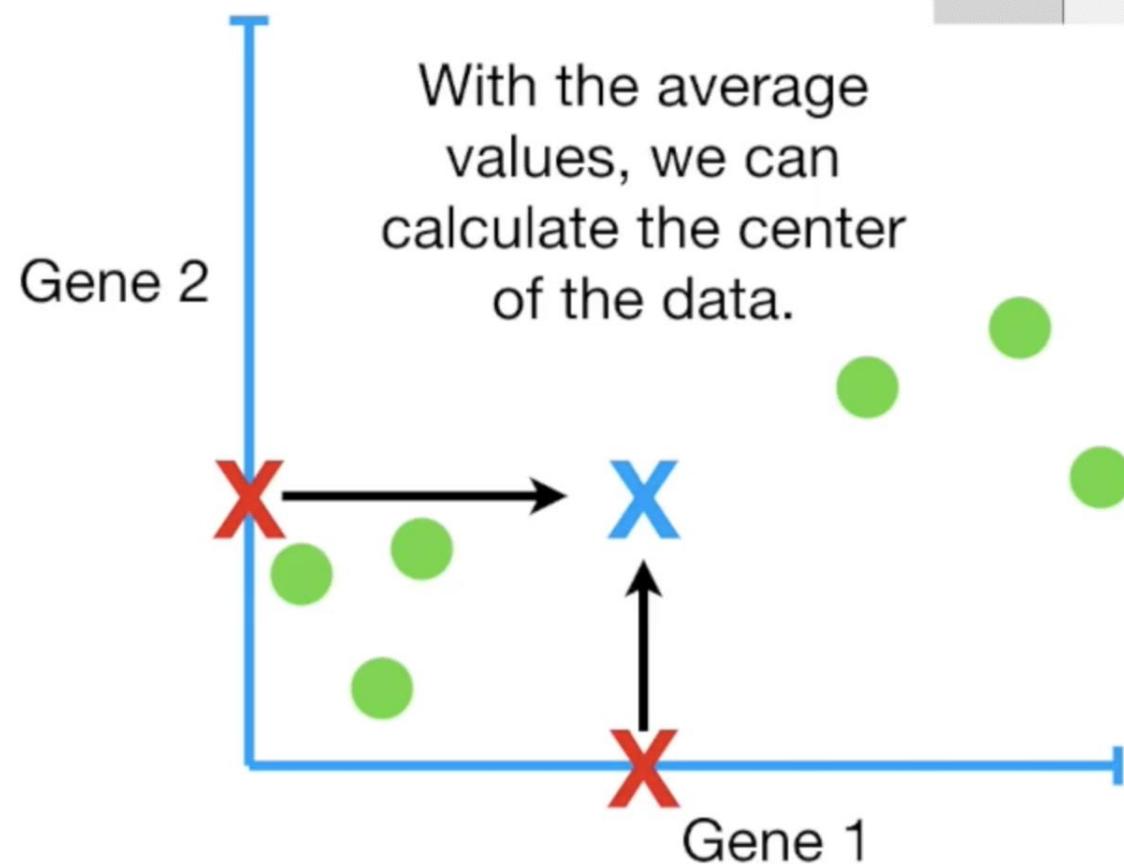
Variables!

	Mouse 1	Mouse 2	Mouse 3	Mouse 4	Mouse 5	Mouse 6
Gene 1	10	11	8	3	2	1
Gene 2	6	4	5	3	2.8	1

We can plot 2 variables in 2D.

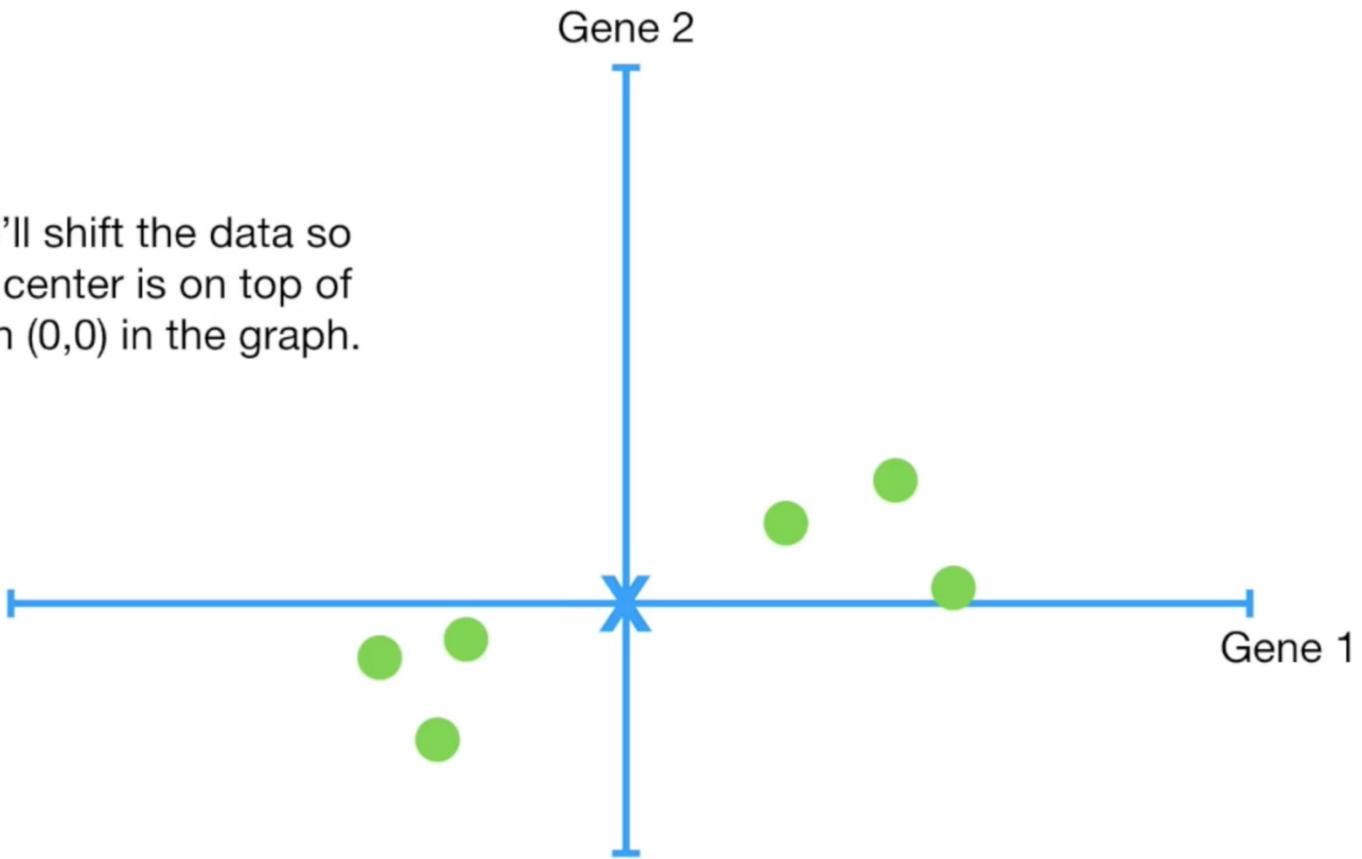
Average

	Mouse 1	Mouse 2	Mouse 3	Mouse 4	Mouse 5	Mouse 6
Gene 1	10	11	8	3	2	1
Gene 2	6	4	5	3	2.8	1

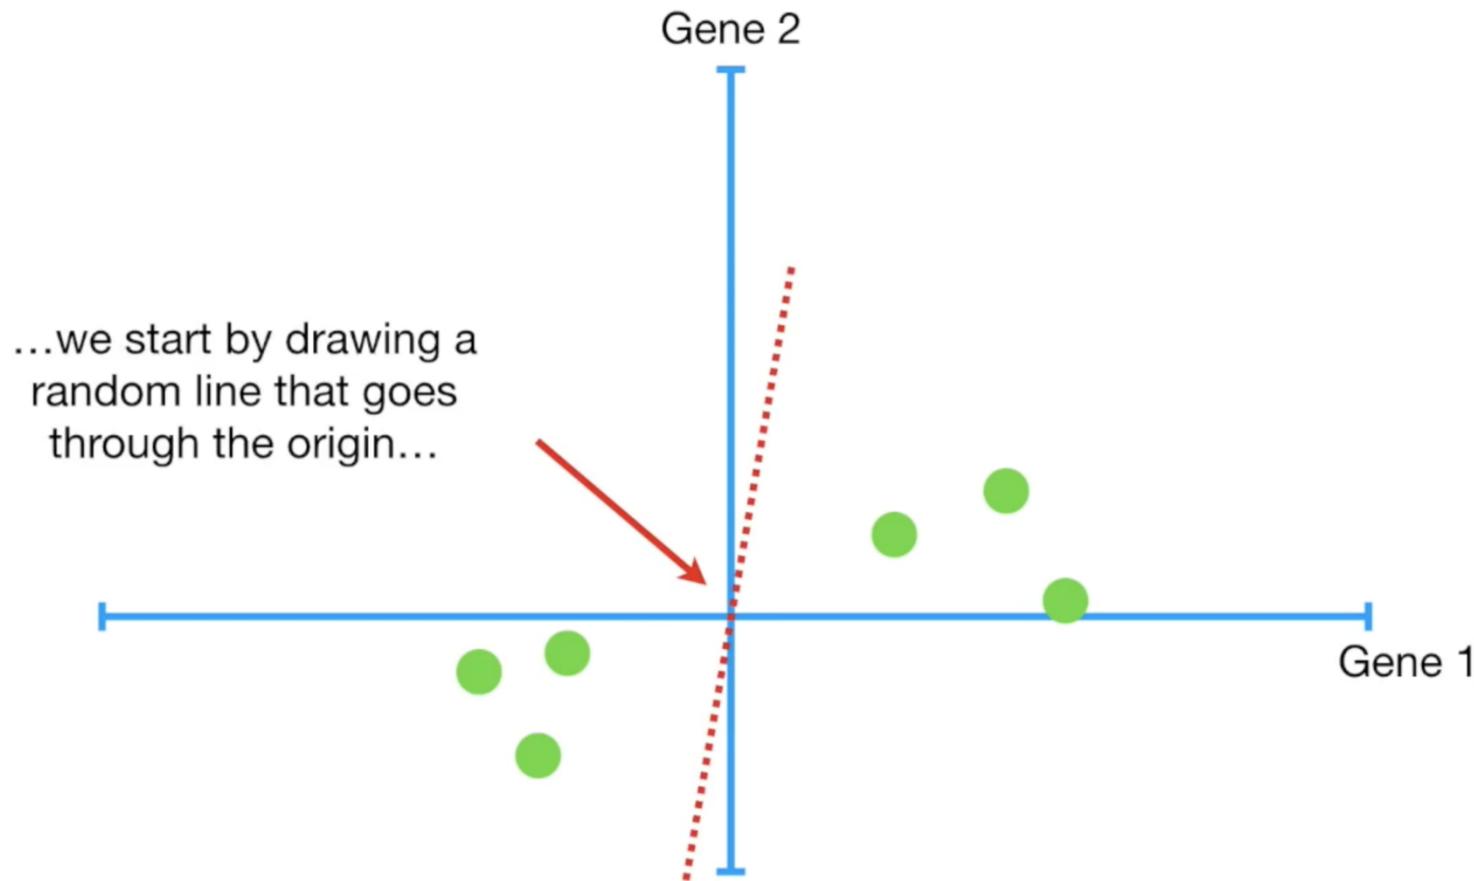


Shift the data

Now we'll shift the data so that the center is on top of the origin (0,0) in the graph.

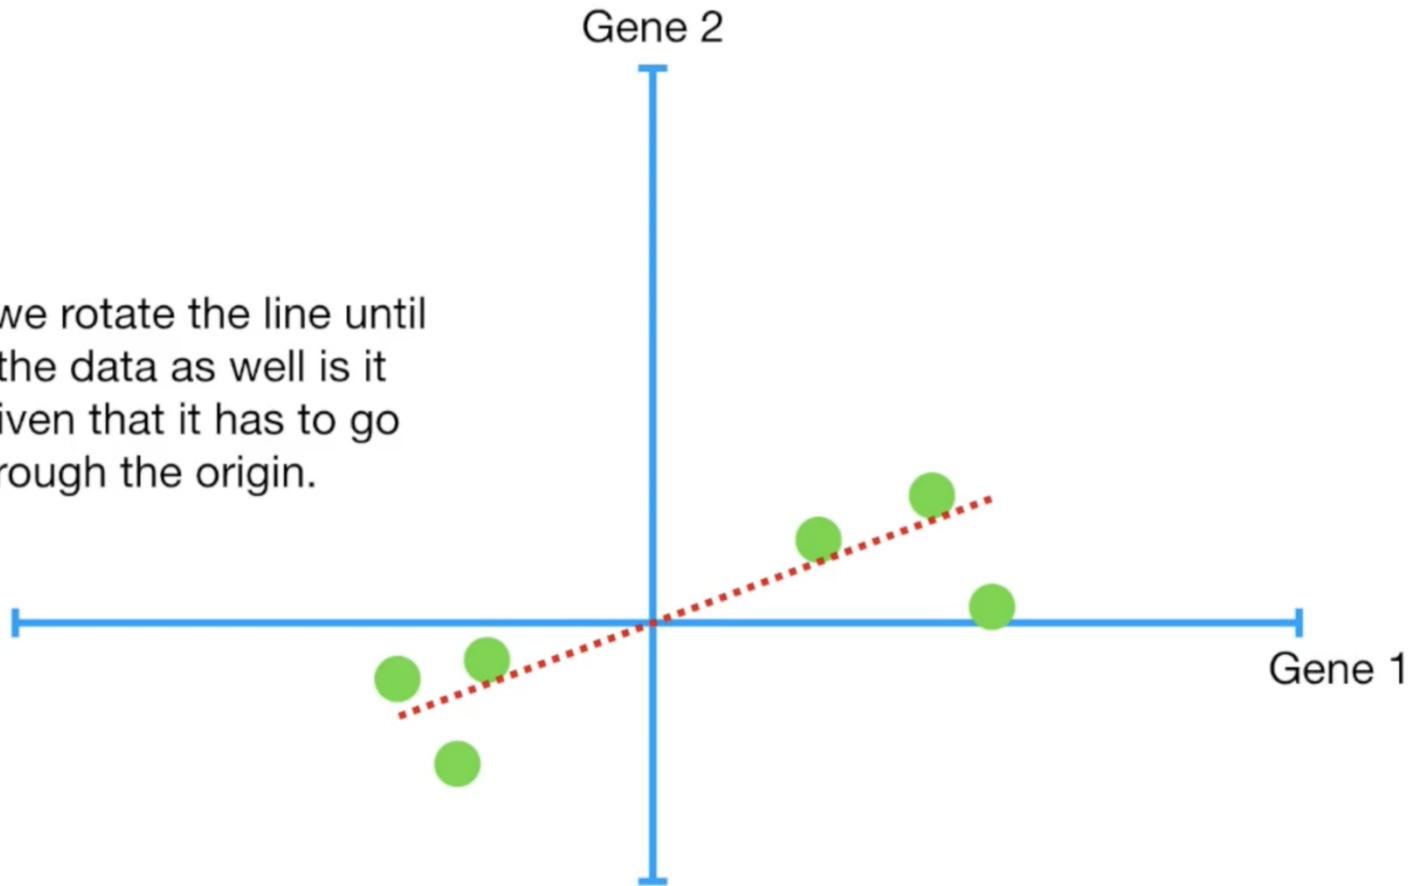


Draw a line crossing the origin

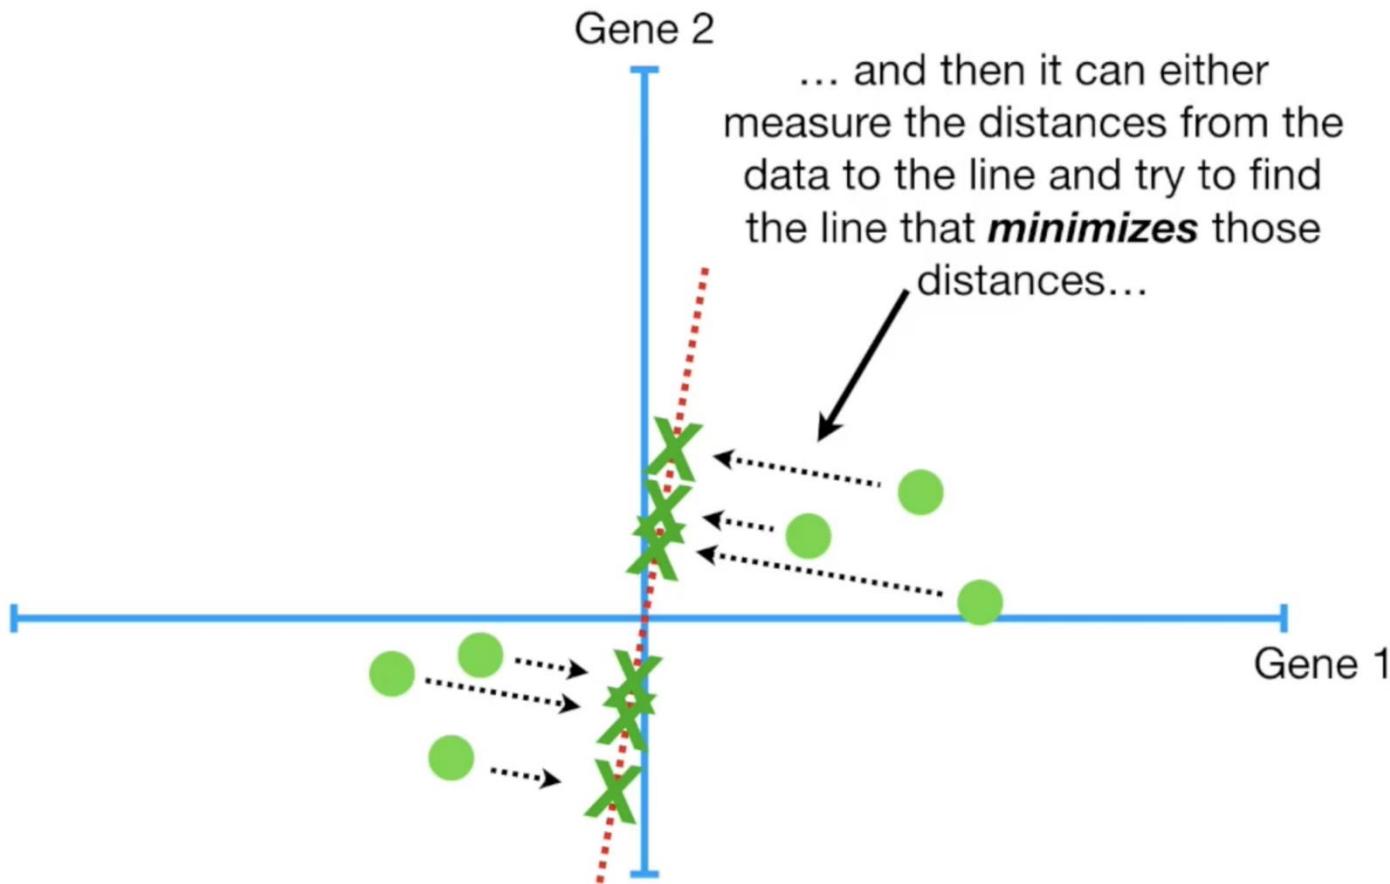


Goal: Fit the data as best as a line can

...then we rotate the line until it fits the data as well as it can, given that it has to go through the origin.

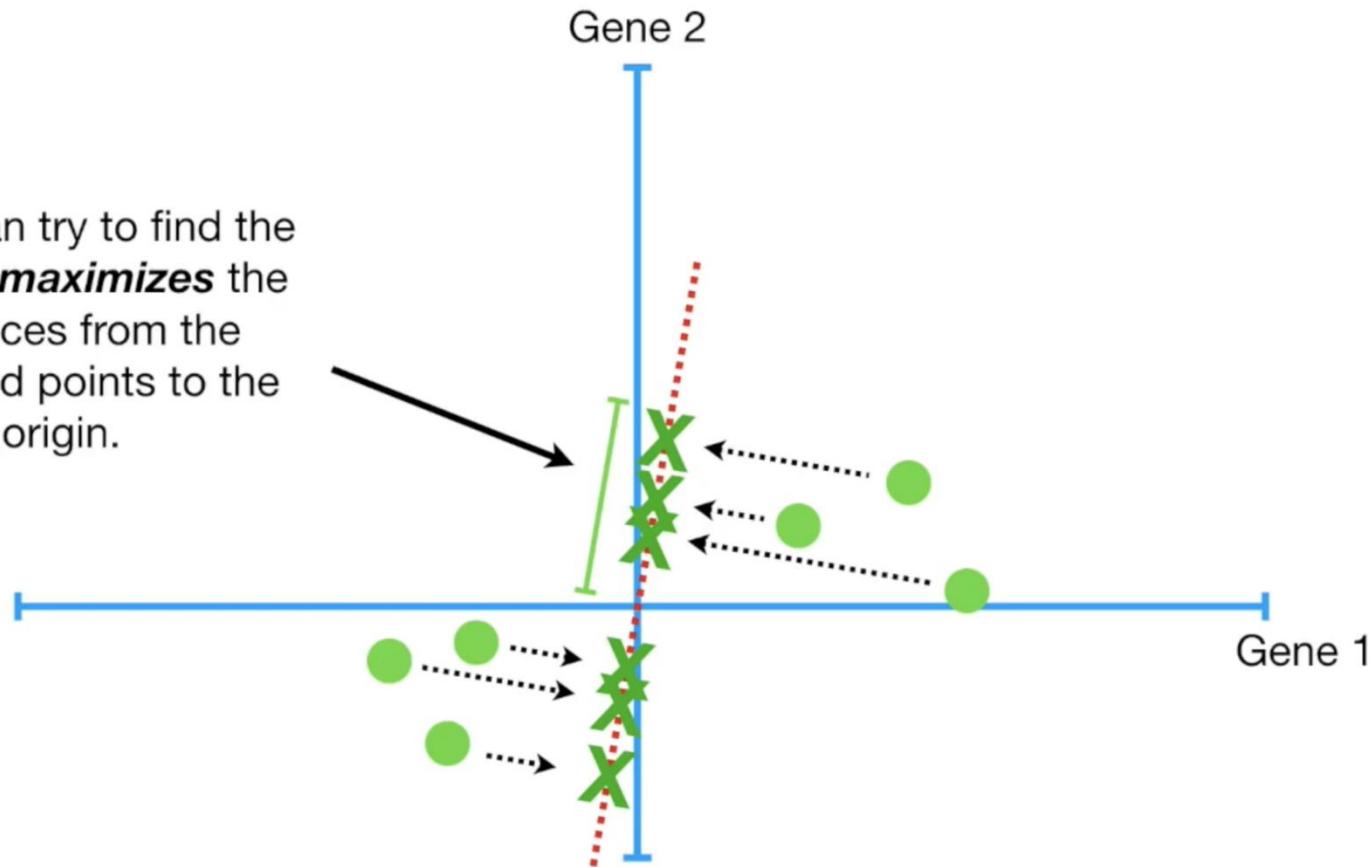


How do we know we have the best fit?

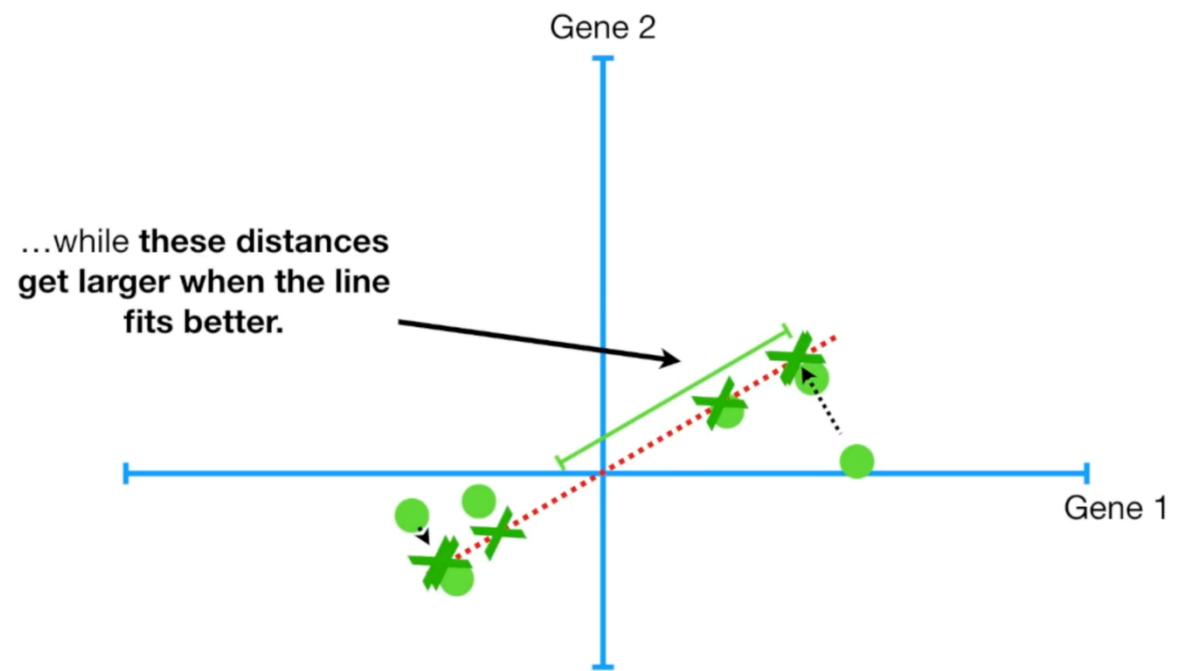
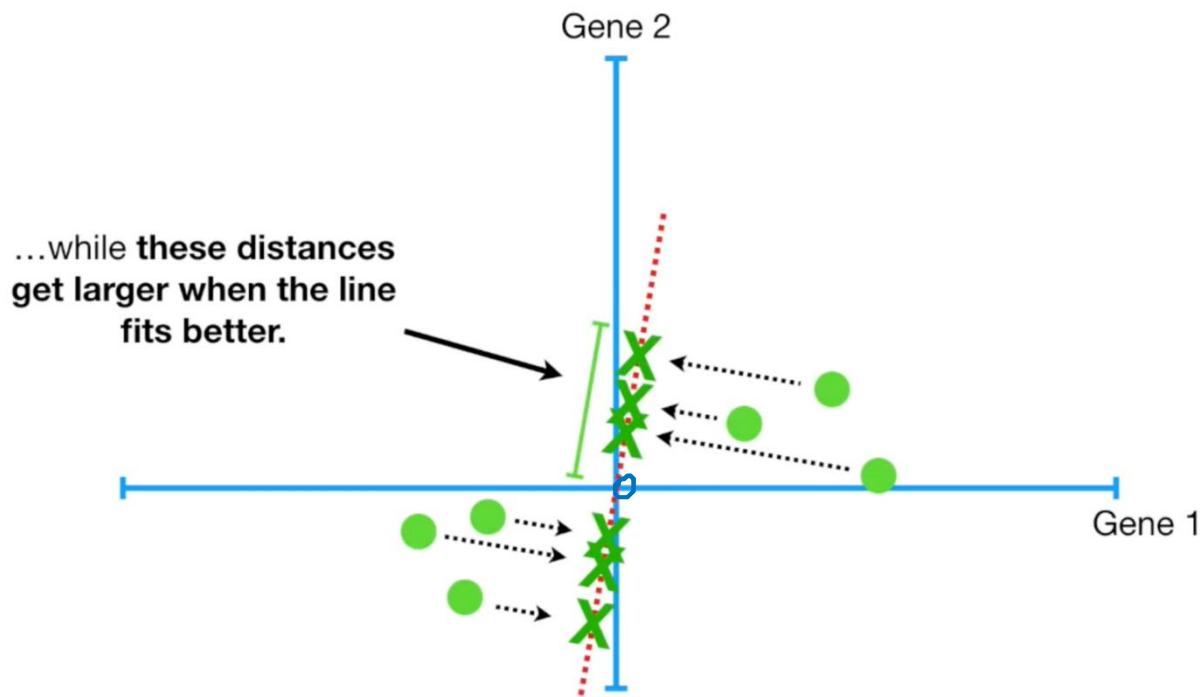


How we know the best fit?

...or it can try to find the line that **maximizes** the distances from the projected points to the origin.



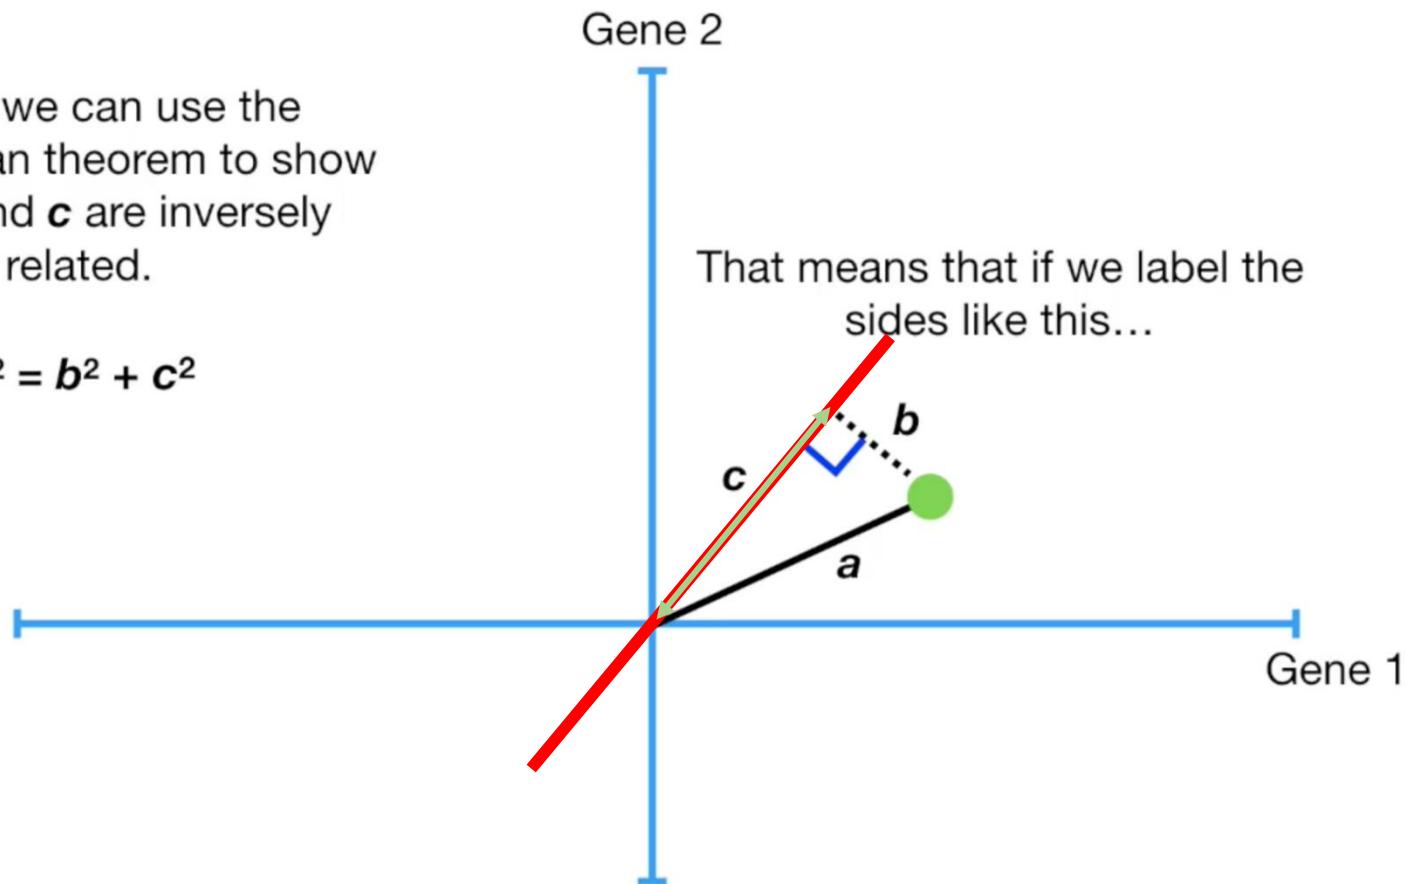
Intuition



Why?

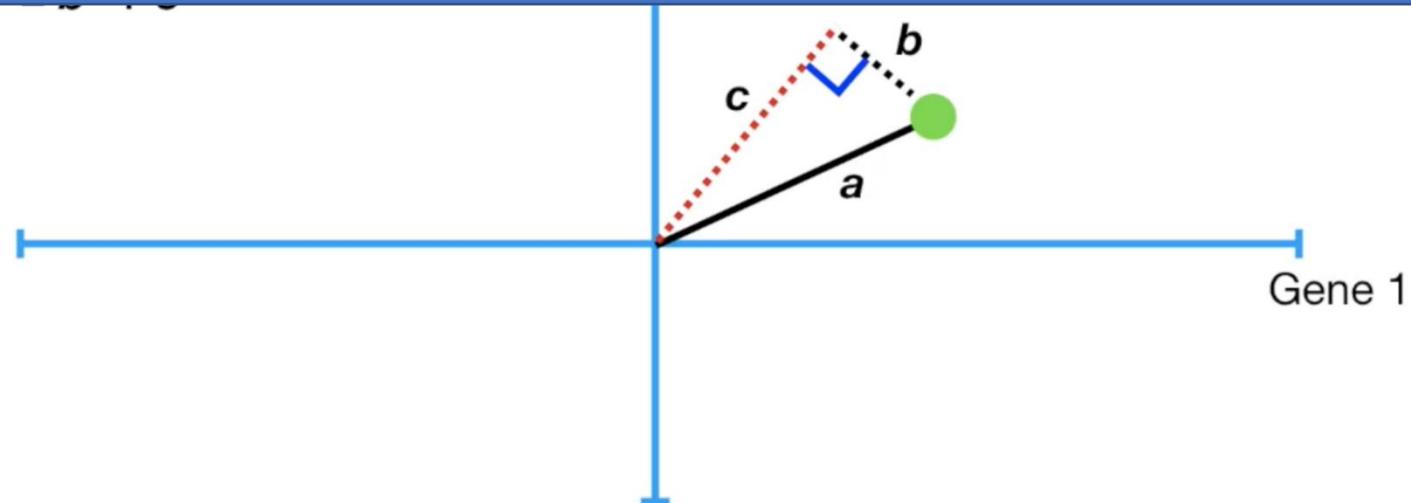
...then we can use the Pythagorean theorem to show how **b** and **c** are inversely related.

$$a^2 = b^2 + c^2$$

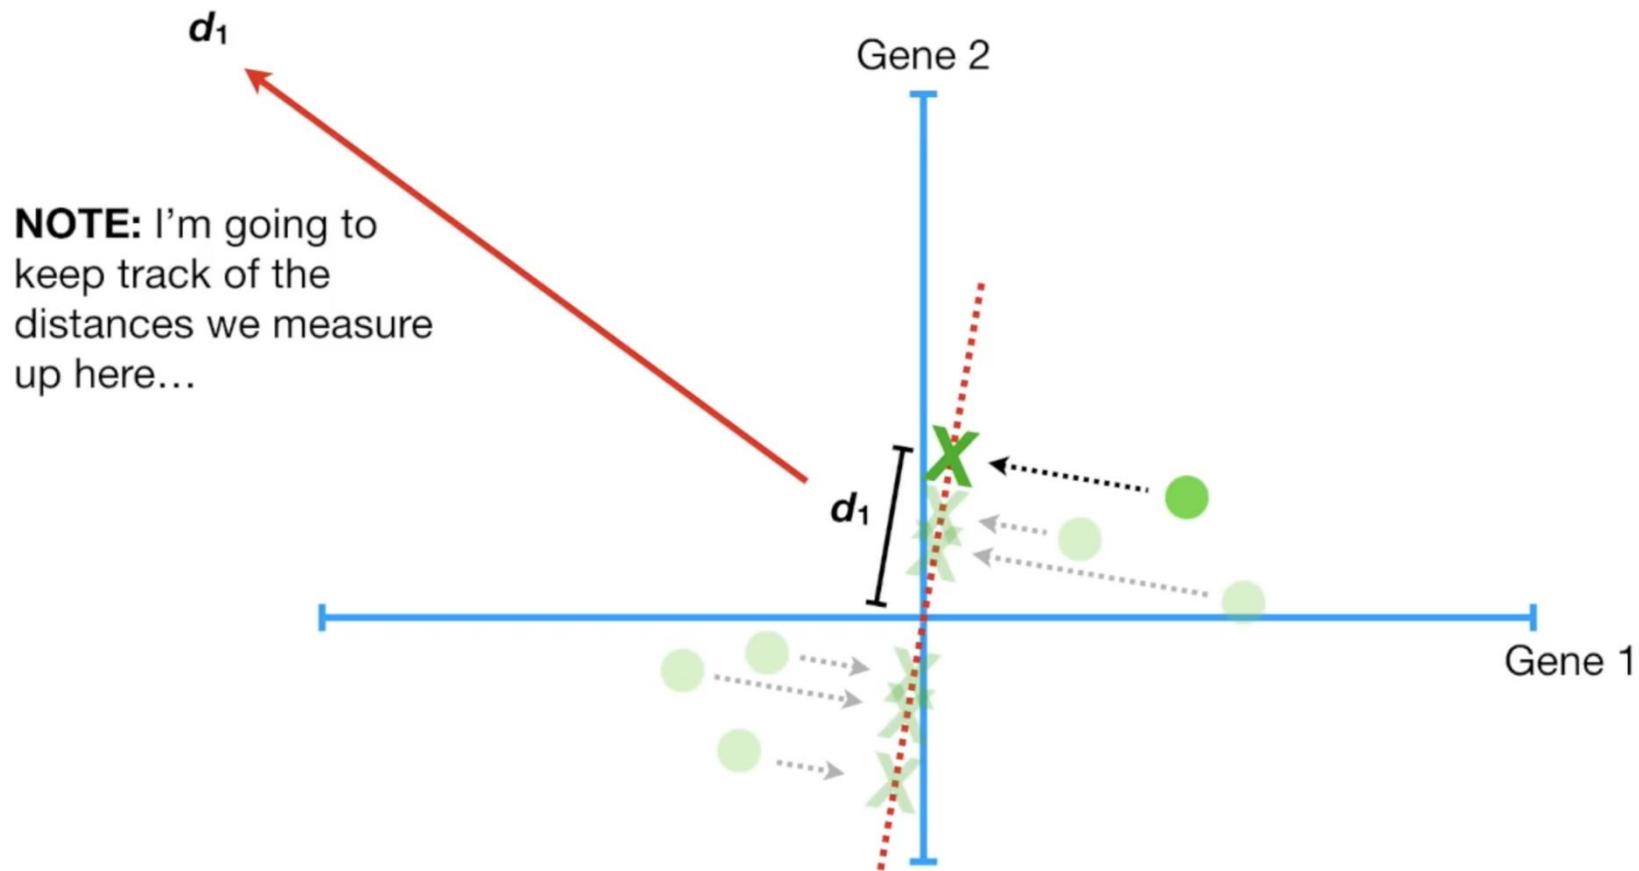


Why?

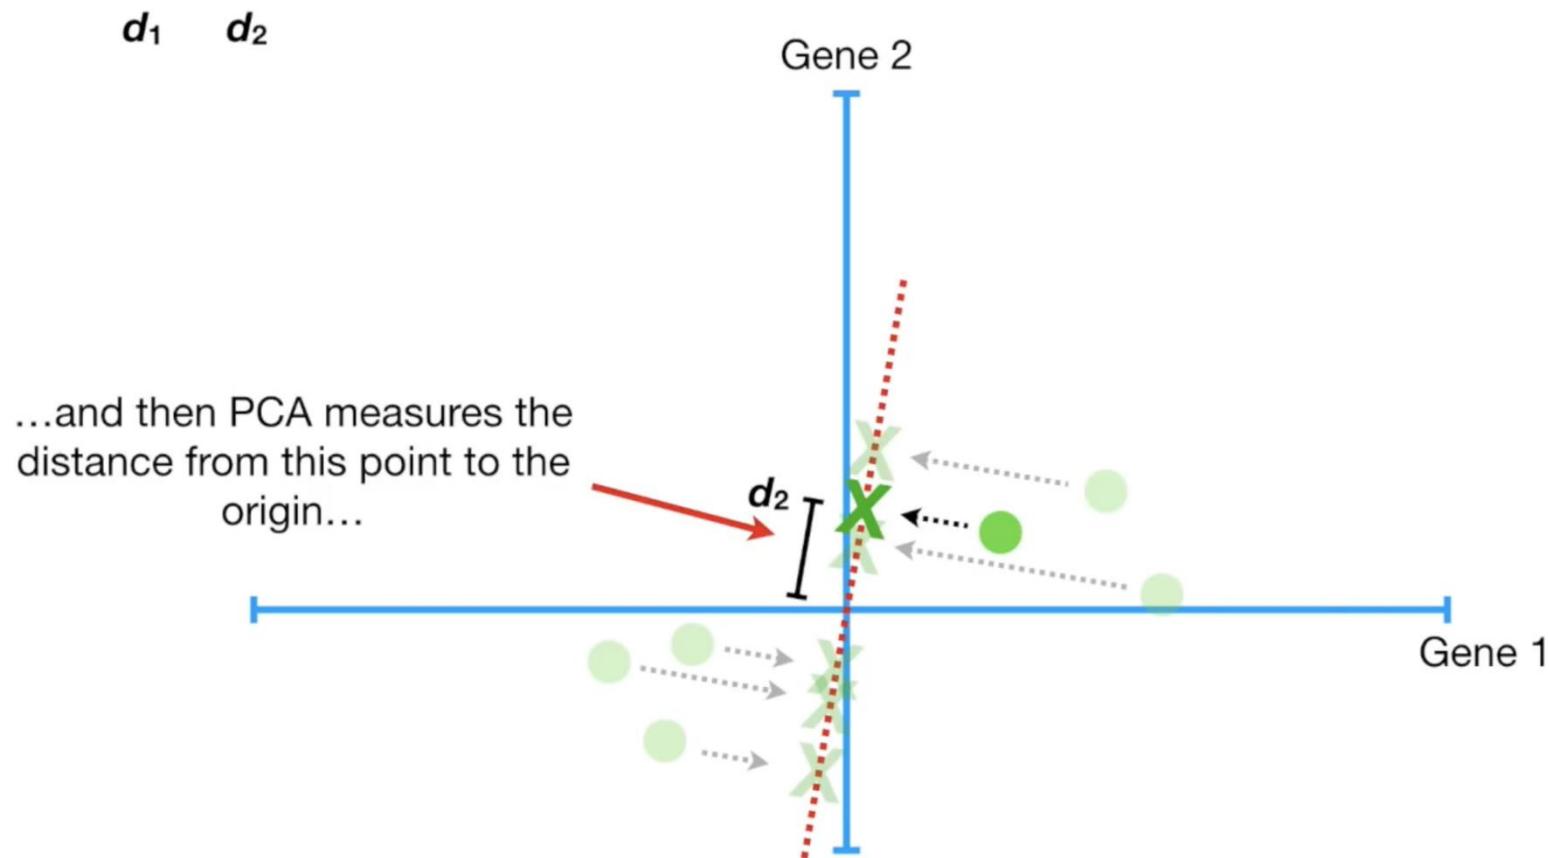
Objective: Minimize **b** or Maximize **c**



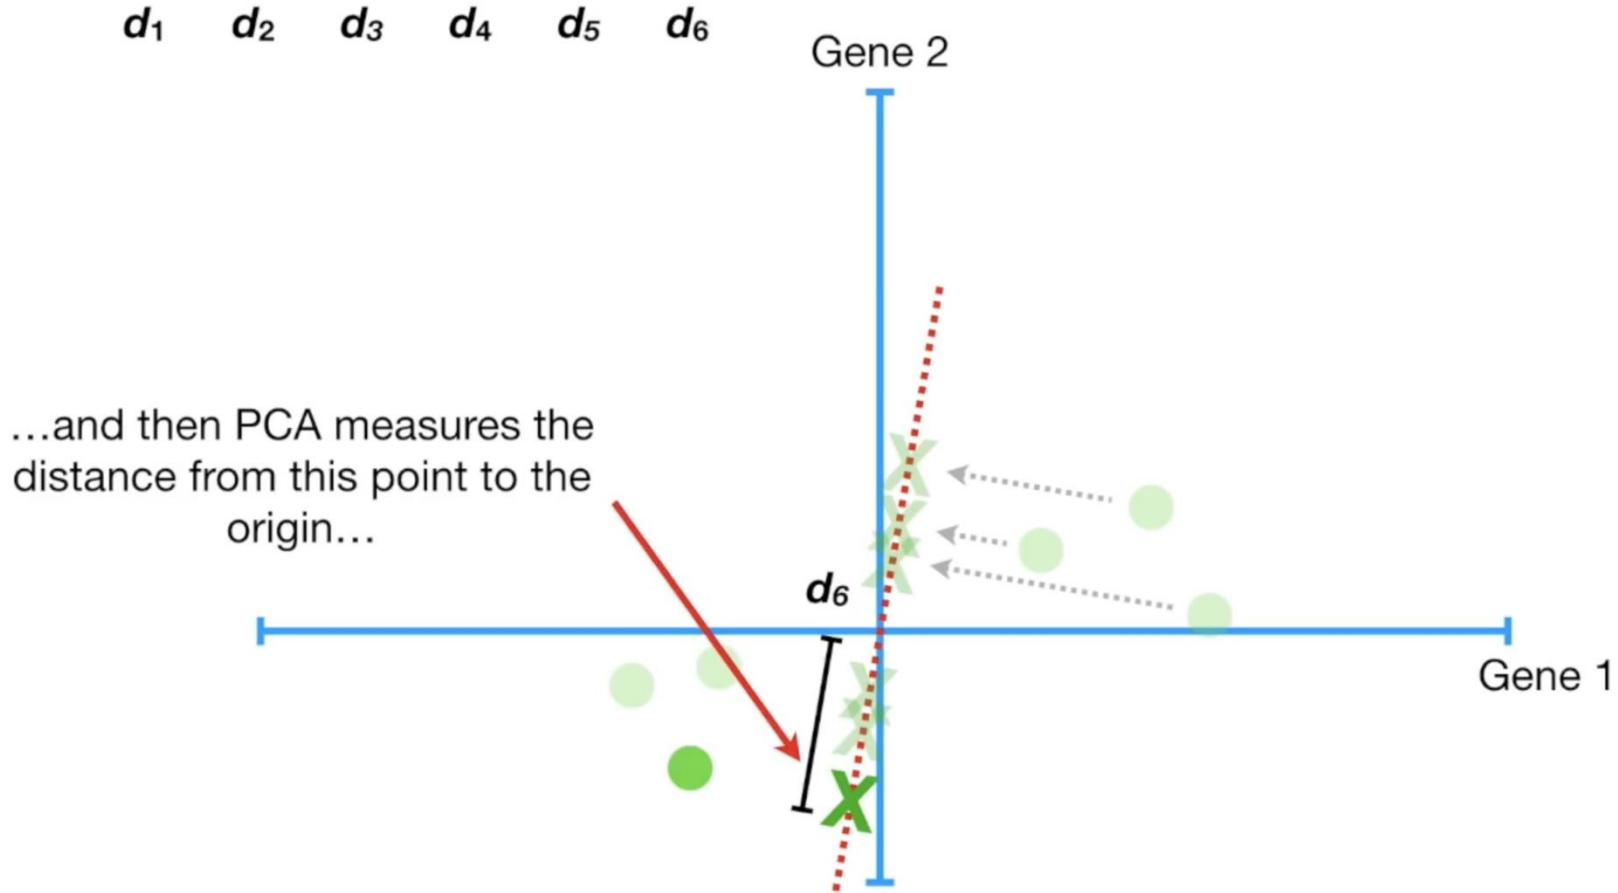
Distances



Distances



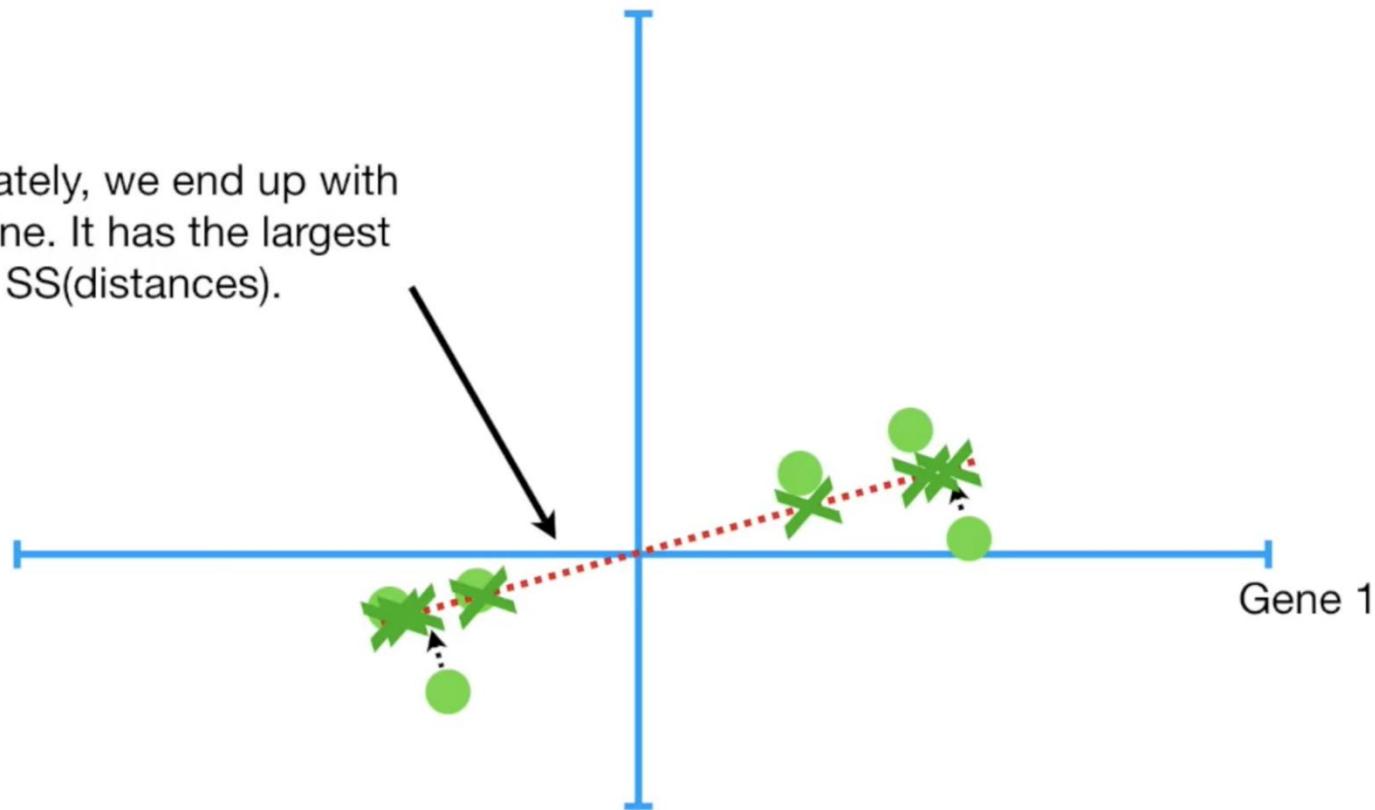
Distances



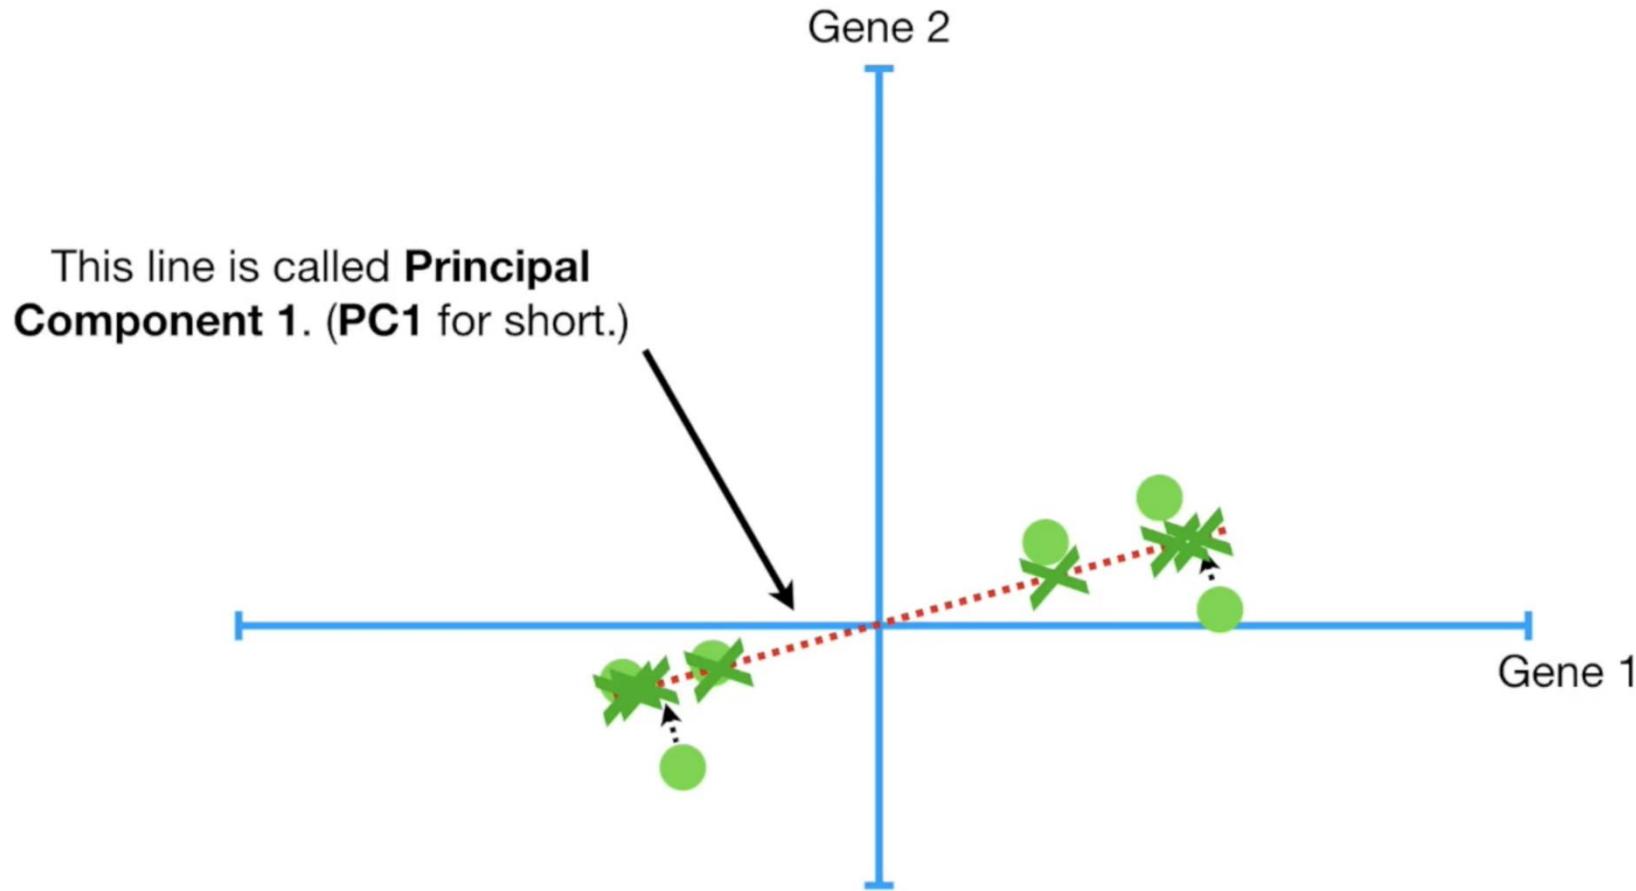
Maximize the Sum of squared distances

$$d_1^2 + d_2^2 + d_3^2 + d_4^2 + d_5^2 + d_6^2 = \text{sum of squared distances} = \text{SS(distances)}$$

Ultimately, we end up with this line. It has the largest SS(distances).

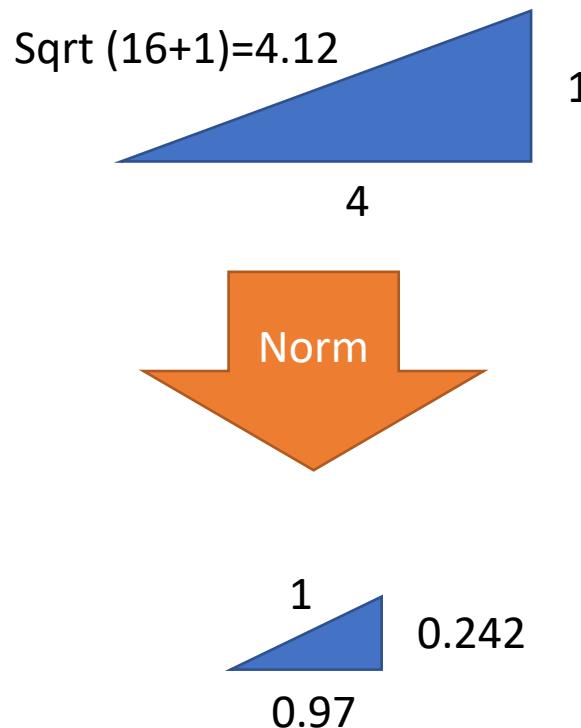


First Principal Component

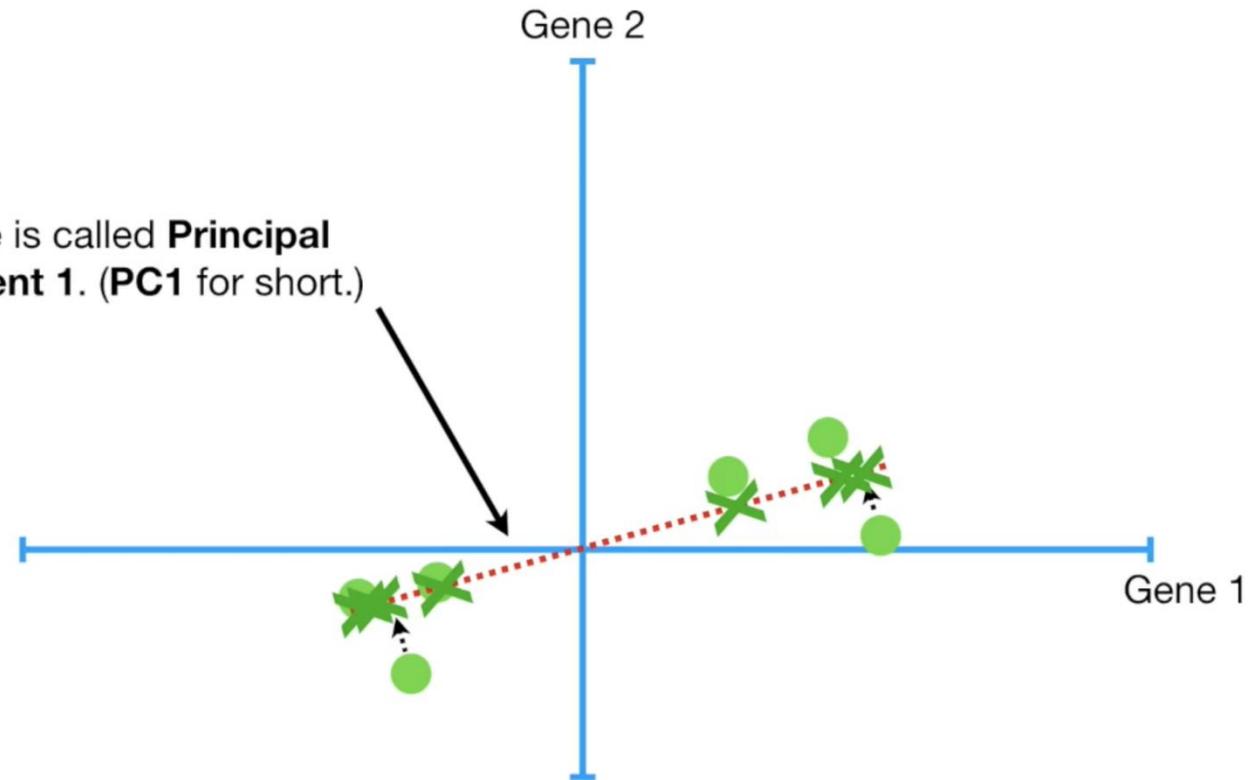


First Principal Component

The slope of the PC1 = 0.25



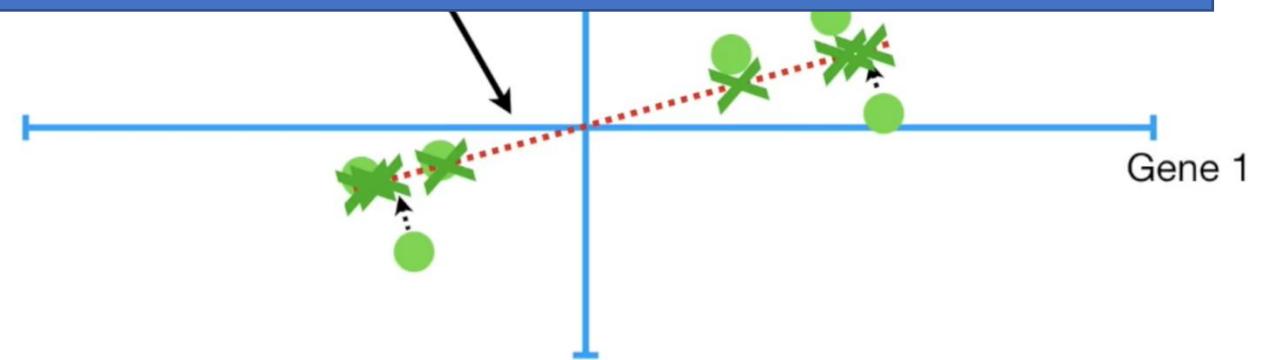
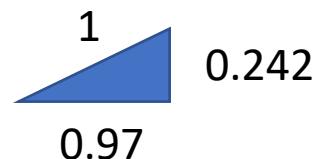
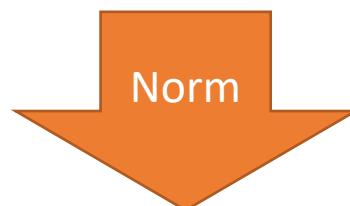
This line is called **Principal Component 1.** (PC1 for short.)



First Principal Component

Eigenvector for PC1 = [0.97, 0.242]

Other name: Singular vector

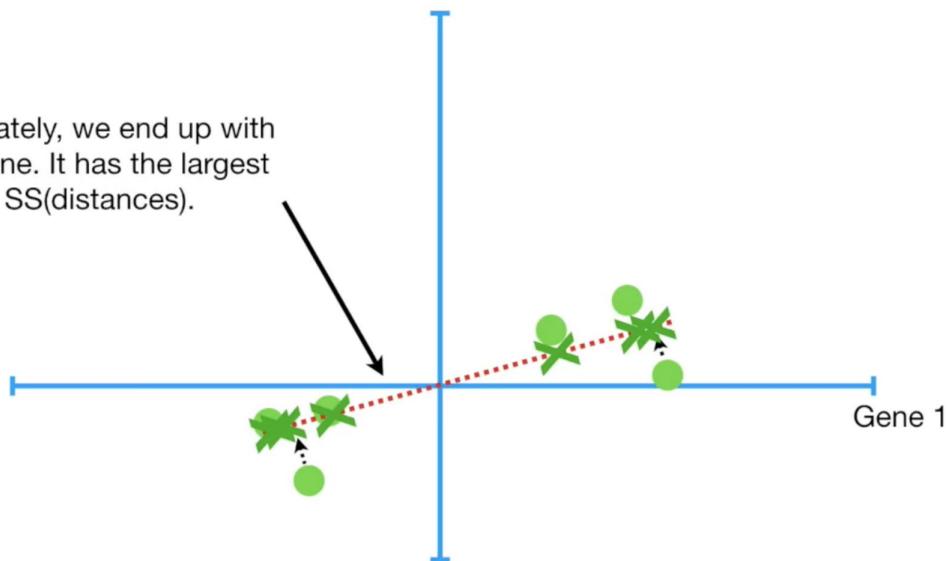


First Principal Component

Sum of squared distances for PC1 = **Eigenvalue** for PC1
Sqrt(Eigenvalue for PC1) = Singular Value for PC1

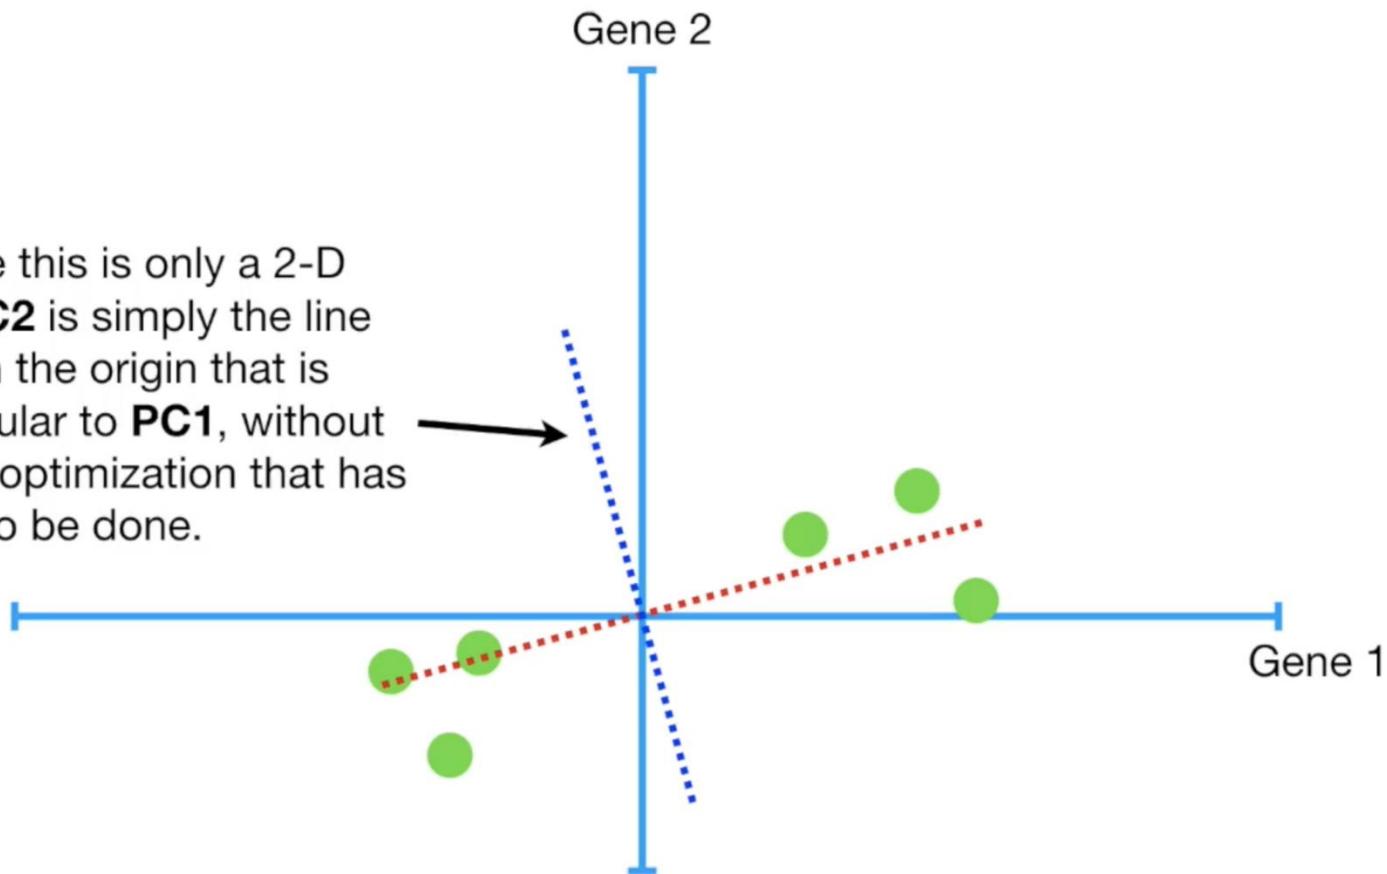
$$d_1^2 + d_2^2 + d_3^2 + d_4^2 + d_5^2 + d_6^2 = \text{sum of squared distances} = \text{SS}(distances)$$

Ultimately, we end up with this line. It has the largest SS(distances).

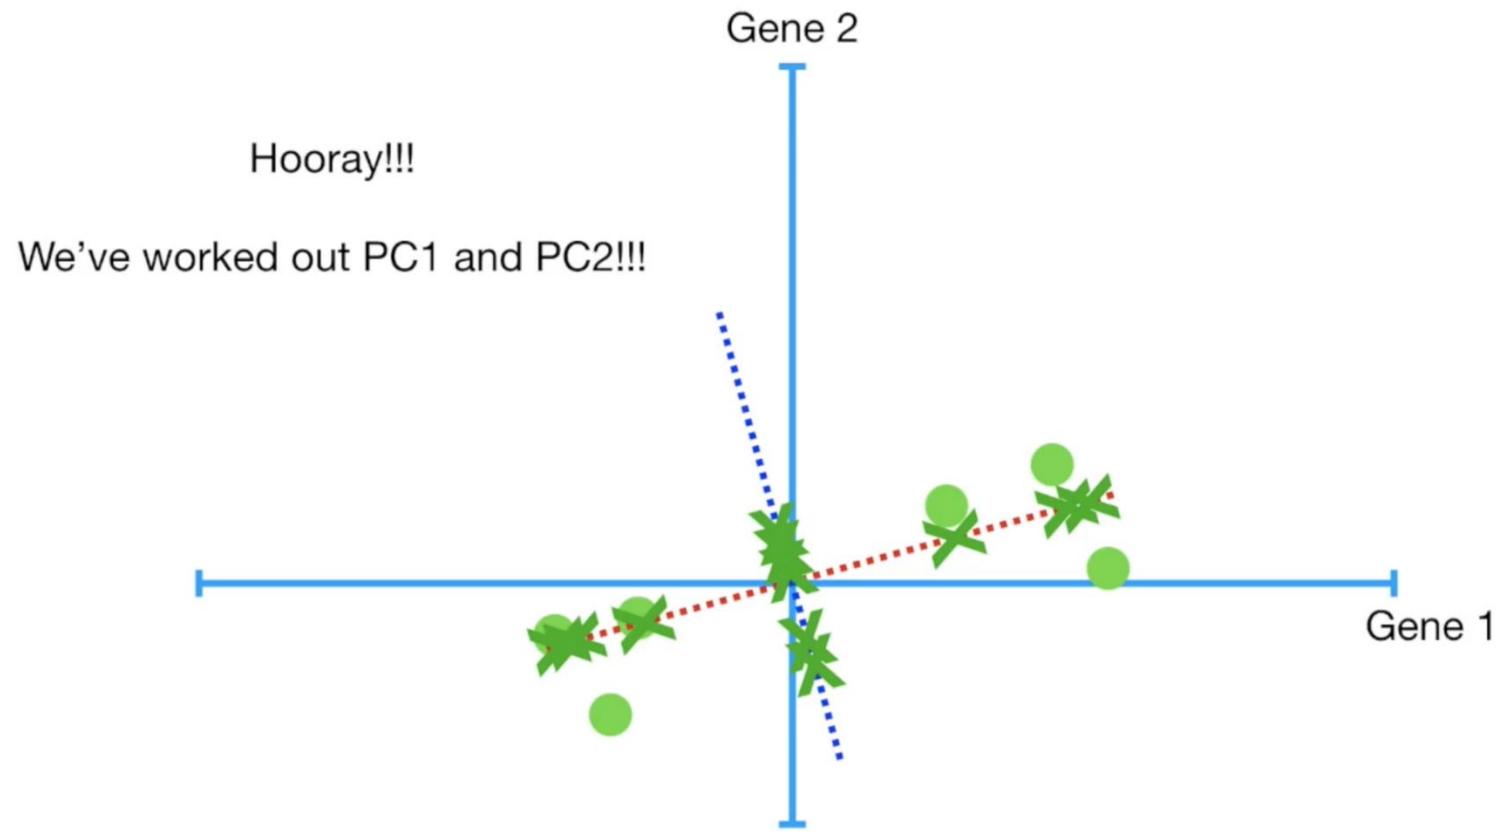


PC2

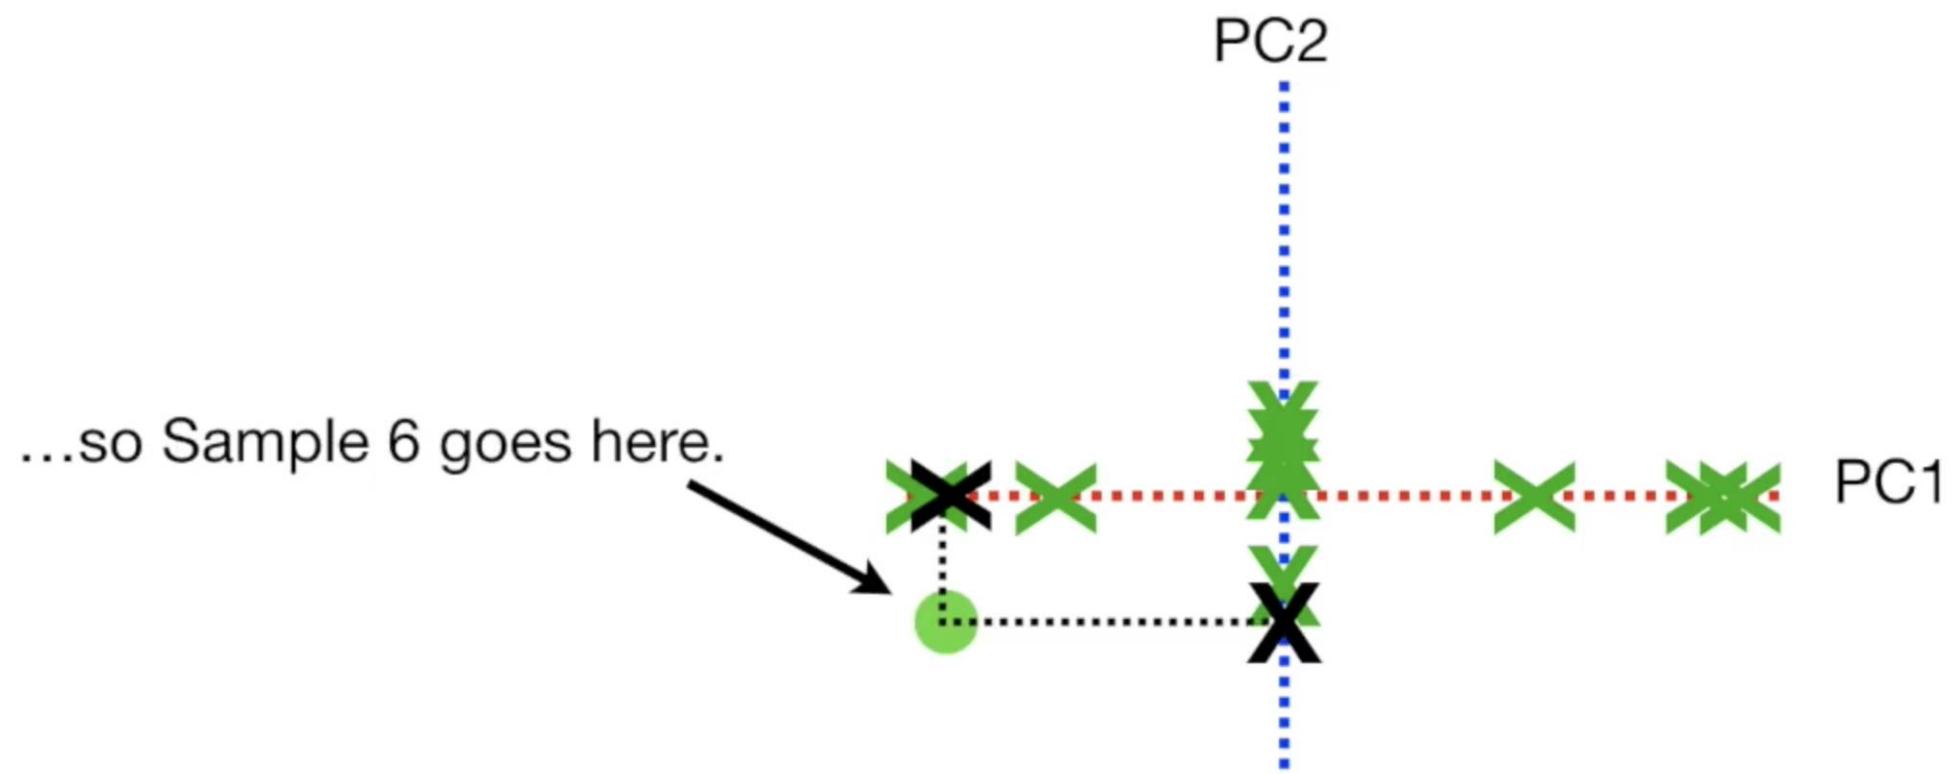
Because this is only a 2-D graph, **PC2** is simply the line through the origin that is perpendicular to **PC1**, without any further optimization that has to be done.



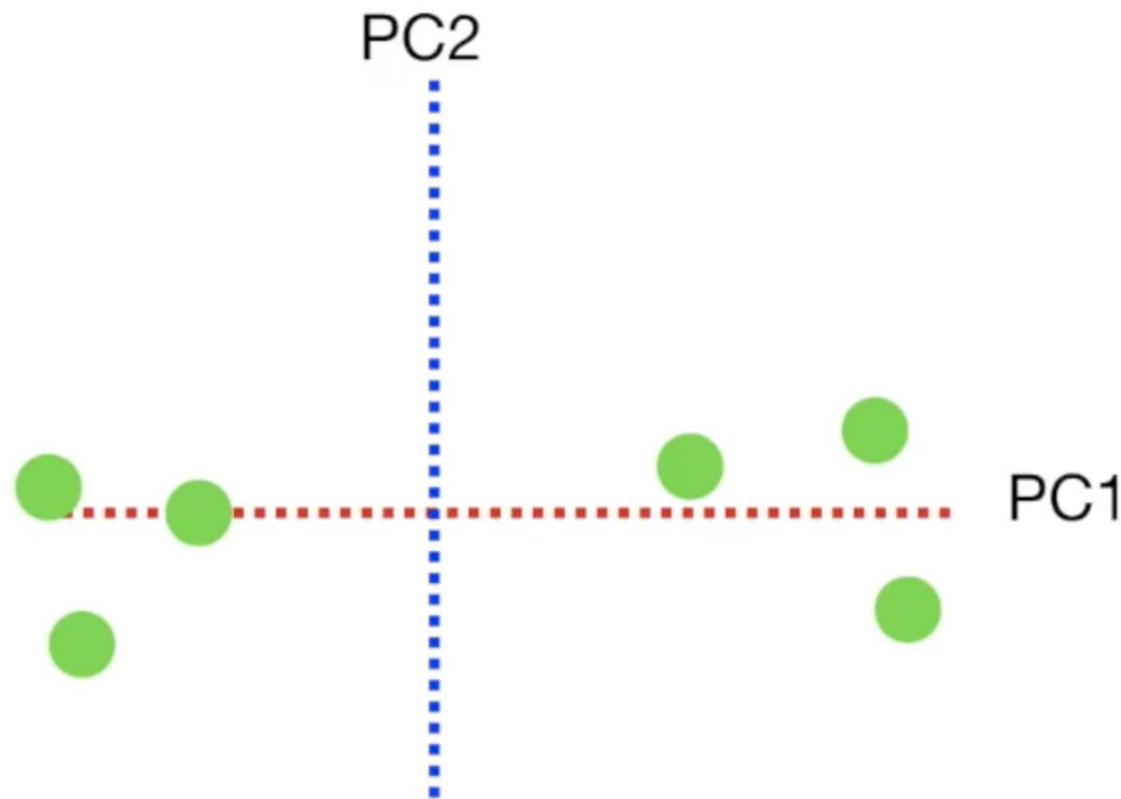
Projection - PC1 and PC2



Reconstruction

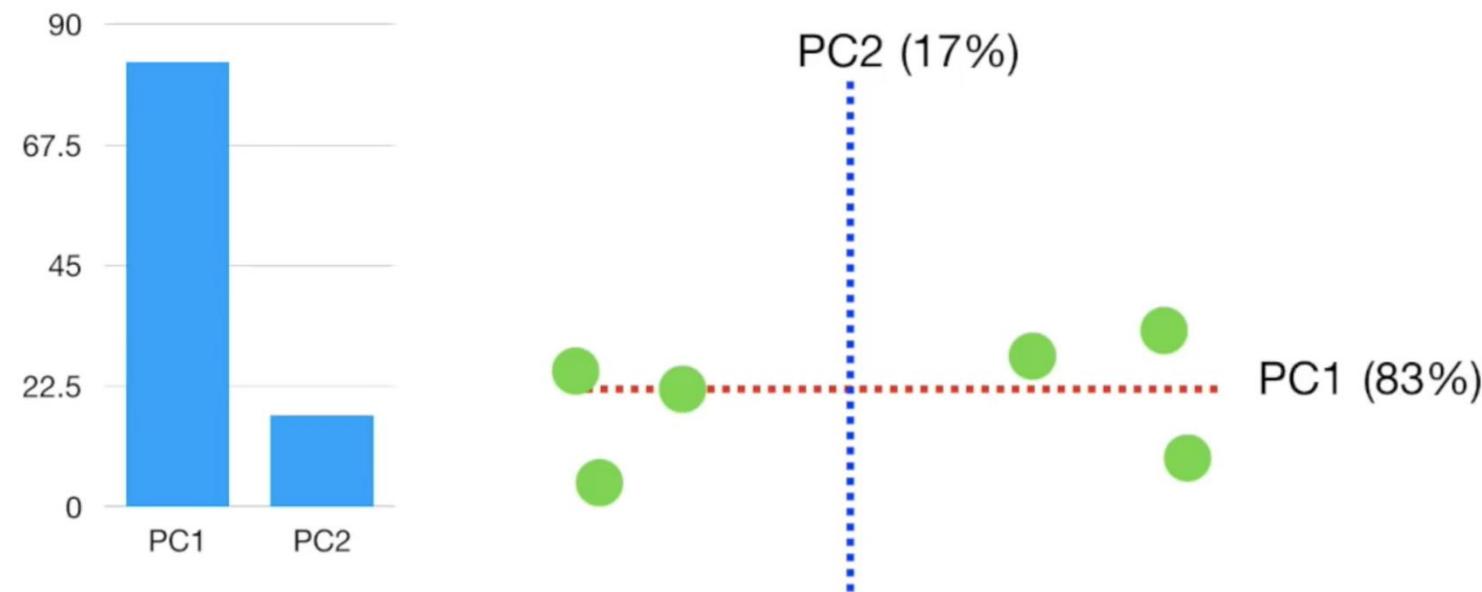


Reconstruction



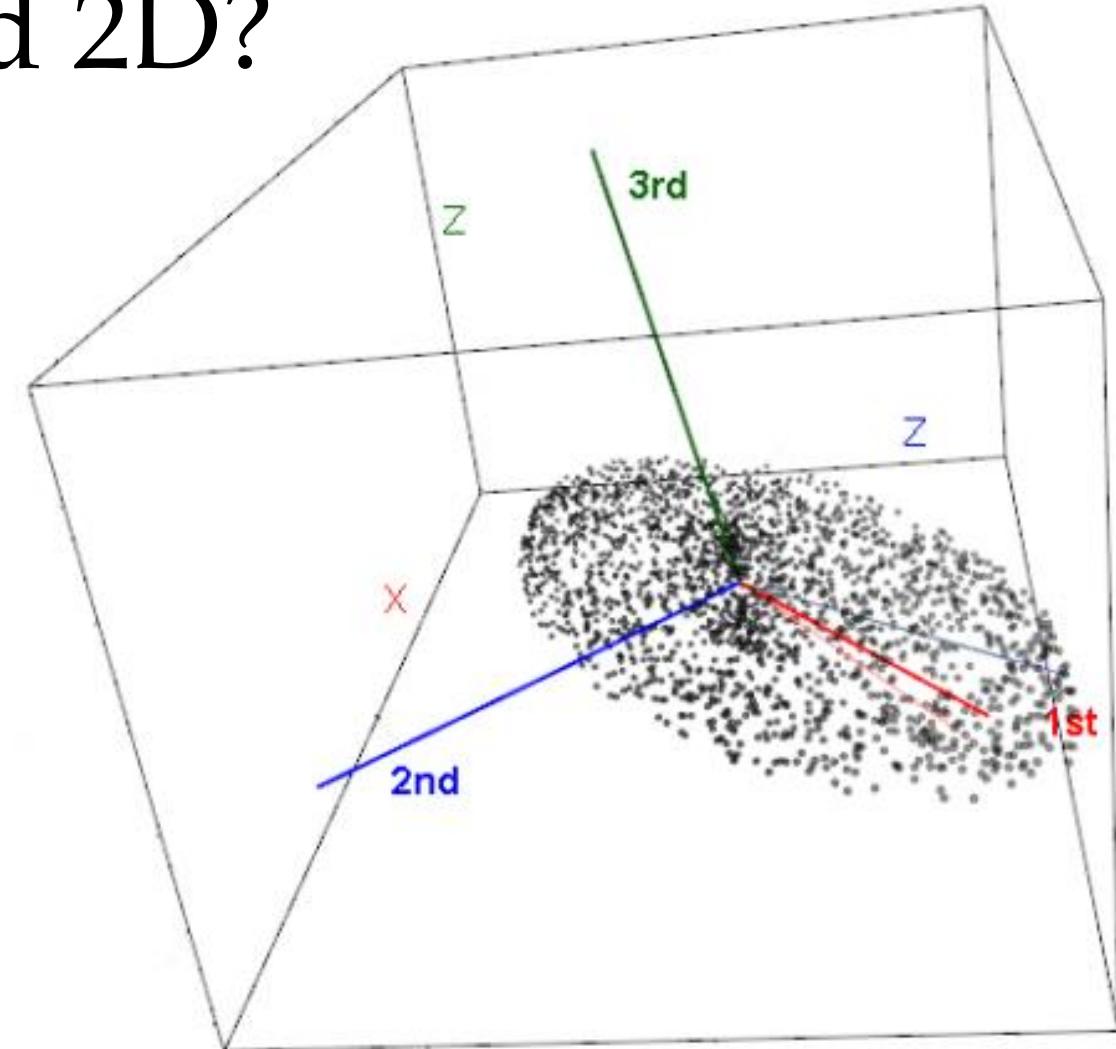
Scree Plot

TERMINOLOGY ALERT!!!! A Scree Plot is a graphical representation of the percentages of variation that each PC accounts for.

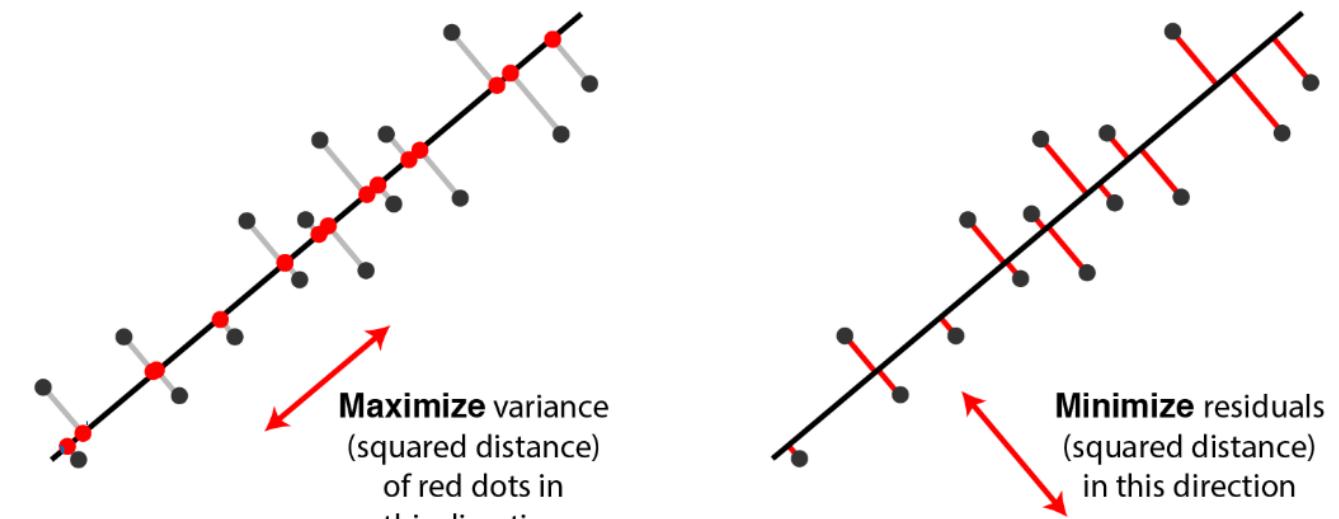
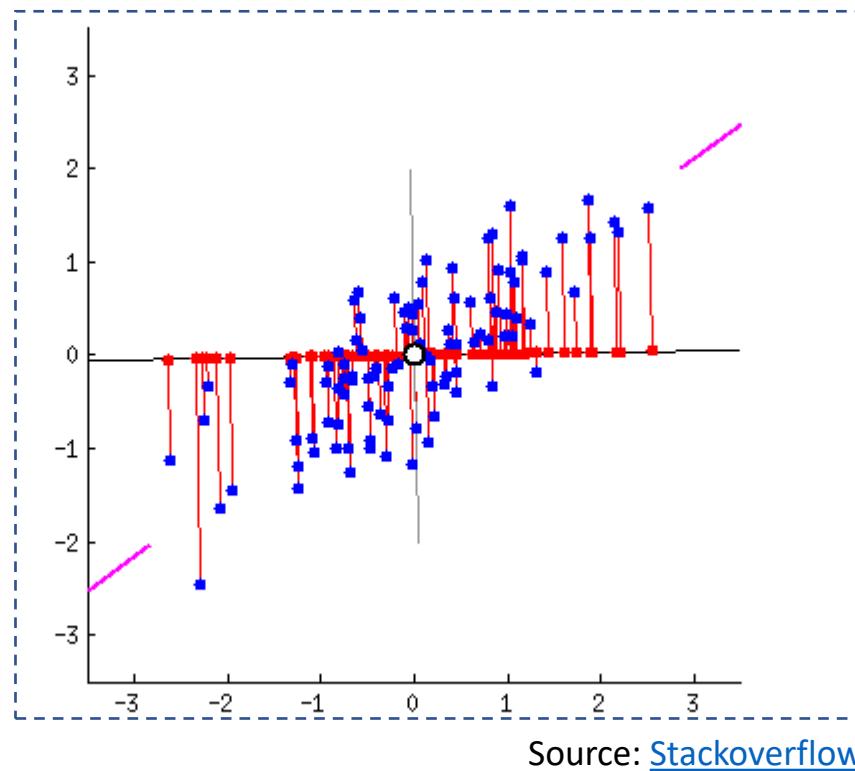


What happens beyond 2D?

- Rigidly rotate the coordinate axes to new (principal axes) such that:
 - principal axis 1 corresponds to the highest variance,
 - axis 2 has the next highest variance,
 - . . . ,
 - and axis p has the lowest variance.
 - All principal axes are uncorrelated (orthogonal).



In Summary...



Source: Alex Williams

- Project D-dimensional data onto an M-dimensional subspace that retains as much information as possible
- Informally: information = diversity = variance

PCA Algorithm Overview

1. Standardize the data (zero mean) and compute the covariance matrix
2. Obtain the Eigenvectors and Eigenvalues of the covariance matrix
3. Choose the k eigenvectors that correspond to the k largest eigenvalues
4. Construct the projection matrix \mathbf{W} from the selected k eigenvectors.
5. Transform the original dataset \mathbf{X} via \mathbf{W} to obtain a k -dimensional feature subspace \mathbf{Y} . ($=\mathbf{X}@\mathbf{W}$)



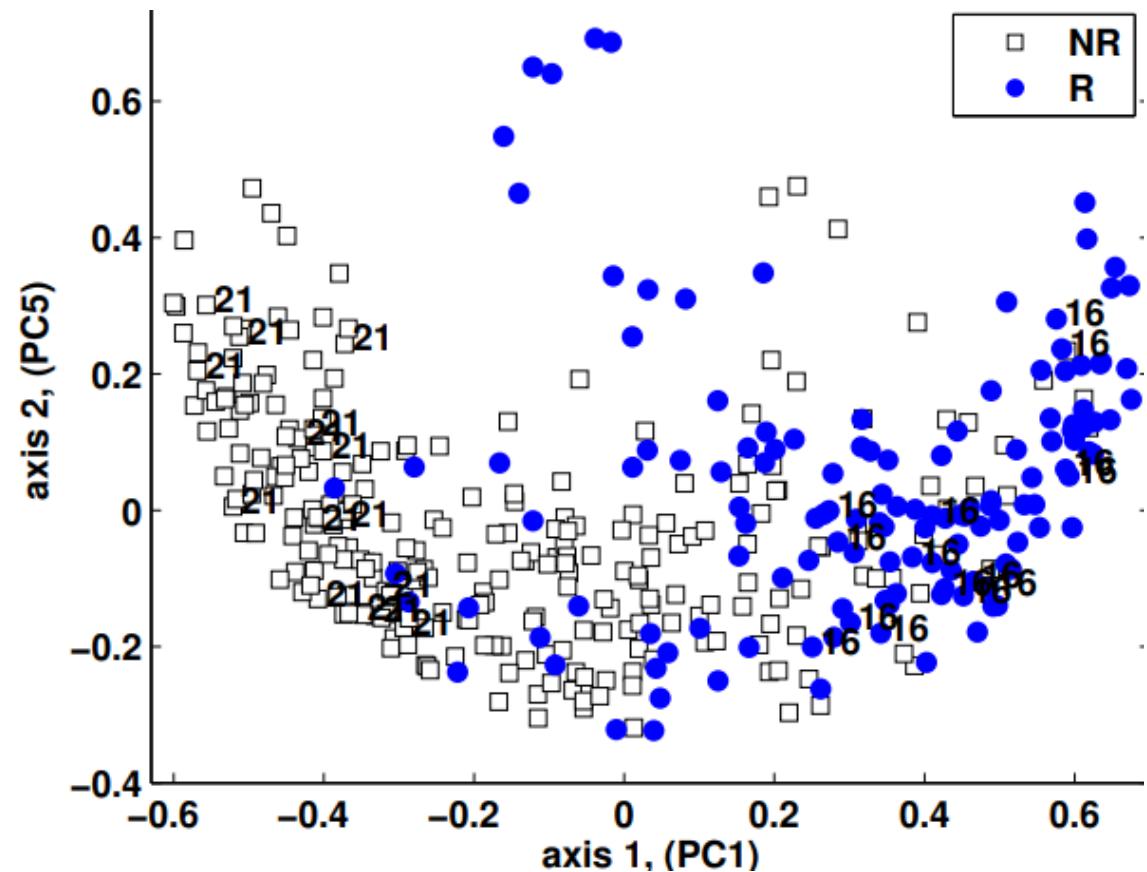
interpretation of data by examining principal components



discard components for **dimensionality reduction**

Visualization of Principle Components

- PCA is commonly applied to datasets to assess information content.
- Example:
 - PCA applied on patient dataset and samples projected on top principal components.
 - Patients with known treatment outcomes are color identified (responders vs non-responders).

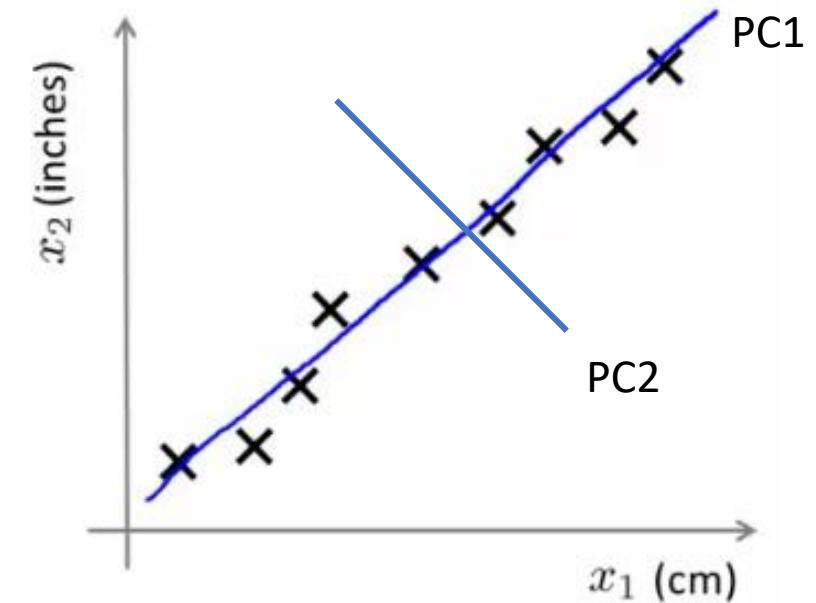
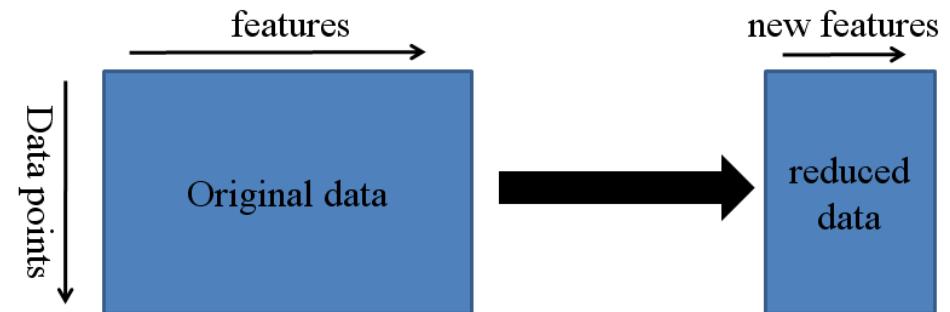


Source: Khodayari-Rostamabad (2011)

Dimensionality Reduction

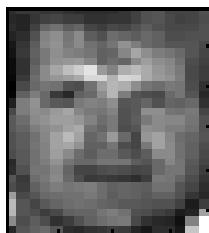
- Projecting onto principal components can improve model/algorithm performance.
 - Reduces the number of parameters and required complexity

- Example:
 - PC1 captures most of the variance (information).
 - PC2 limited variance and projections onto PC2 would look similar (or the same) for all samples.

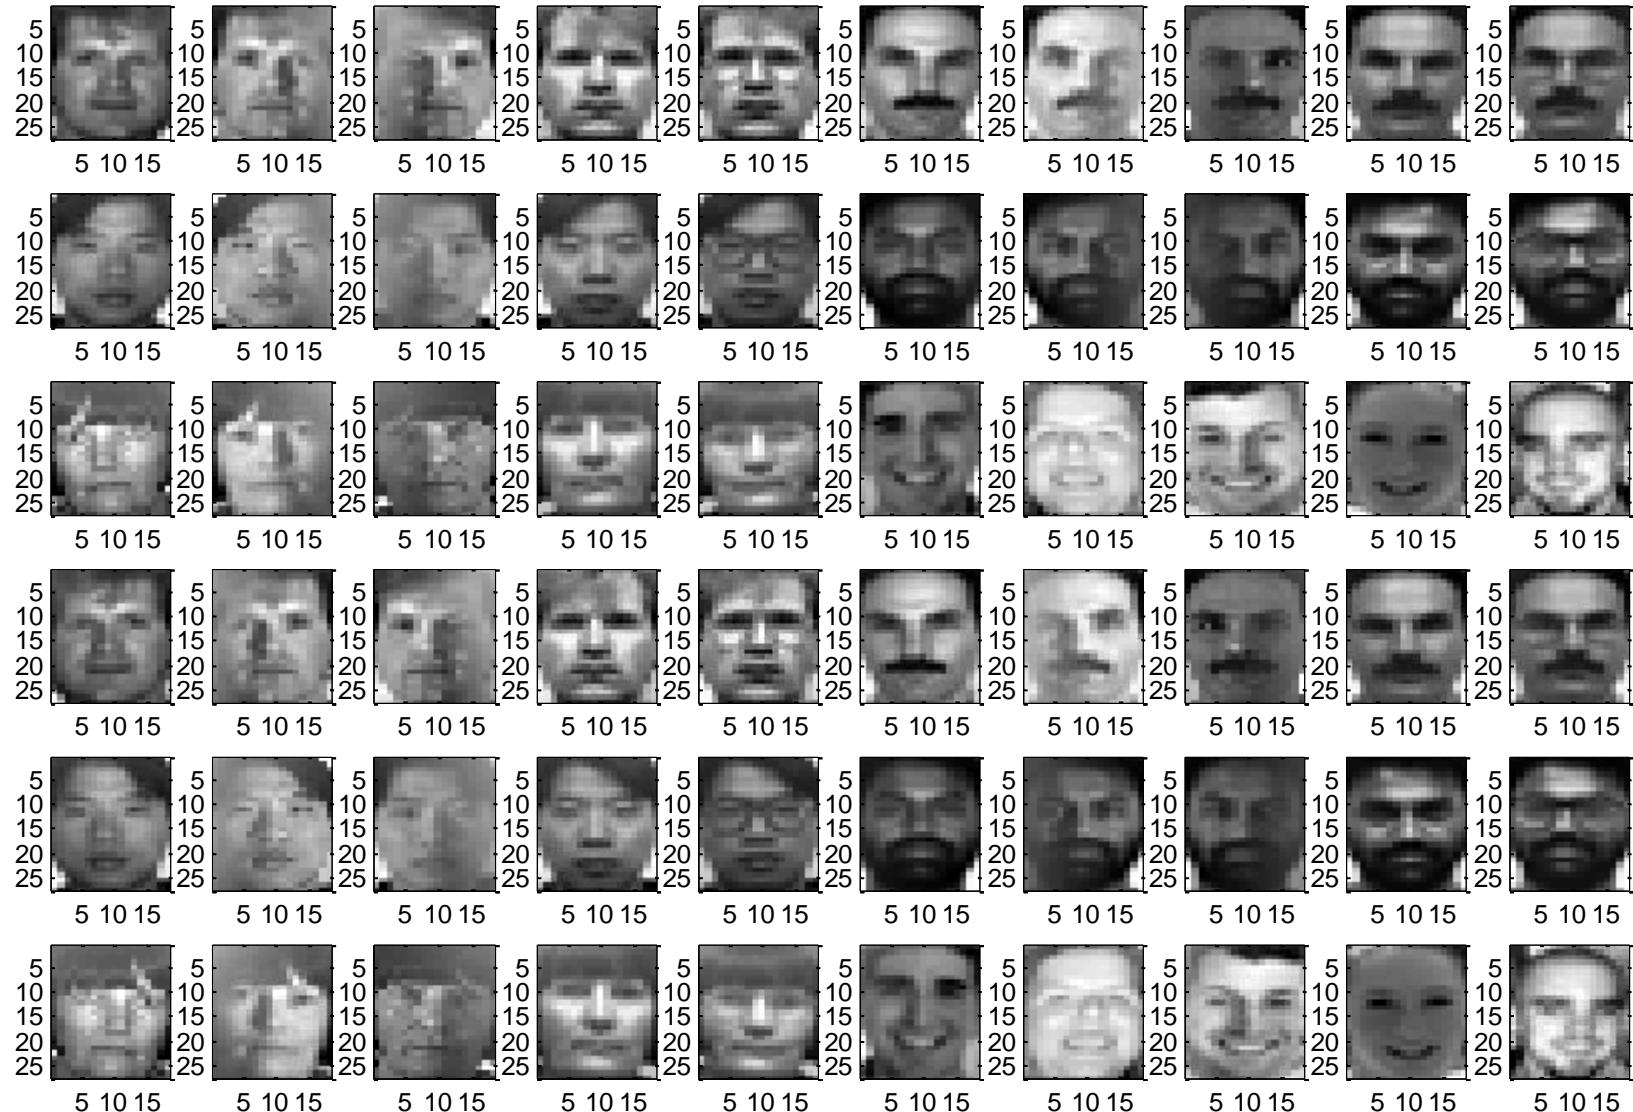


What about Images?

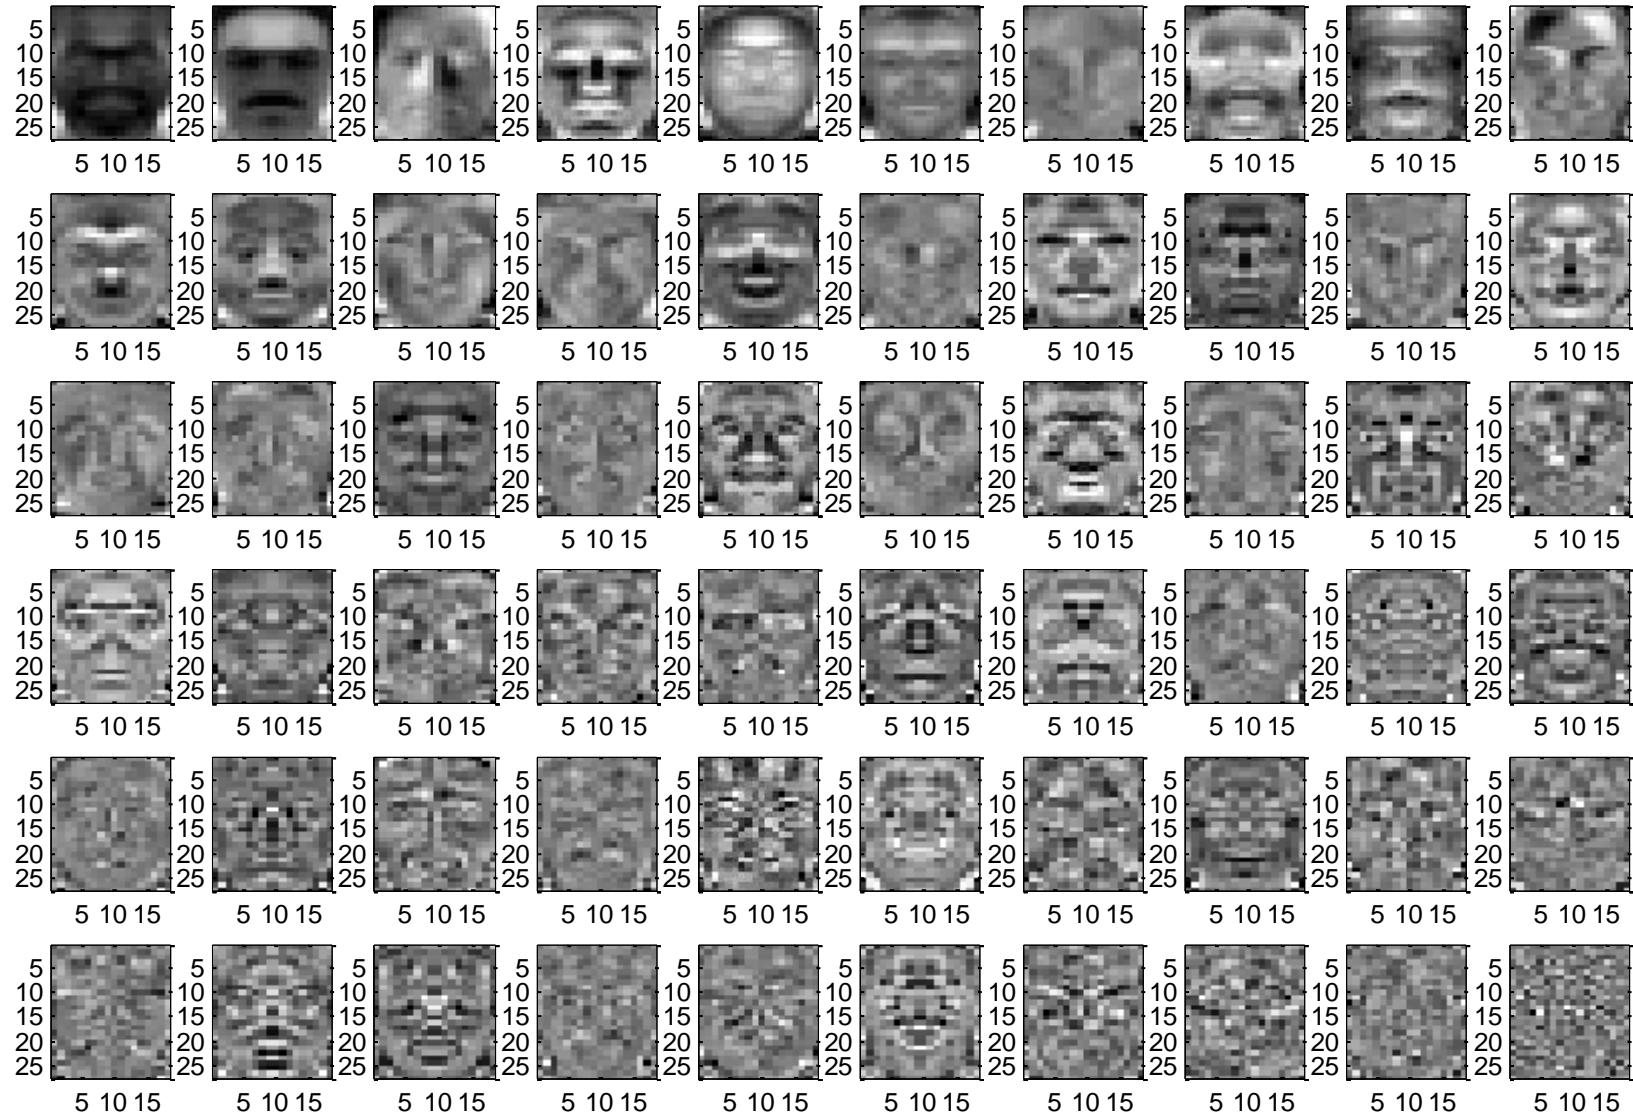
- Images?
 - Yes, we can use PCA! (Turk, Pentland, 1991)
- Each 150 pixel × 150 pixel image can be represented by a 22,500-long vector



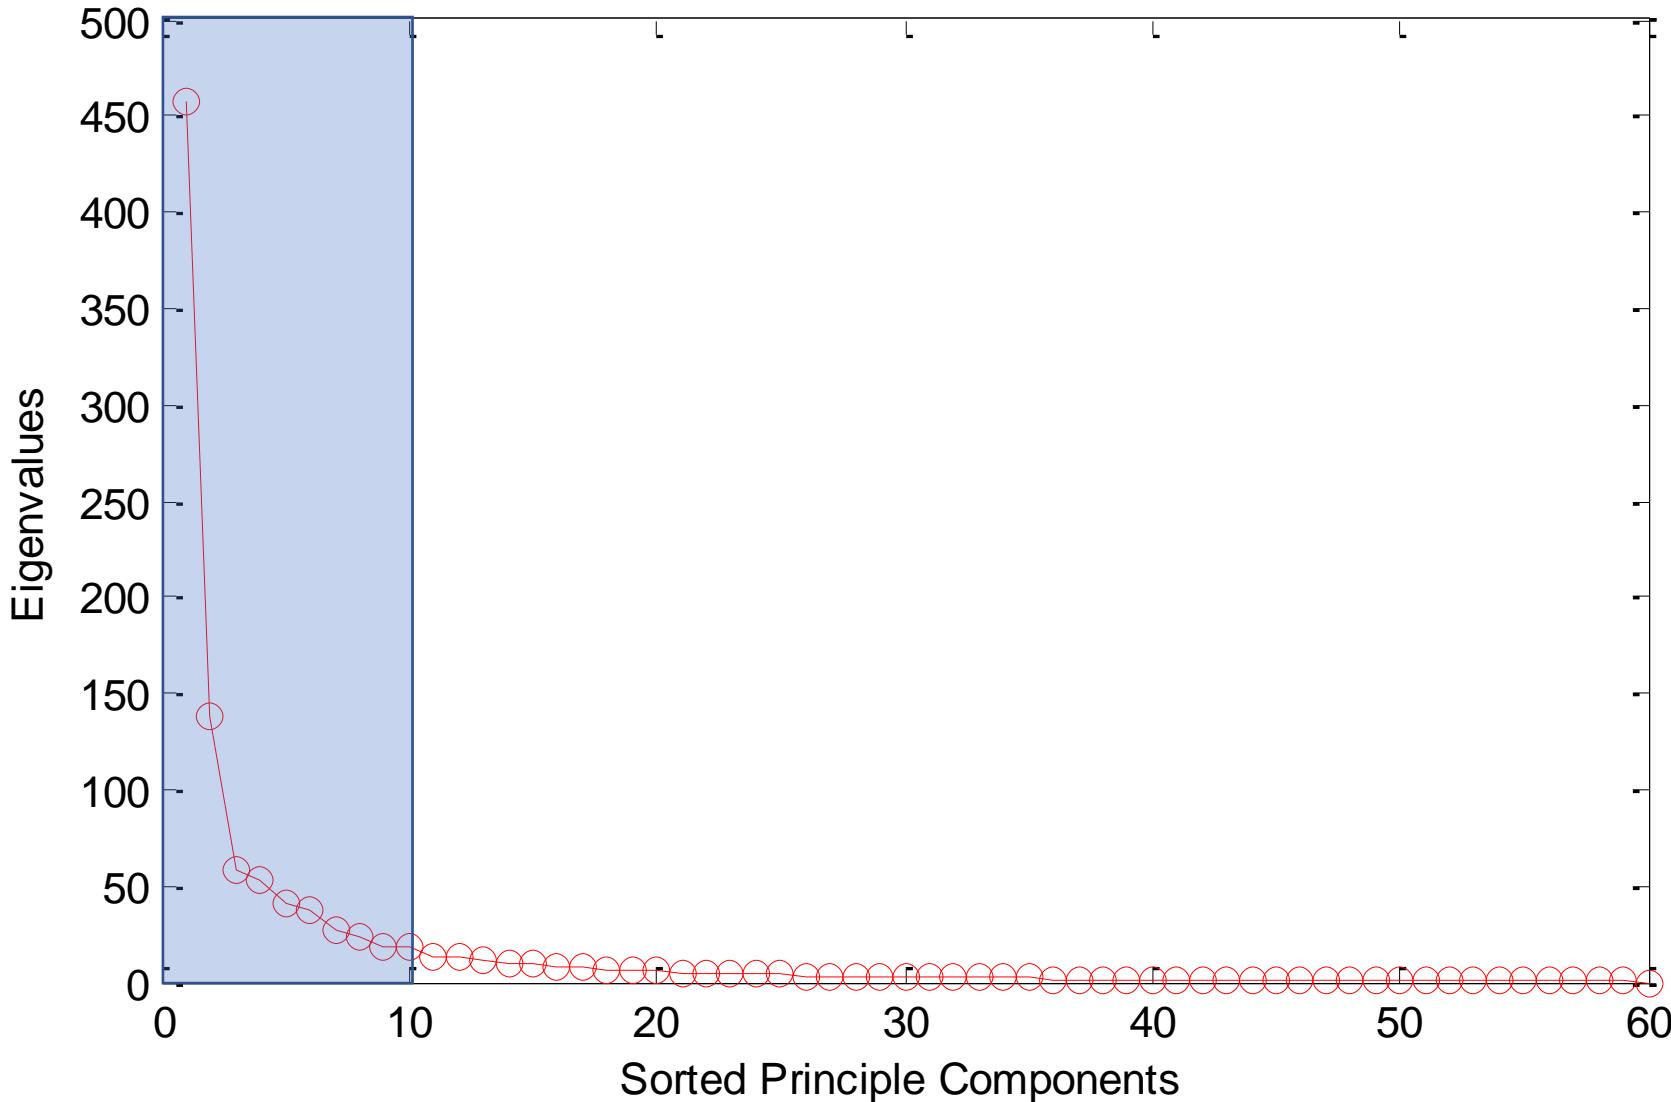
Normalized Face Data



The Eigenfaces



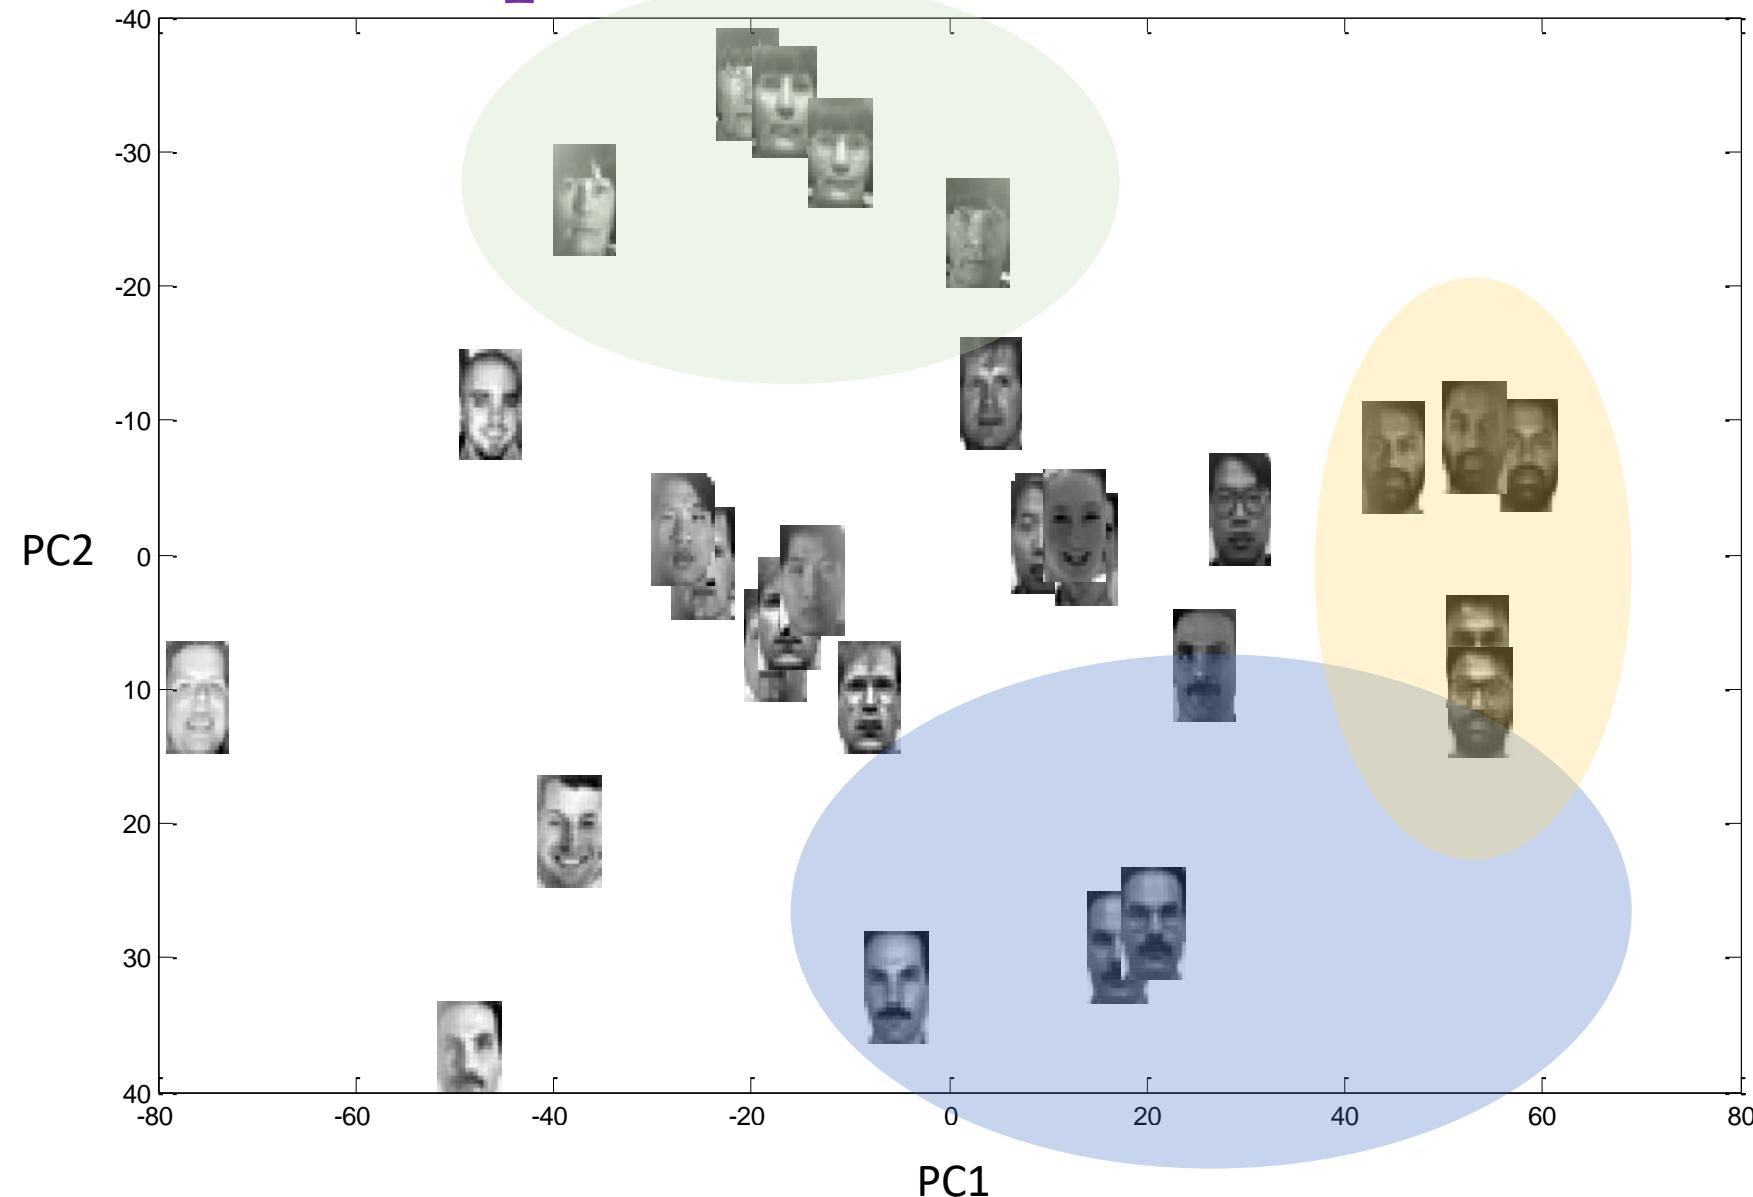
Eigenvalues for each Eigenface



Reconstructing a Face from PCs

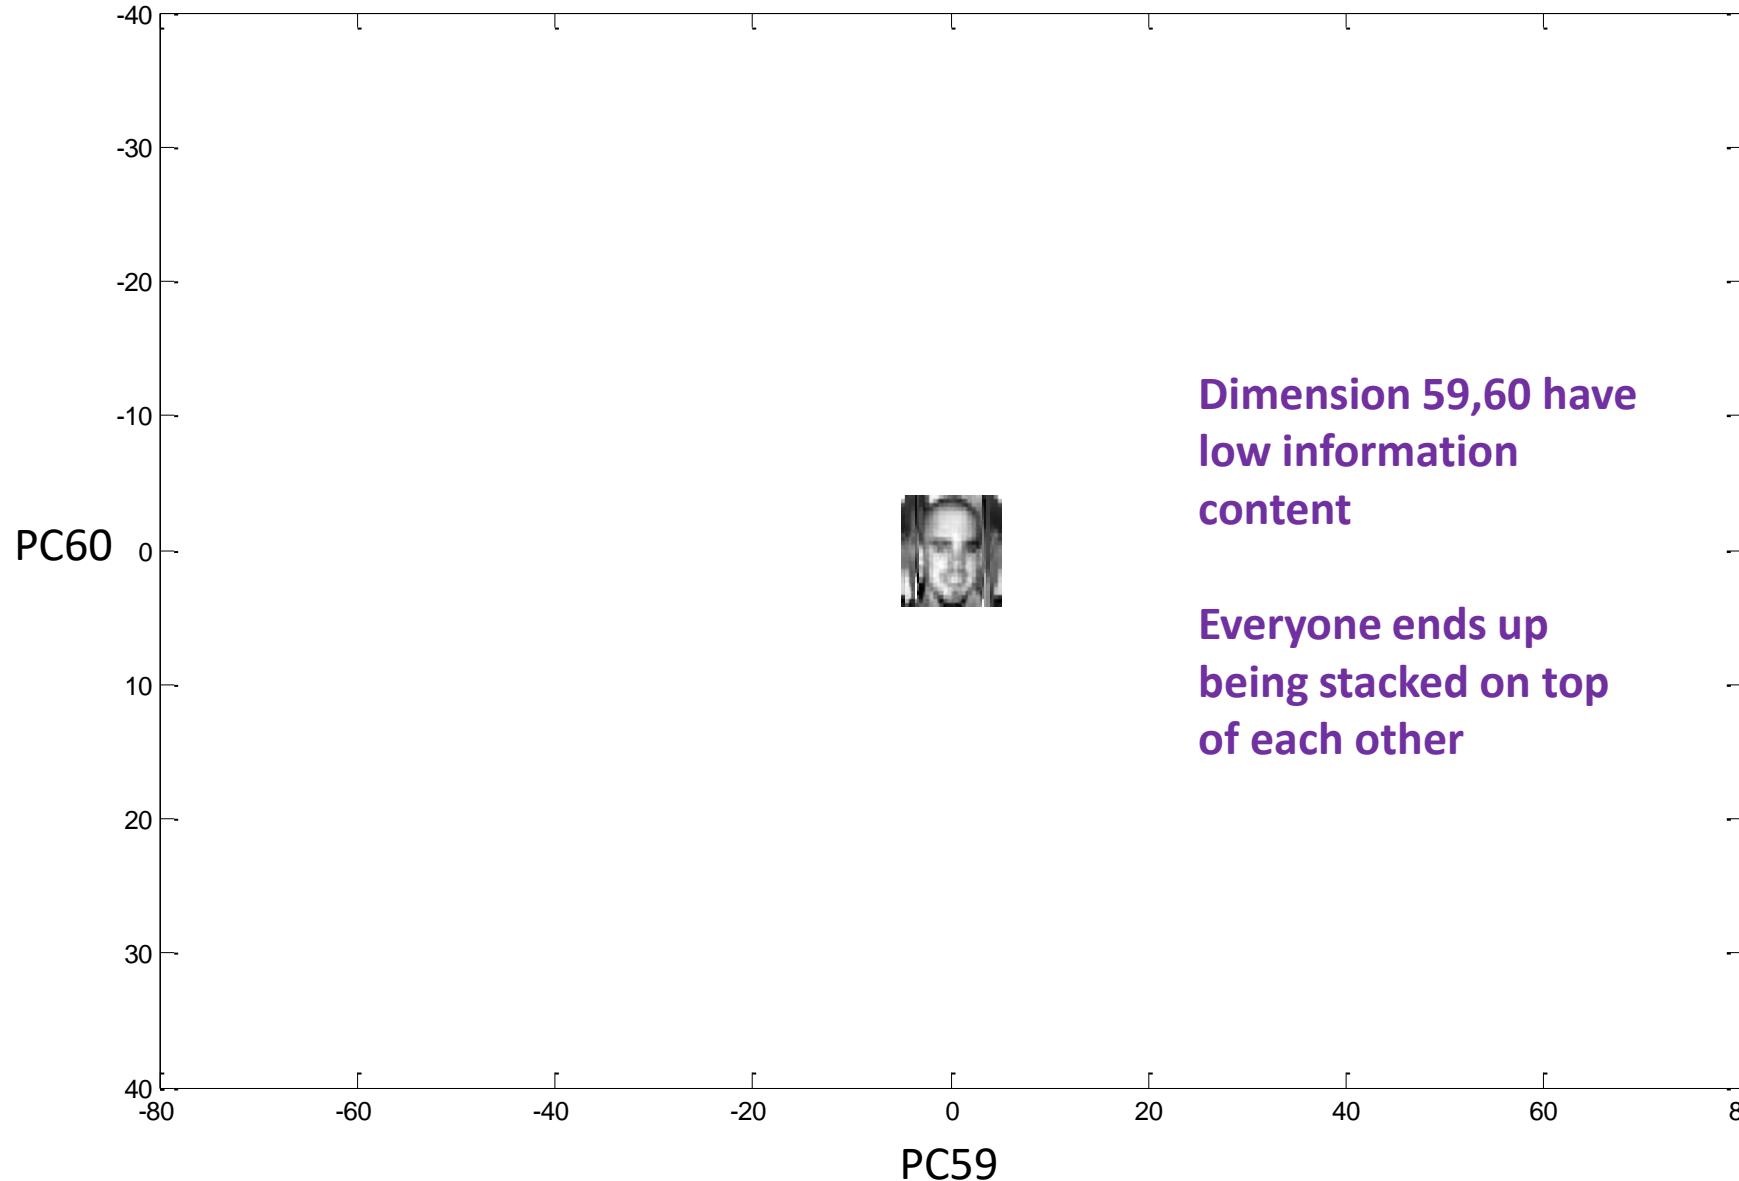


Faces in PC-space (Dims 1,2)



Are there
any patterns
presented?

Faces in PC-space (Dims 59,60)



Data Security

- To reconstruct data we require the mean and standard deviation used for normalization
- Can be used to **secure your datasets** as described for a sample dataset on Kaggle:

“It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to **confidentiality issues, we cannot provide the original features** and more background information about the data. **Features V1, V2, ... V28 are the principal components obtained with PCA**, the only features which have not been transformed with PCA are 'Time' and 'Amount’.”

Source: [Kaggle Credit Card Fraud](#)

Limitations of PCA

- High dimensional data can be problematic. For example, when you have short (small n) and wide data (large p), PCA may not work well.
- We can speed up eigenfaces decomposition by exploiting the smaller dimension of $m \times n$, to do this we need a way to work with rectangular matrices...

Singular Value Decomposition (SVD) to the rescue!

Next Time

- Week 7 Tutorial, Thursday and Friday
 - Project 3 is on PCA and is due 11 March
- Lecture 8 – Dimensionality Reduction Part 2
 - Vector Calculus
 - Singular Value Decomposition (SVD)
 - Interpreting Features
 - Reconstruction
 - Applications
- Guest Lecture on 15 March at 10:00
 - Dr. Sophie Lohmann: “Limits of measurement - who are we measuring?”
 - Zoom link <https://utoronto.zoom.us/j/86722516215>