

APS1070

Foundations of Data Analytics and
Machine Learning

Winter 2022

Week 10:

- *Polynomial Regression*
- *Optimization and Convexity*
- *Regularization*
- *Classification*
- *Neural Networks*

Sinisa Colic and Samin Aref



Slide Attribution

These slides contain materials from various sources. Special thanks to the following authors:

- Lisa Zhang
- Roger Grosse
- Jason Riordon


Last Time

- Linear Regression
 - Empirical Risk Minimization
 - Maximum Likelihood Estimation
 - Negative Log-likelihood

- Today we will continue with **nonlinear regression**.

Nonlinear Regression

Age Prediction




<https://www.how-old.net/>

Stock Market Prediction

Daily Chart – AT&T



Nonlinear Regression



Agenda

- Polynomial Regression
- Convexity and Optimization
- Logistic Regression
- Gradient Descent
- Regularization
- Multiclass Classification
- Neural Networks



Theme:
**Nonlinear Regression
(and Classification)**

Nonlinear Regression

Recap: Linear Regression

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters:

$$\theta_0, \theta_1$$

Cost Function:

$$J(\theta_0, \theta_1) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal:

$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1) \quad \left. \vphantom{\underset{\theta_0, \theta_1}{\text{minimize}}} \right\}$$


(1) direct solution

$$\frac{dJ}{d\theta} = 0$$

OR

(2) gradient descent

$$\theta \leftarrow \theta - \alpha \frac{\partial J}{\partial \theta}$$



Recap: Vectorization

Hypothesis:

$$h_{\theta}(\mathbf{X}) = \mathbf{X}\theta$$

Parameters:

$$\theta$$

Cost Function:

$$J(\theta) = \frac{1}{2N} \|\mathbf{y} - \hat{\mathbf{y}}\|^2$$

Goal:

$$\underset{\theta}{\text{minimize } J}$$


(1) direct solution

$$\frac{dJ}{d\theta} = 0$$



OR

(2) gradient descent

$$\theta \leftarrow \theta - \alpha \frac{\partial J}{\partial \theta}$$




Recap: Convexity



Nonlinear Regression


➤ Suppose we want to model the following data



Given noisy **sample data** we want to find a **hypothesis** for what **generated the data**

➤ Cannot be fit with a linear model...

Nonlinear Regression



➤ **One option** is to fit a low-degree polynomial:

$$\hat{y} = \theta_3 x^3 + \theta_2 x^2 + \theta_1 x + \theta_0$$

➤ This is known as **polynomial regression**

Q: Does this mean we have to derive a whole new algorithm?

Feature Mapping

- Implement a polynomial transformation (feature mapping) by replacing input with polynomial of increasing order:

$$\psi(x) = \begin{pmatrix} 1 \\ x \\ x^2 \\ x^3 \end{pmatrix}$$

- Hence our hypothesis can be written as:

$$\hat{y} = \theta_3 x^3 + \theta_2 x^2 + \theta_1 x + \theta_0 (1)$$

$$\hat{y} = \boldsymbol{\theta}^T \psi(x)$$

The derivations and algorithms from last lecture remain the same! Why?


Polynomial Regression

- This doesn't require changing the algorithm, just **pretend $\psi(x)$ is the input vector**.

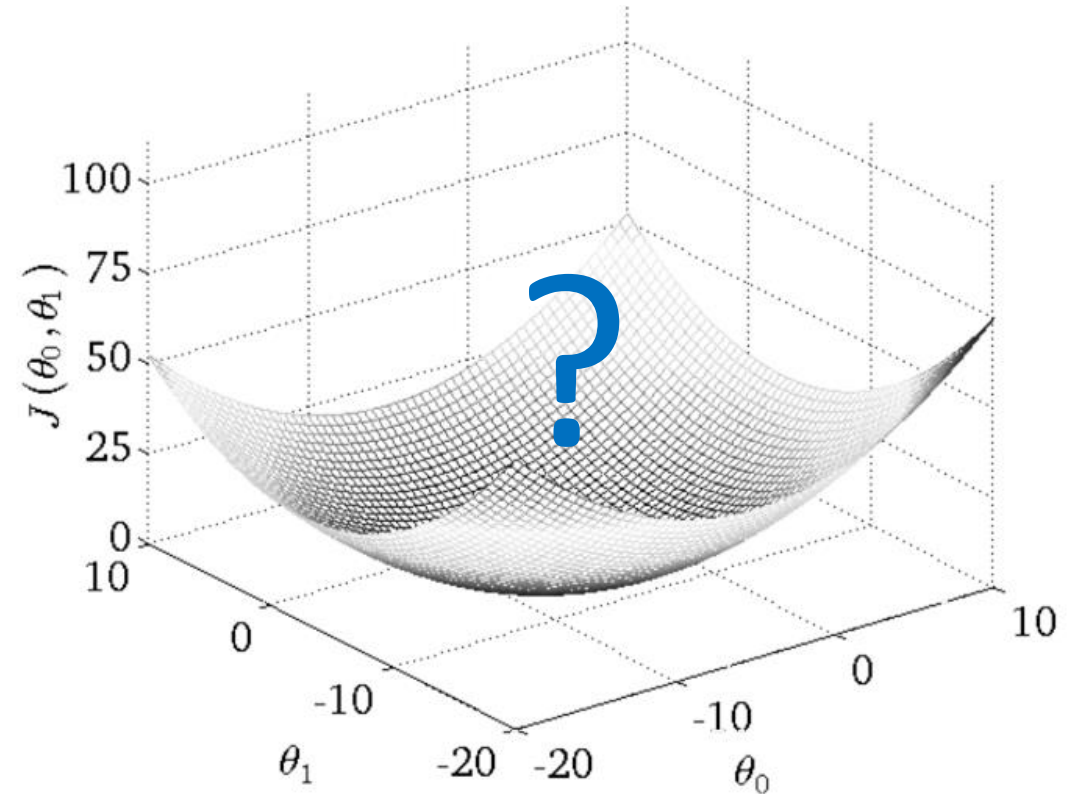
$$\hat{y} = \theta^T \psi(x) \quad \psi(x) = \begin{pmatrix} 1 \\ x \\ x^2 \\ x^3 \end{pmatrix}$$

- Feature maps let us fit nonlinear models
- Before deep learning, most of the effort in building a practical machine learning system was **feature engineering**.

Q: Convexity of Polynomial Regression?



$$\hat{y} = \theta^T \psi(x)$$



Direct Solution

- Polynomial regression is really a linear regression problem with some feature engineering.

$$\hat{y} = \theta^T \psi(x) \quad \xrightarrow{\text{minimize cost}} \quad \theta = (\psi^T \psi)^{-1} \psi^T y$$


$J(\theta) = \frac{1}{2N} \|\mathbf{y} - \hat{\mathbf{y}}\|^2$

Requires that all columns are linearly independent

- Require that $\psi^T \psi \in \mathbb{R}^{D \times D}$ to be invertible. This is the case if and only if $\text{rank}(\psi) = D$.


Fitting Polynomial ($M = 0$)

$$\hat{y} = w_0$$




Fitting Polynomial ($M = 1$)

$$\hat{y} = w_0 + w_1 x$$




Fitting Polynomial (M = 3)

$$\hat{y} = w_0 + w_1x + w_2x^2 + w_3x^3$$



Fitting Polynomial (M = 9)


$$\hat{y} = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_9x^9$$



Fitting Polynomial (M = 9)



$$\hat{y} = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_9x^9$$

OVERFIT




Generalize to New Samples

- We could give the hypothesis a higher complexity, or capacity to fit the data, but this may not generalize well to **new samples**



Generalization

➤ Training and test error as a function of # parameters:




Generalization


➤ Training and test error as a function of # training examples:

Fixed model: Polynomial of degree 2

①



②



Training error: Zero


Test error: $\gg 0$

Their difference: Huge

Training error: Some nonzero value a

Test error: $a + \epsilon$

Their difference: ϵ




Regularization

- The degree of the polynomial is a hyperparameter, just like k in KNN. We can tune it using a validation set.
- But restricting the size of a model is a crude solution, since you'll never be able to learn a more complex model, even if the data support it.
- **Another approach:** keep the model flexible, but **regularize** it
 - **Regularizer:** a function that quantifies how much we prefer one hypothesis vs another

L² Regularization

Observation:

polynomials that overfit often have large coefficients



$$y = 0.1x^5 + 0.2x^4 + 0.75x^3 - x^2 - 2x + 2$$

$$y = -7.2x^5 + 10.4x^4 + 24.5x^3 - 37.9x^2 - 3.6x + 12$$

L² Regularization

- **Another reason** we want parameters (weights) to be small:
 - Suppose inputs x_1 and x_2 are nearly identical for all training examples. The following two hypotheses make nearly the same predictions:


$$\theta_1 = \begin{pmatrix} 1 \\ 1 \end{pmatrix} \qquad \theta_2 = \begin{pmatrix} -9 \\ 11 \end{pmatrix}$$

- But the second model might make weird predictions if the test distribution is slightly different (e.g. x_1 and x_2 match less closely).

L² Regularization

- We can **encourage the parameters to be small** by adding an L^2 penalty (regularizer) to our cost function:

$$\frac{1}{2} \|\boldsymbol{\theta}\|^2 = \frac{1}{2} \sum_j \theta_j^2$$


 hyperparameter to be tuned

$$\mathcal{J}_{reg} = \mathcal{J} + \frac{\lambda}{2} \sum_j \theta_j^2$$

- The **regularized cost function** makes a tradeoff between fit to the data and the norm of the weights vector.

L² Regularization

➤ The geometric picture:



L² Regularization

$$\mathcal{J}(w, b) = \frac{1}{2N} \sum_{i=1}^N \left(y^{(i)} - t^{(i)} \right)^2$$

- We can encourage the weights to be small by choosing as our regularizer the L^2 penalty.

$$\mathcal{R}(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 = \frac{1}{2} \sum_j w_j^2.$$

- Note: to be pedantic, the L^2 norm is Euclidean distance, so we're really regularizing the *squared* L^2 norm.
- The regularized cost function makes a tradeoff between fit to the data and the norm of the weights.

$$\mathcal{J}_{\text{reg}} = \mathcal{J} + \lambda \mathcal{R} = \mathcal{J} + \frac{\lambda}{2} \sum_j w_j^2$$

- Here, λ is a hyperparameter that we can tune using a validation set.

L² Regularization

- Recall the gradient descent update:


$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial \mathcal{J}}{\partial \mathbf{w}}$$

- The gradient descent update of the regularized cost has an interesting interpretation as **weight decay**:


$$\begin{aligned}\mathbf{w} &\leftarrow \mathbf{w} - \alpha \left(\frac{\partial \mathcal{J}}{\partial \mathbf{w}} + \lambda \frac{\partial \mathcal{R}}{\partial \mathbf{w}} \right) \\ &= \mathbf{w} - \alpha \left(\frac{\partial \mathcal{J}}{\partial \mathbf{w}} + \lambda \mathbf{w} \right) \\ &= (1 - \alpha \lambda) \mathbf{w} - \alpha \frac{\partial \mathcal{J}}{\partial \mathbf{w}}\end{aligned}$$

L^1 vs L^2 Regularization

- The L^1 norm (or sum of absolute values) is another regularizer that encourages weights to be exactly zero.




Source: Bishop PRML



L^1 -regularization tends to push parameters to zero

Generalization

➤ Training and test error as a function of regularization parameter λ :



	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
θ_0	0.35	0.35	0.13
θ_1	232.37	4.74	-0.05
θ_2	-5321.83	-0.77	-0.06
θ_3	48568.31	-31.97	-0.05
θ_4	-231639.30	-3.89	-0.03
θ_5	640042.26	55.28	-0.02
θ_6	-1061800.52	41.32	-0.01
θ_7	1042400.18	-45.95	-0.00
θ_8	-557682.99	-91.53	0.00
θ_9	125201.43	72.68	0.01

Ineffective Batch Size

➤ Q: What happens if the batch size is too small? Too large?


➤ **Too large:**

- Computationally expensive
- Average loss might not change very much as batch size grows


➤ **Too small:**

- We optimize a (possibly very) different loss function at each iteration
- Noisy
- SGD may actually take longer


Gradient Descent



Stochastic Gradient Descent




Mini-Batch Gradient Descent



Classification

Overview

- **Classification:** predicting a discrete-valued target
- **Binary classification:** number of target values is 2 (binary-valued)
- Examples:
 - predict where a patient has a disease given presence or absence of various symptoms
 - classify e-mails as spam or non-spam
 - predict whether a financial transaction is fraudulent



Binary Classification

- We can start with our linear function of \mathbf{x} , but now we introduce a threshold :

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x} + b$$

$$\hat{y} = \begin{cases} 1 & \text{if } z \geq r \\ 0 & \text{if } z < r \end{cases}$$

Binary Classification

- Eliminating the threshold

- Can make the threshold $r = 0$ without loss of generality.

$$\boldsymbol{\theta}^T \mathbf{x} + b \geq 0 \quad \longleftarrow \quad \boldsymbol{\theta}^T \mathbf{x} + \underbrace{b - r}_{\triangleq b'} \geq 0$$

- Eliminating the bias


- Add a dummy feature x_0 which always takes the value 1. The weight θ_0 is equivalent to the bias.

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x} \qquad \hat{y} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

The Geometric Picture

- Here we're visualizing the logical NOT example
- Training examples are points
- Hypothesis are half-spaces whose boundaries pass through the origin
- This boundary is the decision boundary
 - In 2-D it's a line, but think of it as a hyperplane
- If the training examples can be separated by a linear decision rule, they are linearly separable.

➤ Input Space, or Data Space:




NOT

x_0	x_1	t
1	0	1
1	1	0


The Geometric Picture

- Hypotheses are points
- Training examples are half-spaces whose boundaries pass through the origin
- The region satisfying all the constraints is the feasible region; if this region is nonempty, the problem is feasible

How do you optimize?





Weight Space



$$w_0 > 0$$

$$w_0 + w_1 < 0$$

Example 1

NOT		
 w_0	 w_1	
x_0	x_1	t
1	0	1
1	1	0

Example 2

AND

x_0	x_1	x_2	t
1	0	0	0
1	0	1	0
1	1	0	0
1	1	1	1

How do you optimize?



Logistic Regression

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

- Because the prediction $\hat{y} \in [0, 1]$, we can interpret it as the estimated probability that the label is positive ($y = 1$).
- **Cross-entropy loss captures** this intuition:

$$\begin{aligned}\mathcal{L}_{CE}(\hat{y}, y) &= \begin{cases} -\log \hat{y} & \text{if } y = 1 \\ -\log(1 - \hat{y}) & \text{if } y = 0 \end{cases} \\ &= -y \log \hat{y} - (1 - y) \log(1 - \hat{y})\end{aligned}$$





Logistic Regression with Cross-entropy Loss

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

- Being 99% confident in the wrong answer is much more wrong than being only 90% confident.

$$\mathcal{L}_{CE}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$



Logistic Regression with Cross-entropy Loss

➤ Computation Problem:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- If $\hat{y} = \sigma(z)$ gets too close to 0 or 1, it may cause very subtle and hard-to-find bugs:

$$\hat{y} = \sigma(z) \Rightarrow \hat{y} \approx 0$$

$$\mathcal{L}_{CE}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y}) \Rightarrow \text{computes } \log 0$$

- Instead, we combine the activation function and the loss into a single logistic-cross-entropy function.

$$\mathcal{L}_{LCE}(\sigma(z), y) = y \log(1 + e^{-z}) + (1 - y) \log(1 + e^z)$$

Example:

$$\mathcal{L}_{CE}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\mathcal{L}_{LCE}(\sigma(z), y) = y \log(1 + e^{-z}) + (1 - y) \log(1 + e^z)$$

Weight Updates

- Comparison of gradient descent updates:

- Linear regression:

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \frac{\alpha}{N} \sum_{i=1}^N (\hat{y}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

- Logistic regression:


$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \frac{\alpha}{N} \sum_{i=1}^N (\hat{y}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$$

$$z = \boldsymbol{\theta}^T \mathbf{x} + b$$

Using activation function

$$\hat{y} = \frac{1}{1 + e^{-z}}$$

Loss Summary



Multiclass Classification

Multiclass Classification


- What about classification task with more than two categories?
- Targets form a discrete set $\{1, \dots, K\}$
- It's often more convenient to represent them as one-hot vectors or a one-of-K encoding:

$$y = [0, \dots, 0, 1, 0, \dots, 0]$$

entry k is a 1, all others are 0



4 → 4 2 → 2 3 → 3
4 → 4 9 → 9 0 → 0
5 → 5 7 → 7 1 → 1
9 → 9 0 → 0 3 → 3
6 → 6 7 → 7 4 → 4



Multiclass Classification

- Now there are D input dimensions and K output dimensions, so we need $K \times D$ parameters, which we arrange as a matrix $\boldsymbol{\theta}$.
- Also, we have a K -dimensional vector \mathbf{b} of biases.
- Linear predictions:

$$z_k = \sum_j \theta_{kj} x_j + b_k$$

- Vectorized:

$$\mathbf{z} = \boldsymbol{\theta} \mathbf{x} + \mathbf{b}$$

Activation Function

- A natural activation function to use is the softmax function, a multivariable generalization of the logistic (sigmoid) function:

$$\hat{y}_k = \text{softmax}(z_1, \dots, z_K)_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$$

Softmax makes differences larger – pushes values close to 1 to 1, values close to 0 to 0

- The input are called the logits.
- Outputs are positive and sum to 1 (interpreted as probabilities)

Loss Function

- If a model outputs a vector of class probabilities, we can use cross-entropy as the loss function:

Binary

$$\mathcal{L}_{CE}(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

Multi-class

$$\begin{aligned}\mathcal{L}_{CE}(\hat{\mathbf{y}}, \mathbf{y}) &= - \sum_{k=1}^K y_k \log \hat{y}_k \\ &= -\mathbf{y}^T \log \hat{\mathbf{y}}\end{aligned}$$

Softmax Regression

➤ Softmax regression:

$$\mathbf{z} = \boldsymbol{\theta}\mathbf{x} + \mathbf{b}$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{z})$$


$$\mathcal{L}_{CE}(\hat{\mathbf{y}}, \mathbf{y}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

➤ Gradient descent updates:


$$\frac{\partial \mathcal{L}_{CE}}{\partial \mathbf{z}} = \hat{\mathbf{y}} - \mathbf{y}$$

Summary

Classification



Regression



Practice in Google Colab

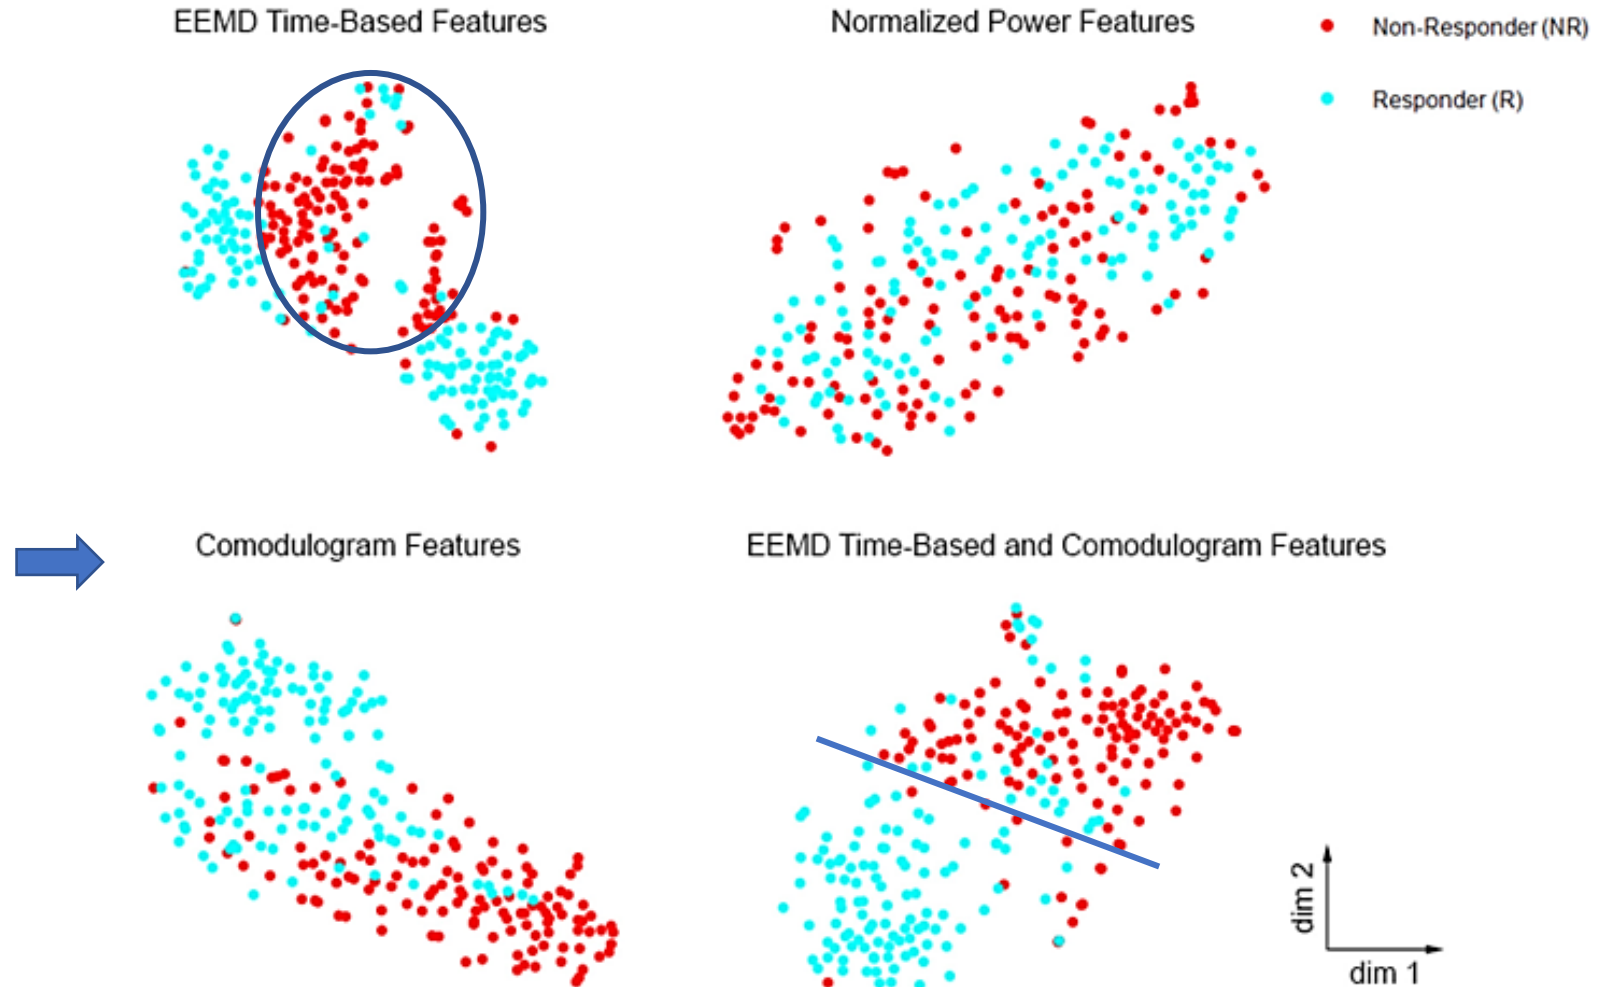

Neural Networks

Challenges with Feature Maps

- Feature maps can be useful for solving nonlinear regression and classification problems.
- Have several limitations:
 - The feature maps must be selected in advance
 - Not always easy to pick a good feature map and can take a long time to craft
 - In high dimensions the feature representations can explode

Example of Feature Engineering:

Objective: Predict **responders** from **non-responders** given raw electroencephalogram data.




Challenges with Feature Maps

We need an algorithm that can learn good features for nonlinear regression and classification.

Motivating Example: Tumor Classification

Objective: Classify an image of a biopsy as cancerous (malignant) or benign (tumor)


- Pathologists/radiologists train for years to do this!
- How would you solve this problem?



Source: [Levenson et al., 2015](#)

Maybe We Can Use Pigeons

- **Train a Pigeon!**
- A 2015 study suggests that the common pigeon can reliably distinguish between benign versus malignant tumors and, in doing so, could help researchers develop better cancer screening technologies.



Source: [Levenson et al., 2015](#)

Maybe We Can Use Pigeons

Training Algorithm:

1. Show an image of a magnified biopsy to pigeon
2. Pigeon pecks at one of two answer buttons on sides for malignant/benign
3. Correct classifications are rewarded with food pellets





Video: <https://www.youtube.com/watch?v=flzGjnJLyS0>

Why are we talking about pigeons?

We need to answer similar questions in training a pigeon/artificial neural network:

- How will we reward correct responses?
- How do we train the neural network efficiently?
- How do we know the pigeon **didn't just memorize** the images we showed it?
- What are the ethics of trusting a pigeon to detect cancer?


How do Pigeons Work?



Source: [Letsmaketech](https://www.lets maketech.com)


Simplified Neuron Anatomy

- **Dendrites:** are connected to other cells that provide information.
- **Cell body:** consolidates information from dendrites.
- **Axon:** an extension from the cell body that passes information to other neurons.
- **Synapse:** the area where the axon of one neuron and the dendrite of another connect.

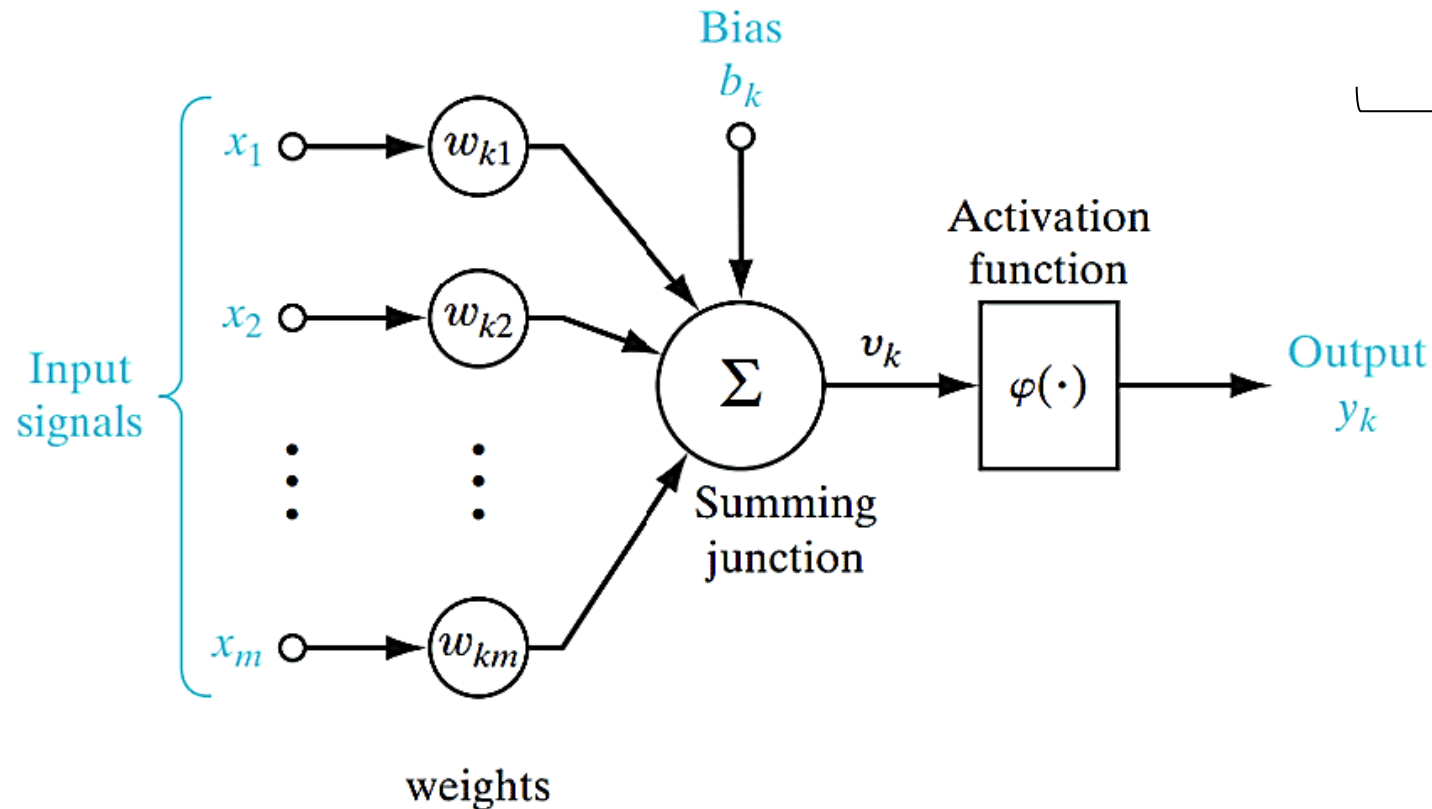


Artificial Neural Network

- Maybe we're not quite ready for pigeon doctors, but we can use the next best thing... an artificial pigeon (**artificial neural network**)



Artificial Neural Network




$$\hat{y} = \varphi \left(\underbrace{\sum_{i=0}^p w_i x_i = \mathbf{w}^T \mathbf{x}} \right)$$

does this look familiar?

Artificial Neural Network

$$y = \varphi \left(\sum_{i=0}^p w_i x_i = \mathbf{w}^T \mathbf{x} + \mathbf{b} \right)$$



Logistic Regression!

What will we have if we exclude the activation function?

Training / Learning Parameters

- In order to train an ANN we should define the error on our predictions.
- This is the same as with linear regression and logistic regression:

Mean Squared Error
(regression)

$$\frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|^2$$


Cross-Entropy Loss
(classification)

$$-\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K y_{n,k} \log(\hat{y}_{n,k})$$

implement gradient descent to learn parameters!


Gradients

- For this simple version of neural networks all the gradient calculations are the same as we've seen earlier...
- It probably wouldn't be a surprise to find out that to train this simple ANN is a convex problem (not true for all ANNs).




parameters (w) can be multidimensional

Achieving Nonlinearity






Add More Neurons?



This is a 2-layer neural network. We do not count the input layer, so the number of layers equal number of sets of weights and biases.

2-layer ANN with Linear Activation




Expressive Power

- Adding more linear layers does not help increase the capacity of the model.
- Any sequence of linear layers can be equivalently represented with a single linear layer.



$$\hat{\mathbf{y}} = \underbrace{\mathbf{W}^{(1)}\mathbf{W}^{(2)}\mathbf{W}^{(3)}}_{\mathbf{W}'}\mathbf{x}$$

- Deep linear networks are no more expressive than **linear regression!**




Need an Activation Function



- Neural Networks can be viewed as a **way of learning features**
- The goal being that the final layer is presented with **linearly separable feature data**




2-layer ANN with Nonlinear Activation



Expressive Power

- Multilayer feed-forward neural nets with nonlinear activation functions are **universal function approximators**.
- They can approximate any function arbitrarily well, but this comes at a cost...

Cost Function is Non-Convex!



Nonlinear
activations
introduce
non-convex
surface!


Tuning Neural Networks

Changing ANN Architecture:

- Number of hidden units
- Weights
- Activation Functions


Applying Different Training Techniques:

- Number of iterations
- Learning rate and adaptive learning rate
- Momentum
- Batching of Data
- Regularization
- Dropout
- Feature Augmentation
- Many more...



Take-home Exercise: Discuss on Piazza

Q: Determine the gradients for a 2-layer artificial neural network with sigmoid activation on the hidden and output layers. The error is computed using mean squared error loss.



Next Time

- Week 10 Q&A session: Thursday and Friday
 - Project 4 is due on April 1st
- Week 11 Lecture – Deep Learning (and More)
 - Neural Network Architectures
 - Automatic Differentiation
 - Discrete Optimization