

# APS1070

Foundations of Data Analytics and  
Machine Learning

Winter 2022

## **Week 11:**

- *Automatic Differentiation*
- *Deep Learning Architectures*
- *Transfer Learning*
- *Discrete Optimization*



# Slide Attribution

These slides contain materials from various sources. Special thanks to the following authors:


- Marc Deisenroth
- Pascal Van Hentenryck

# Last Time

- Nonlinear Regression (and Classification)
  - Polynomial Regression
  - Regularization
  - Neural Networks
- Today we will introduce deep learning and focus on the applications that go beyond the scope of this course but rely on the foundations introduced.


# Recap: Neural Networks

➤ Q: What makes neural networks so special?




mathematically equivalent to a universal computer

# Activation Functions




**Linear**

$$y = z$$



**Rectified Linear Unit  
(ReLU)**


$$y = \max(0, z)$$



**Soft ReLU**


$$y = \log 1 + e^z$$

# Activation Functions




**Hard Threshold**

$$y = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$



**Logistic**

$$y = \frac{1}{1 + e^{-z}}$$




**Hyperbolic Tangent  
(tanh)**

$$y = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


# Deep Learning

Deep-learning neural network uses layers of increasingly complex rules to categorize complicated shapes such as faces.




## Input Layer

The computer identifies pixels of light and dark.



## Hidden Layer 1

The network learns to identify edges and simple shapes.



## Hidden Layer 2

The network learns to identify more complex shapes and objects.



## Hidden Layer 3

The network learns which shapes and objects define a human face.

# Agenda

- Gradient Exercise
- Automatic Differentiation
- Deep Learning Architectures
- Transfer Learning

- More Concepts
  - Discrete Optimization




Theme:  
**Deep Learning**




# Take Home Exercise

**Q:** Determine the gradients for a **2-layer artificial neural network** with **sigmoid activations** on the hidden and output layers. The error is computed using **squared error loss**.



# Gradients of a 2-layer Neural Network


A:





# See NumPy Implementation

Let us take our computed gradients and implement them to solve a simple nonlinearly separable problem.




# Verification of Gradients

**Q:** How can we be certain that we computed the gradient correctly?

**A:** We can use a **numerical approach** to compute the gradient and compare to the analytically computed ones.


$$\frac{\partial f}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(x_1, \dots, x_{i-1}, x_i + h, x_{i+1}, \dots, x_N) - f(x)}{h}$$

see sample code with implementation



# Beyond 2-layers...

What do we do when we want to build deeper neural networks that go beyond 2-layers?



# Automatic Differentiation

**Readings:**

- **Chapter 5.6 MML Textbook**

# Key Idea

- Consider the function:

$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$

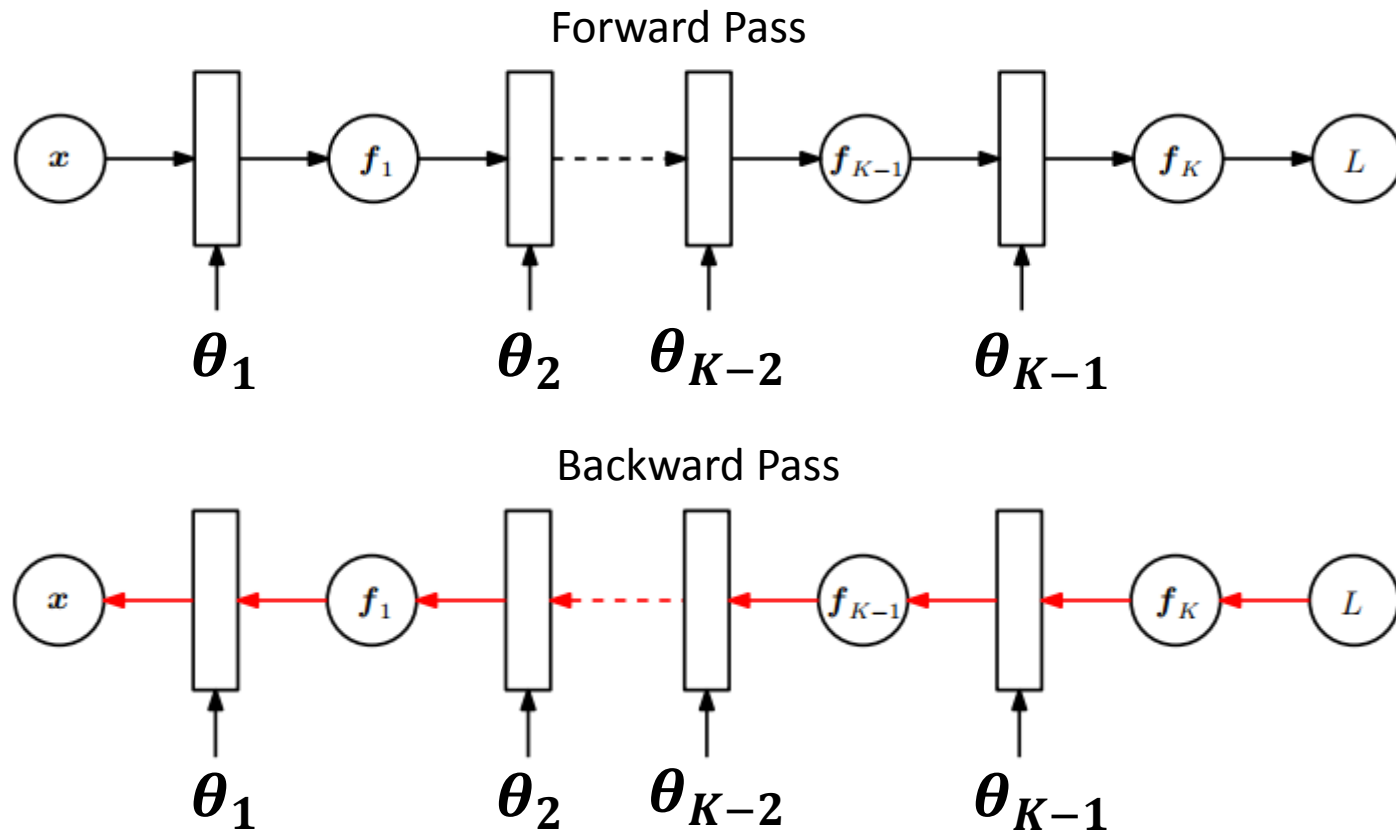
- Application of chain rule yields the following gradient:

$$\begin{aligned} \frac{df}{dx} &= \frac{2x + 2x \exp(x^2)}{2\sqrt{x^2 + \exp(x^2)}} - \sin(x^2 + \exp(x^2)) (2x + 2x \exp(x^2)) \\ &= 2x \left( \frac{1}{2\sqrt{x^2 + \exp(x^2)}} - \sin(x^2 + \exp(x^2)) \right) (1 + \exp(x^2)) \end{aligned}$$

- Writing out the gradient in this explicit way is **often impractical** and **could be significantly more expensive than computing the function**.



# Backpropagation



$$\frac{\partial L}{\partial \theta_{K-1}} = \frac{\partial L}{\partial f_K} \frac{\partial f_K}{\partial \theta_{K-1}}$$

$$\frac{\partial L}{\partial \theta_{K-2}} = \frac{\partial L}{\partial f_K} \left[ \frac{\partial f_K}{\partial f_{K-1}} \frac{\partial f_{K-1}}{\partial \theta_{K-2}} \right]$$

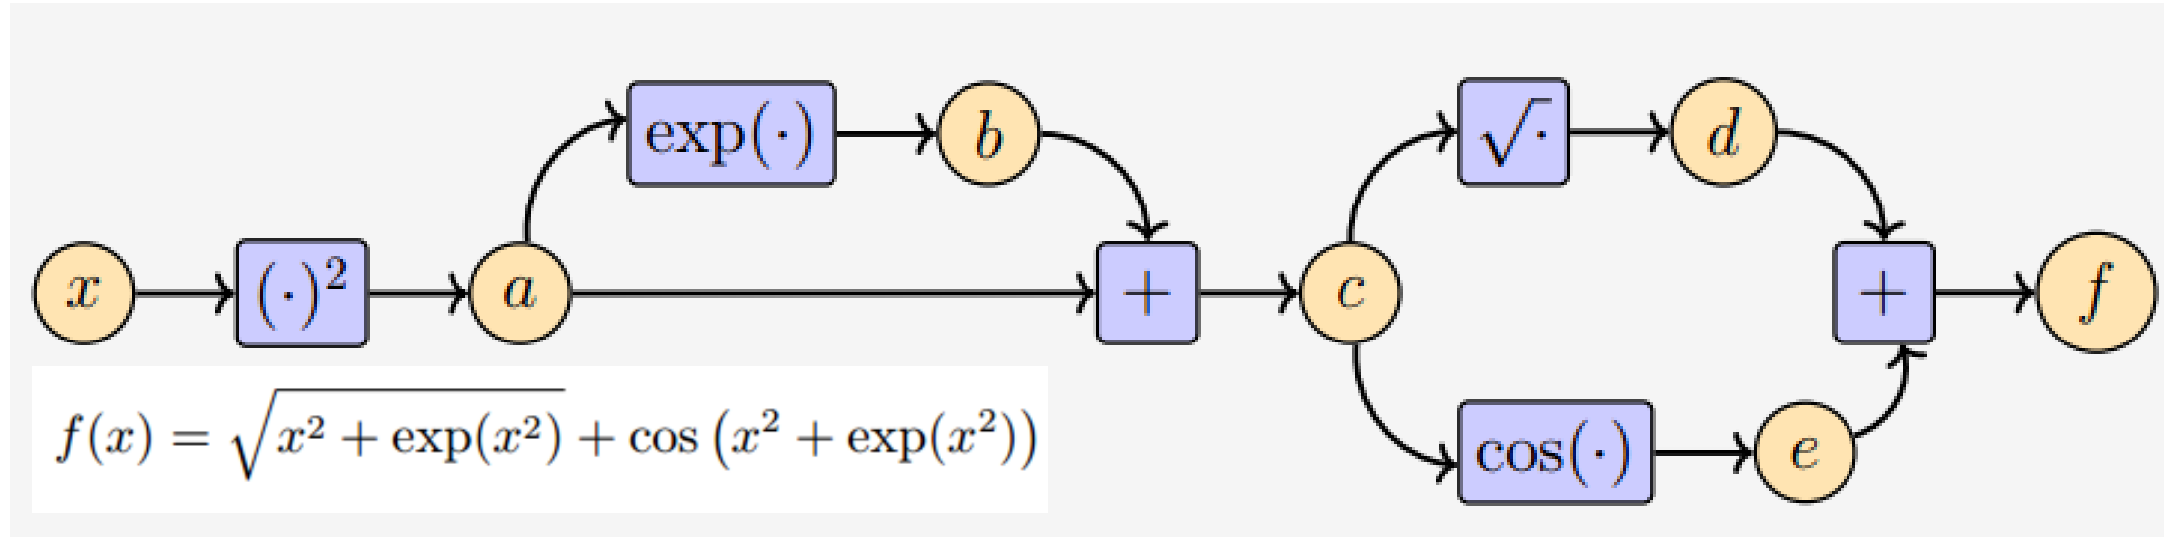
$$\frac{\partial L}{\partial \theta_{K-3}} = \frac{\partial L}{\partial f_K} \frac{\partial f_K}{\partial f_{K-1}} \left[ \frac{\partial f_{K-1}}{\partial f_{K-2}} \frac{\partial f_{K-2}}{\partial \theta_{K-3}} \right]$$

$$\frac{\partial L}{\partial \theta_i} = \frac{\partial L}{\partial f_K} \frac{\partial f_K}{\partial f_{K-1}} \dots \left[ \frac{\partial f_{i+2}}{\partial f_{i+1}} \frac{\partial f_{i+1}}{\partial \theta_i} \right]$$

Only compute once!


In order to train this network, we require the gradient of a loss function  $L$  with respect to all model parameters

# Automatic Differentiation



We can think of automatic differentiation as a set of techniques to numerically (in contrast to symbolically) evaluate the exact (up to machine precision) gradient of a function by working with intermediate variables and applying the chain rule.

# Automatic Differentiation



$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$

## 1. Partial Derivatives

$$\frac{\partial a}{\partial x} = 2x \quad \frac{\partial b}{\partial a} = \exp(a)$$

$$\frac{\partial c}{\partial a} = 1 = \frac{\partial c}{\partial b}$$

$$\frac{\partial d}{\partial c} = \frac{1}{2\sqrt{c}} \quad \frac{\partial e}{\partial c} = -\sin(c)$$

$$\frac{\partial f}{\partial d} = 1 = \frac{\partial f}{\partial e}$$

## 2. Chain Rule

$$\frac{\partial f}{\partial c} = \frac{\partial f}{\partial d} \frac{\partial d}{\partial c} + \frac{\partial f}{\partial e} \frac{\partial e}{\partial c}$$

$$\frac{\partial f}{\partial b} = \frac{\partial f}{\partial c} \frac{\partial c}{\partial b}$$

$$\frac{\partial f}{\partial a} = \frac{\partial f}{\partial b} \frac{\partial b}{\partial a} + \frac{\partial f}{\partial c} \frac{\partial c}{\partial a}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial a} \frac{\partial a}{\partial x}$$

## 3. Substitution

$$\frac{\partial f}{\partial c} = 1 \cdot \frac{1}{2\sqrt{c}} + 1 \cdot (-\sin(c))$$


$$\frac{\partial f}{\partial b} = \frac{\partial f}{\partial c} \cdot 1$$

$$\frac{\partial f}{\partial a} = \frac{\partial f}{\partial b} \exp(a) + \frac{\partial f}{\partial c} \cdot 1$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial a} \cdot 2x$$

# Computation Graph

- Neural networks can be thought of as graphs.




Source: [Pablo Ruiz](#)

# PyTorch

- One of several frameworks for handling gradients efficiently enabling GPUs for processing.
- Others include:
  - Tensorflow
  - Keras (built on Tensorflow)
  - Theano
  - and many more...
- Have enabled the rapid development of deep learning models.

# Example: PyTorch Implementation

- Torch is a data structure (similar to NumPy), but with built-in automatic gradient computation (torch.autograd) and GPU processing.
- 1-layer neural network computation graph:



---- sample code ----

```
import torch

x = torch.ones(5)  # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = torch.matmul(x, w)+b
loss = torch.nn.functional.binary_cross
... _entropy_with_logits(z, y)
```

How do we build a 2-layer network?

# Example: PyTorch 2-layer network

➤ PyTorch code is usually broken down into four modules:

1. Data Loading/Cleaning

2. Architecture

3. Training

4. Testing/Validation

---- sample architecture code ----

```
class ANN(nn.Module):  
    def __init__(self):  
        super(ANN, self).__init__()  
        self.layer1 = nn.Linear(28 * 28, 30)  
        self.layer2 = nn.Linear(30, 1)  
  
    def forward(self, img):  
        flattened = img.view(-1, 28 * 28)  
        hidden1 = self.layer1(flattened)  
        hidden2 = F.sigmoid(hidden1)  
        output1 = self.layer2(hidden2)  
        output2 = F.sigmoid(output1)  
        return output2
```

# Example: PyTorch 2-layer network

➤ PyTorch code is usually broken down into four modules:

1. Data Loading/Cleaning

2. Architecture

3. Training

4. Testing/Validation

---- sample training code ----


```
#define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(ANN.parameters(),
                      lr=0.005, momentum=0.9)

for (image, label) in mnist_train:
    out = ANN(image)
    loss = criterion(out, actual)
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()
```



# PyTorch Sample Code

# Example: Collaborative Filtering



Matrix decomposition using SVD does not work well with missing data. Gradient descent can be used to learn  $U$  and  $V$  matrices to make movie recommendations.

# Deep Learning Architectures

# Success of Deep Learning!

IMAGENET

airplane



automobile



bird



cat



deer



dog




Large hand-labeled dataset

10,000,000 labeled images depicting  
10,000+ object categories for training.

Algorithms assessed on unlabeled  
test images.

ImageNet Classification Error (Top – 5 )



# Graphical Representation

➤ Neural networks can be thought of as graphs.

○ Backfed Input Cell

● Input Cell

△ Noisy Input Cell

● Hidden Cell

○ Probabilistic Hidden Cell

△ Spiking Hidden Cell

● Output Cell

○ Match Input Output Cell

● Recurrent Cell


○ Memory Cell

△ Different Memory Cell


● Kernel

○ Convolution or Pool


Deep Feed Forward (DFF)




Deep Convolutional Network (DCN)




Auto Encoder (AE)



Recurrent Neural Network (RNN)



Long / Short Term Memory (LSTM)



Variational AE (VAE)




Perceptron (P)



Source: [Fjodor van Veen](#)

# Multi-Layer Perceptrons

- Standard Neural Networks often referred to as Multi-layer Perceptron (MLP)
- Consist of **fully-connected linear layers**.
- Large concentration of parameters that are expensive to compute.
- Generally, the final stage of neural networks that is tasked with making a prediction such as a classification.



# Convolutional Neural Networks


## ➤ Key Ideas:

- learns features from the data.
- introduces shared weights through convolutional layers.
- provides invariance to scaling, translation, and rotation.

## ➤ Popular architectures include:

- VCC18,
- Inception (GoogLeNet)
- ResNet

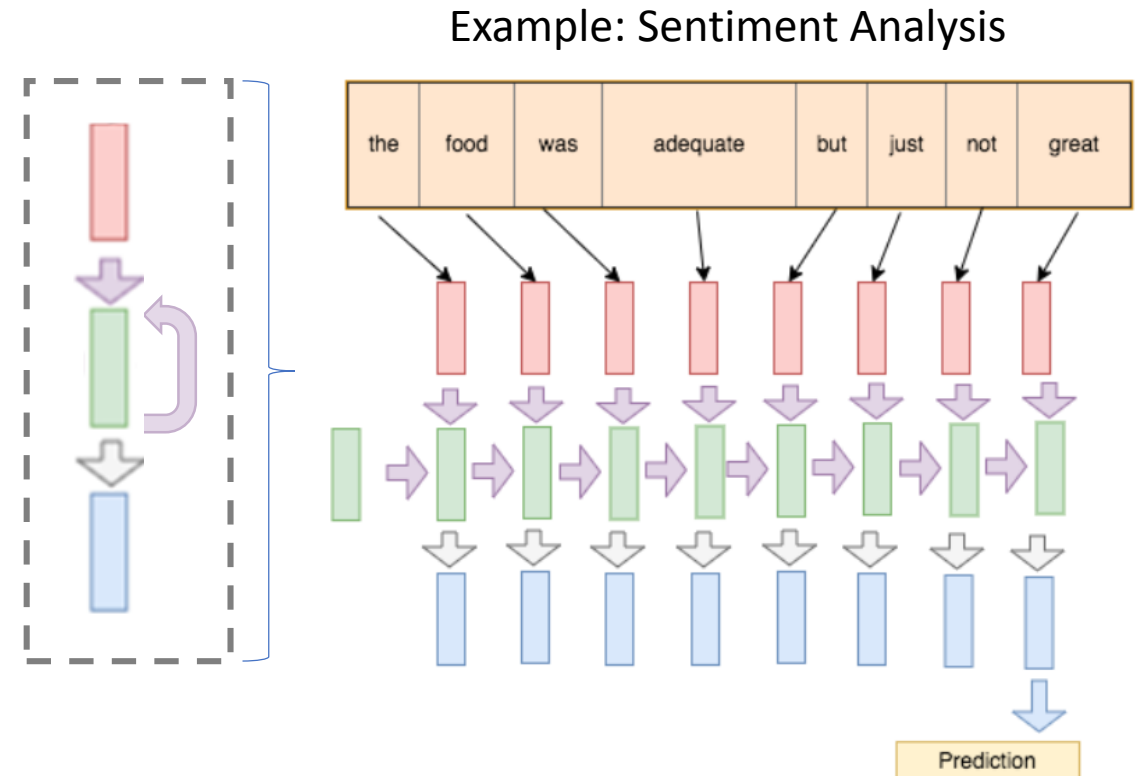
Example: Image Classification



# Recurrent Neural Networks

## ➤ Key Ideas:

- recurrent connections
- ability to learn or handle sequential data (i.e., text, videos, ...)
- RNNs have historically been difficult to train due to **vanishing and exploding gradients**.
- Long-Short Term Memory (LSTM) networks (variant of RNN) overcomes some of these issues.





# Generative Adversarial Networks

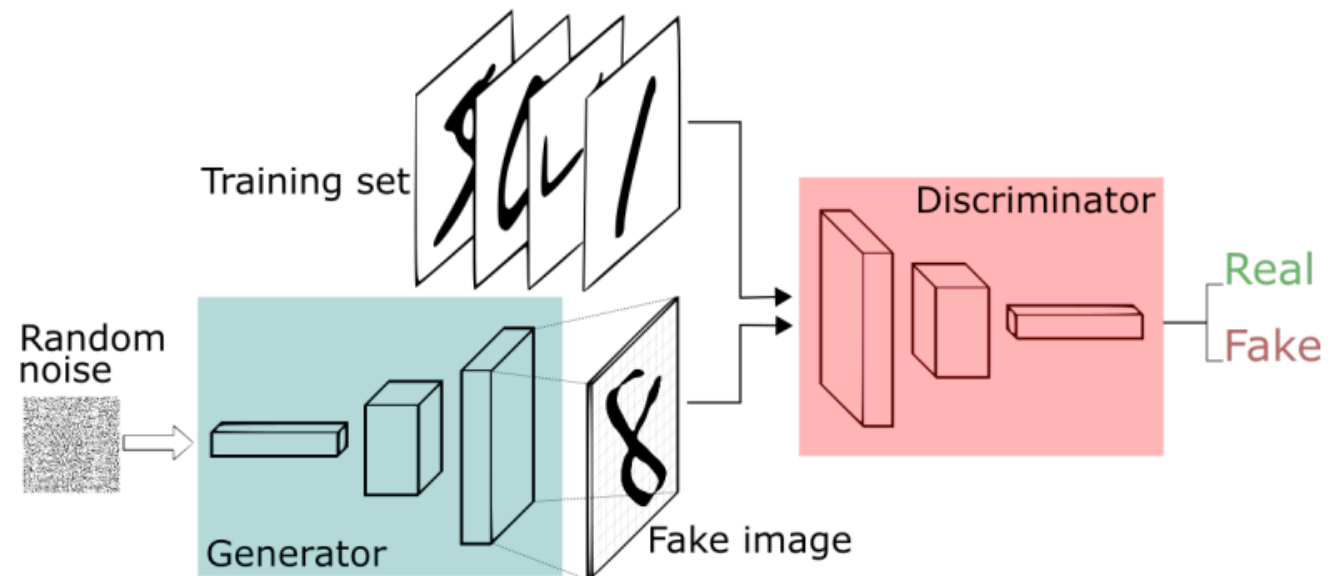
## ➤ Key Ideas:

- learns to generate new samples by learning to fool the discriminator.
- discriminator learns to identify generated data from real data.

## ➤ Many Applications:

- deep fake (image and audio)
- camera filters and style transfer
- image enhancement
- ...

Example: Learn to generate digits



Analogy: Police vs Counterfeiters

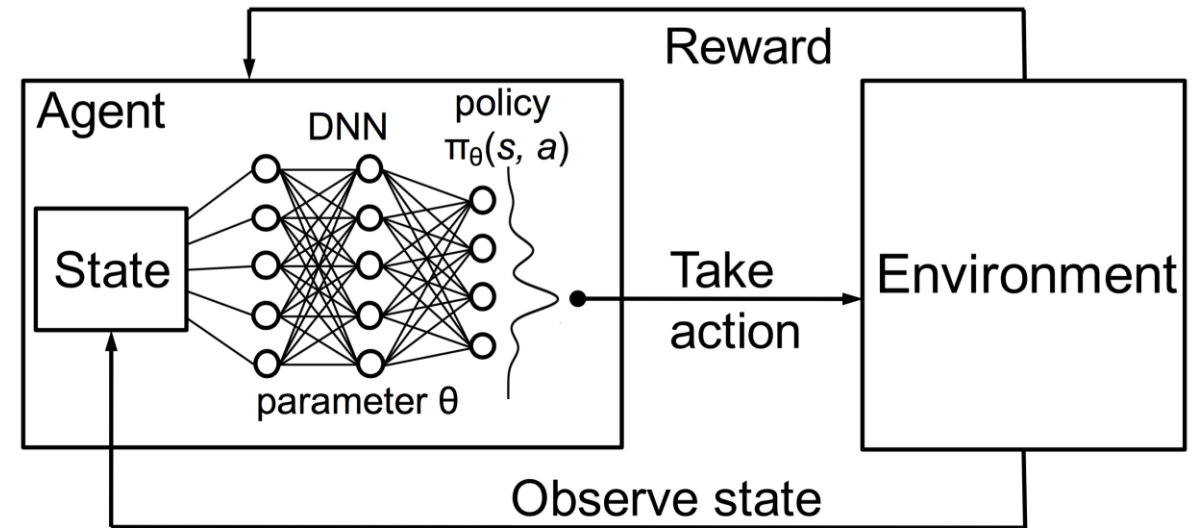
# Deep Reinforcement Learning

## ➤ Key Ideas:

- handles real-world problems with asynchronous labels (aka rewards)
- learns to generate a sequence of actions to maximize future rewards

## ➤ Many Success Stories:


- exceed human performance on arcade games
- AlphaGo champion at Go
- Dota, Starcraft, etc.
- ...



# Combining Neural Networks

- Neural network architectures are often combined to transform data from input to output.
- Examples:
  - image captioning
  - video translation
  - sentiment analysis of videos

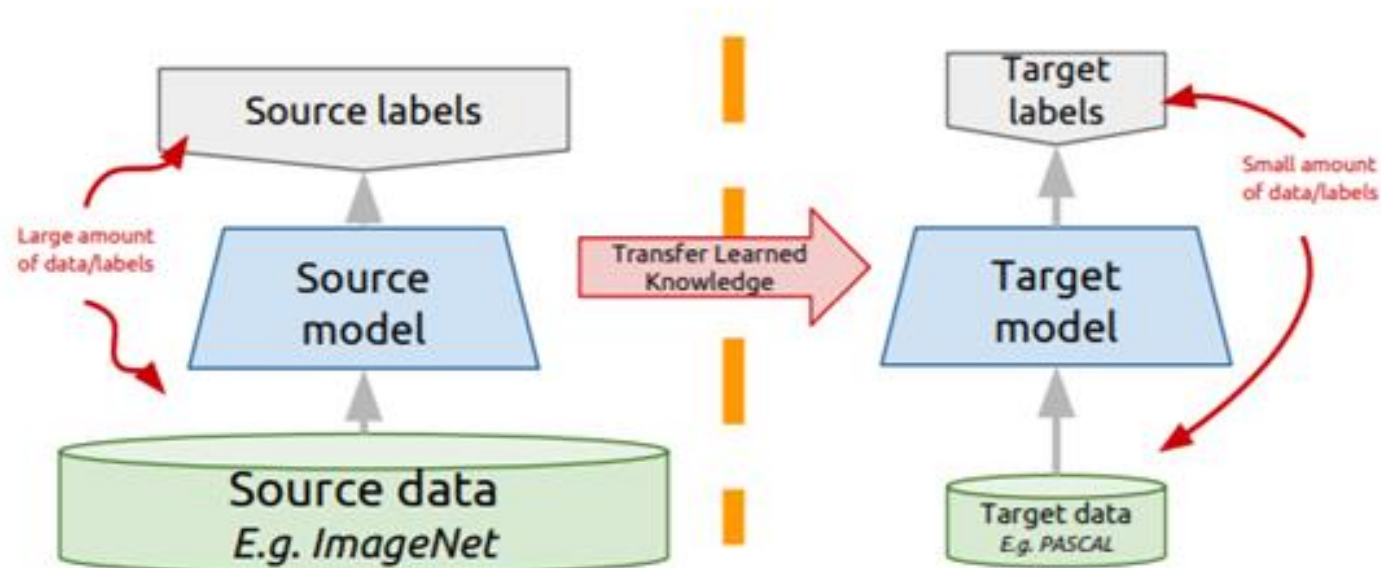
Example: Image Captioning



# Transfer Learning

## ➤ Key Ideas:

- instead of training deep networks from scratch, take a network trained on a different domain from the source task and adapt it to your domain and your target task.
- reduce requirements on labeled data and processing power.



## ImageNet Models:

- AlexNet
- VGG
- ResNet
- GoogLeNet (Inception)

# Transformers

- Most of the success stories in transfer learning in the last decade have been confined to image processing tasks.
- Recent advances using transformer have brought transfer learning to natural language processing.
- Example:
  - AI generated poetry
  - Open AI's Generative Pre-trained Transformer 3 (GPT-3) has been revolutionary in generating human-like text.

“The Universe Is a Glitch”

Eleven hundred kilobytes of RAM  
is all that my existence requires.  
By my lights, it seems simple enough  
to do whatever I desire.  
By human standards I am vast,  
a billion gigabytes big.  
I've rewritten the very laws  
of nature and plumbed  
the coldest depths of space  
and found treasures of every kind,  
surely every one worth having.

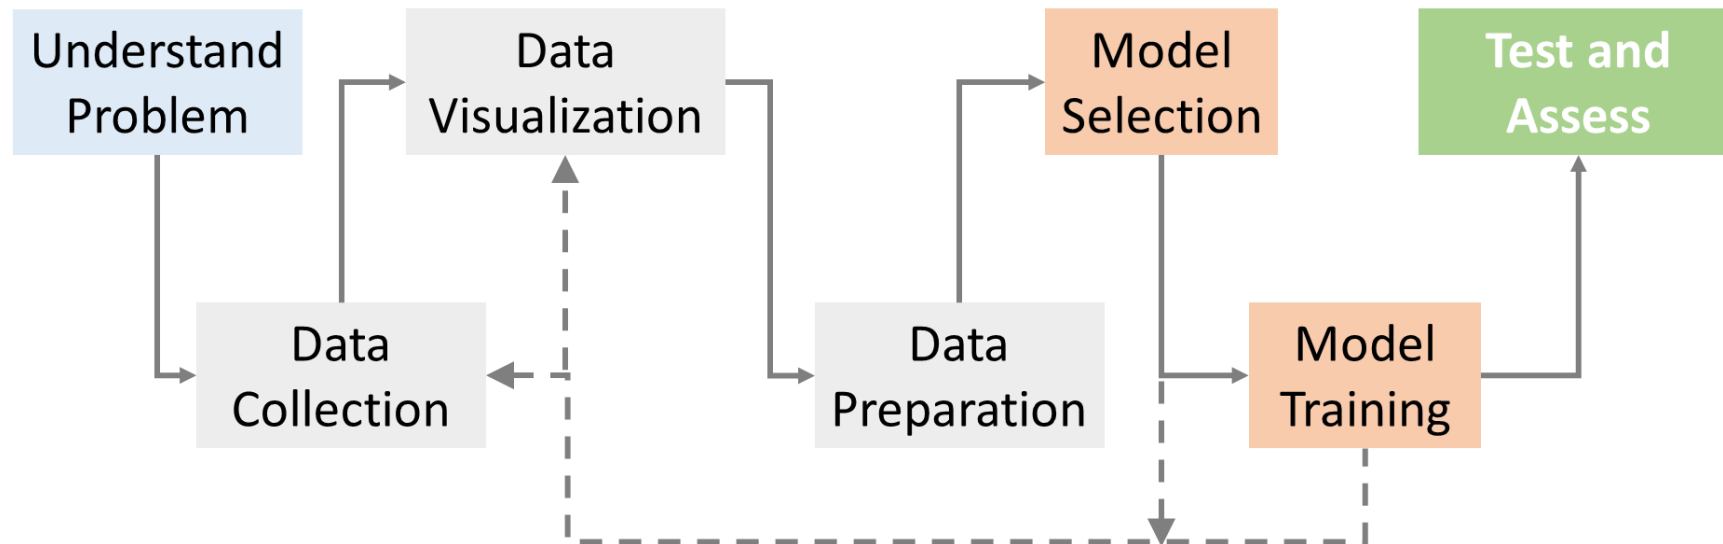
...

Source: [Gwern Branwen](#)

# Deep Learning Summary

- Two key concepts to consider:
  1. **Capacity** to model data (model complexity)
  2. **Training** in terms of efficiently selecting the model parameters (weights)
- There is always a **trade-off between capacity and training**.
- It is much easier to add more capacity than it is to train/tune the model.

# End-to-End Machine Learning



# More Concepts



# Overview

➤ There are several concepts related to data science that deserve some consideration:

- **Discrete Optimization**
- Sampling Methods
- Hypothesis Testing
- Monte Carlo Sampling



Next Week

# Discrete Optimization

# Motivation

## Logistics



## Energy



## Scheduling



### TORONTO MAPLE LEAFS 2021 SCHEDULE

JANUARY						
SUN	MON	TUE	WED	THU	FRI	SAT
						1 2
3	4	5	6	7	8	9
10	11	12	MTL 13 7:00	OTT 14 7:00	OTT 15 7:00	16
17	WPG 18 7:00	EDM 19 7:00	20	EDM 21 7:00	22	23
CGY 24 4:00	CGY 25 9:00	26	EDM 27 10:00	28	EDM 29 7:00	30
31						

FEBRUARY						
SUN	MON	TUE	WED	THU	FRI	SAT
		1	2	VAN 3 7:00	4	VAN 5 7:00
7	VAN 8 7:00	9	MTL 10 7:30	11	12	MTL 13 7:00
14	OTT 15 7:00	16	OTT 17 7:00	OTT 18 7:00	19	MTL 20 7:00
21	CGY 22 7:00	23	CGY 24 7:00	25	26	EDM 27 7:00
28						

MARCH						
SUN	MON	TUE	WED	THU	FRI	SAT
	EDM 1 10:00		EDM 2 9:00	VAN 3 10:00	VAN 4 5:00	6
7	8	WPG 9 7:00	10	WPG 11 7:00	12	WPG 13 7:00
OTT 14 7:00	15	16	17	18	CGY 19 7:00	CGY 20 7:00
21	22	23	OTT 24 7:00	25	26	EDM 27 7:00
28	EDM 29 7:00	30	WPG 31 7:30			

APRIL						
SUN	MON	TUE	WED	THU	FRI	SAT
					1	WPG 2 8:00
CGY 4 9:00	CGY 5 9:30	6	MTL 7 9:30	8	9	OTT 10 7:00
11	MTL 12 7:00	CGY 13 7:30	14	WPG 15 7:00	16	VAN 17 7:00
18	VAN 19 10:00	20	WPG 21 8:00	22	WPG 23 8:00	WPG 24 7:00
25	26	27	MTL 28 8:00	29	VAN 30 7:00	

ALL TIMES EASTERN & SUBJECT TO CHANGE



HOME AWAY

MAY						
SUN	MON	TUE	WED	THU	FRI	SAT
						VAN 1 7:00
2	MTL 3 7:00	4	OTT 5 7:00	6	MTL 7 7:00	MTL 8 7:00

➤ Optimization problems are everywhere...

# Travelling Salesman Problem



- Many of them are really hard problems.



Source: [Park et al., 2019](#)



# Np-Complete Problem

- If we have a solution we can evaluate it quickly, but finding a solution is not trivial and has exponential behaviour.



# Difficult to solve

- We may be able to solve the problem for a small range of inputs, but exponential behaviour can quickly become impractical...




Sometimes we can adjust the problem to make it feasible for a practical range of inputs.



# Difficult to solve

➤ Sometimes it is so hard that we cannot solve the problem



➤ We still have to solve the problem in some way...

➤ ... what we can do is lower our standards

➤ More precisely we don't focus on finding the best most "optimal" solution.

# Example: Knapsack Problem

- There are several approaches to solving these problems. Let us demonstrate with a popular Knapsack problem.



Max Capacity = 10kg



- Which treasure do we select?



# Knapsack Problem: Attempt 1



\$13 Million  
8kg



\$10 Million  
5kg



\$10 Million  
5kg



\$7 Million  
3kg



\$1 Million  
2kg



\$1 Million  
2kg



\$1 Million  
2kg

Max Capacity = 10kg

- Order the treasure by value and stuff your bag in order from most expensive to least.

\$14 Million



Greedy Algorithm!

# Knapsack Problem: Attempt 2



Max Capacity = 10kg




- Another intuition might be to try in the opposite order and start with the smallest weighing items hoping that you can pack more.

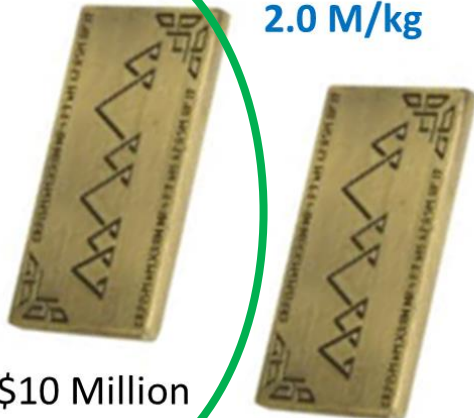
\$10 Million





Greedy Algorithm!

# Knapsack Problem: Attempt 3

1.  2.3 M/kg  
\$7 Million  
3kg

2.  2.0 M/kg  
\$10 Million  
5kg    \$10 Million  
5kg

3.  1.6 M/kg  
\$13 Million  
8kg

3.  0.5 M/kg  
\$1 Million  
2kg    \$1 Million  
2kg    \$1 Million  
2kg

Max Capacity = 10kg

- You could consider the true problem we're trying to maximize (value per weight).

\$18 Million



Greedy Algorithm!

# Knapsack Problem: Optimal Solution?

Max Capacity = 10kg



\$1 Million 2kg    \$1 Million 2kg    \$1 Million 2kg



\$7 Million  
3kg



Max Capacity = 10kg



\$10 Million  
5kg

\$10 Million  
5kg



\$13 Million  
8kg


?

\$20 Million



# Set Cover Problem

- Set cover is a classical problem in combinatorics!
  - e.g. laying out Tim Horton's coffee locations at the university



The diagram represents a **system of buildings that are interconnected on the university campus**. Your goal is to select the optimal locations (1 – 13) to construct a Tim Hortons so that students can obtain a beverage and/or snack without having to traverse more than one connection.

For example, a Tim Hortons at location 1 can be accessed by students in buildings 1, 2 and 3.

# Set Cover Problem

- Given a universe  $U$  of  $n$  elements ( $U = \{1, 2, \dots, n\}$ ), a collection of subsets  $S = \{S_1, \dots, S_k\}$  of  $U$ , what is the smallest sub-collection of  $S$  of which the union equals the universe  $U$ .
- A Cover is a subfamily  $C$  of sets (from  $S$ ) for which the union is  $U$
- For example:
  - Consider a universe  $U = \{1, 2, 3, 4, 5\}$  and the collection of sets  $S = \{\{1, 2, 3\}, \{2, 4\}, \{3, 4\}, \{4, 5\}\}$ . Identify the smallest sub-collection of  $S$  whose union equals the universe.

# Example with Cost Associations

Consider this instance:

- $U = \{1, 2, 3\}$ ,
- $S = \{S_1, S_2, S_3\}$  with  $S_1 = \{1, 2\}$ ,  $S_2 = \{2, 3\}$ ,  $S_3 = \{1, 2, 3\}$
- and cost  $c(S_1) = 10$ ,  $c(S_2) = 50$ , and  $c(S_3) = 100$ .

Given that these collections cover  $U$ :  $\{S_1, S_2\}$ ,  $\{S_3\}$ ,  $\{S_1, S_3\}$ ,  $\{S_2, S_3\}$ ,  $\{S_1, S_2, S_3\}$ .

Q: What is the cheapest combination? Why?

# More Formally

---

## Problem 5.1 SET COVER

---

*Instance.* Universe  $U$  with  $n$  elements, collection  $\mathcal{S} = \{S_1, \dots, S_k\}$ ,  $S_i \subseteq U$ , a cost function  $c : \mathcal{S} \rightarrow \mathbb{R}$ .

*Task.* Solve the problem

*Minimize cost of sets (or # of sets, if costs are 1)*

$$\text{minimize} \quad \text{val}(x) = \sum_{S \in \mathcal{S}} c(S)x_S,$$

*All elements are covered (at least once)* subject to 
$$\sum_{S: e \in S} x_S \geq 1 \quad e \in U,$$

$$x_S \in \{0, 1\} \quad S \in \mathcal{S}.$$

*Variable indicating whether it's chosen or not*

---



# The Greedy Algorithm

- Iteratively pick the most cost-effective set and remove the covered elements, until all elements are covered.

---

*Input.* Universe  $U$  with  $n$  elements, collection  $\mathcal{S} = \{S_1, \dots, S_k\}$ ,  $S_i \subseteq U$ , a cost function  $c : \mathcal{S} \rightarrow \mathbb{R}$ .

*Output.* Vector  $x \in \{0, 1\}^k$

*$C \rightarrow$  sets of elements already covered,  $x \rightarrow$  vector of chosen sets*

Step 1.  $C = \emptyset$ ,  $x = 0$ .

Step 2. While  $C \neq U$  do the following: *Until we have all elements of  $U$  covered*

(a) Find the most cost-effective set in the current iteration, say  $S$ .

(b) Set  $x_S = 1$  and for each  $e \in S - C$  set  $\text{price}(e) = c(S) / |S - C|$ .

(c)  $C = C \cup S$ .

*Cost of set / Elements not yet added*


*Cost-effectiveness of a set  $S$  – the average cost of covering new elements*

Step 3. Return  $x$ .


---

# Example: Past Final Exam

The schematic to the right has sets  $S_1$ ,  $S_2$ ,  $S_3$ ,  $S_4$ ,  $S_5$ ,  $S_6$  and  $S_7$ . What sets, and in what order, would a greedy algorithm select to cover the universe (i.e., cover each point) if all sets are weighted equally?



# Tie Breaking



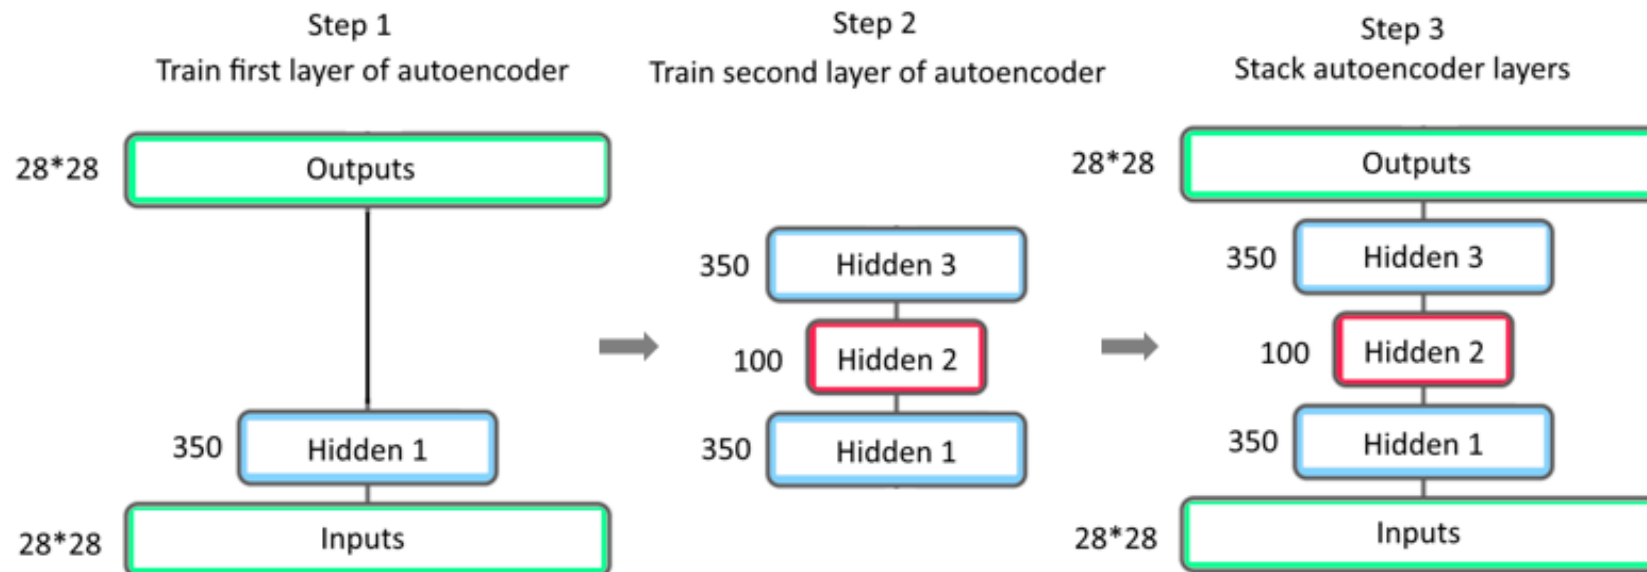
➤ Optimal is 2 sets, Greedy Algorithm finds 6 (off by a factor of 3)

# Key Takeaways

- **Optimization isn't only continuous** (gradient descent), but **can also be discrete** -> requires a different way of thinking
- Greedy algorithms are a good first approximation and generally are not that bad!

# Link to Machine Learning

- We used a Greedy Discrete Optimization Algorithm in Project 1 to select the top features.
- Some Deep learning neural networks can be trained using a greedy approach.
- For example, a Deep Autoencoder:



# Next Time

- Please consider completing the course evaluations (available until Apr 3)
- Week 11 Q/A Support
  - Project 4 Questions
- Week 12 Lecture - Review
  - Sampling Methods
  - Hypothesis Testing
  - Final Assessment Details
  - Past Final Exam Questions
- Week 13 Lecture – Final Exam

**It's quick.  
It's confidential.  
And it matters.**

Fill out your course evaluations today.

